

Lab Assignment 7: Credit Risk Assessment

CSAI2017P: Applied Machine Learning (Lab)

Experiment 9 (Unit VI) • CO2, CO4

Instructor: Dr. Tofik Ali**Due:** announced in the lab**Student Name:** _____**SAP ID:** _____**Batch:** _____**Lab quiz:** LQ7, after submission

Objective

- Build a scoring model where the two error types have very different costs.
- Use ROC and AUC correctly, and know what they hide.
- Examine a model for disparate outcomes across groups.

Datasets used in this assignment

A3 (Statlog German Credit).

Full descriptions, sources and loading snippets for all anchor datasets are in **Anchor-Datasets.pdf**, alongside this brief.

Instructions

- Use a train/test split (80/20 unless stated) and set `random_state` for reproducibility.
- Fit every transformer inside a `Pipeline` so nothing is fitted on the test set.
- Report the metric named in the question. A number without units or context earns no marks.
- Every question asks for a short written observation. Code without commentary caps you at half marks.
- The lab quiz that follows will ask you to read, trace and modify *your own* submitted code.

Submit:

- Notebook: `AML_LA07_<SAPID>.ipynb` with all cells executed and outputs visible.
- Report: 2–3 pages PDF, or a `README.md`, carrying your results table and observations.
- Do not upload datasets. Cite the source and include the loading code.

Section	Level	Choice rule	Marks
A	Easy	Attempt any 2 of 3 (2 marks each)	4
B	Medium	Attempt any 1 of 2 (4 marks each)	4
C	Hard	Compulsory	2
Total			10

Course outcomes assessed by this assignment: **CO2, CO4**. Attempting more than the choice rule allows is not penalised; the best attempts are marked.

Section A (Easy) — attempt any 2 of 3 ($2 \times 2 = 4$)**A1. Load and baseline****[2 marks]***Dataset: A3 (credit-g). Target: class.*

- (a) Fetch the dataset, report its shape and the proportion of bad-credit cases.
- (b) Fit a logistic regression in a pipeline with one-hot encoding and scaling; report accuracy.
- (c) State the majority-class accuracy and compare in two lines.

A2. The cost matrix**[2 marks]***Dataset: A3.*

- (a) Print the confusion matrix for your A1 model.
- (b) The dataset's own documentation says lending to a bad customer costs five times as much as refusing a good one. Compute your model's total cost under that matrix.
- (c) State in two lines whether accuracy was the right metric.

A3. ROC and AUC**[2 marks]***Dataset: A3.*

- (a) Plot the ROC curve and report AUC.
- (b) Mark the operating point that minimises the cost from A2.
- (c) State in two lines what AUC measures that accuracy does not.

Section B (Medium) — attempt any 1 of 2 ($1 \times 4 = 4$)**B1. Cost-sensitive training****[4 marks]***Dataset: A3.*

- (a) Retrain with `class_weight={good:1, bad:5}` and separately by tuning the decision threshold on a validation split.
- (b) Report accuracy, recall on the bad class and total cost for the baseline and both approaches.
- (c) State in 4–6 lines which approach you would ship and why.

B2. Group outcomes**[4 marks]***Dataset: A3.*

- (a) Split the test set by a protected-style attribute available in the data (for example `personal_status` or age band).
- (b) Report approval rate, recall and false-positive rate per group in one table.
- (c) In 6–8 lines, state whether the differences you see are evidence of a problem, and what further information you would need to decide.

Section C (Hard) — compulsory (2)

C1. Explain one decision

[2 marks]

Dataset: A3.

- (a) Pick one applicant your model rejected and list the three features that pushed the score most.
- (b) Write the two-sentence explanation you would give that applicant.

End of Lab Assignment 7

Sources: Müller and Guido, *Introduction to Machine Learning with Python*, companion notebooks.;McKinney, *Python for Data Analysis* (3e), O'Reilly, 2022.;Scikit-learn user guide, accessed 2026-08-04.;Matplotlib and Seaborn documentation, accessed 2026-08-04.