

Applied Machine Learning

Unit 02 — Lecture 02 Notes

Loss Functions for Regression, Classification and Margins

Tofik Ali

Autumn 2026

Read this first

These notes are written so that you can learn the material *without* having been in the lecture. Every concept is developed from scratch, every number you see was produced by running the code shown beside it, and every figure was generated from real computation rather than drawn by hand.

Work through them with a Python session open. The whole lecture rests on one question — *what number are we trying to make small?* — and you will understand it far better by computing a few losses by hand than by reading about them.

Contents

1	How to use these notes	3
1.1	What you need	3
1.2	Learning outcomes	3
2	What a loss function is	4
2.1	Three words that are not synonyms	4
2.2	The two properties a loss must have	4
3	Regression losses	5
3.1	Mean squared error	5
3.2	Mean absolute error	5
3.3	Root mean squared error	5
3.4	The result that matters: MSE fits the mean, MAE fits the median	6
3.5	Huber loss	7
4	Classification losses	10
4.1	Why accuracy cannot be the loss	10
4.2	Binary cross-entropy	10
4.3	Confident and wrong	11
4.4	Categorical and sparse categorical cross-entropy	12
5	Margin losses	14
5.1	Hinge loss	14
5.2	Squared hinge, and the family portrait	15
5.3	Triplet loss	15

6	Choosing a loss	17
7	Summary	17
8	Practice questions	19
9	Solutions	20
10	References and further reading	22

1 How to use these notes

Read in order. Section 5 assumes Section 4, which assumes Section 3.

There are four kinds of box:

- **Key idea** — the one sentence to remember a month from now.
- **Worked example** — a calculation done by hand, with small numbers, before any library is used.
- **Caution** — the specific mistake that section exists to prevent.
- **Try it yourself** — something to run, with the expected output given so you can check.

Code appears in a framed block, and directly beneath it a shaded block showing what that code *actually printed* when these notes were compiled. If you run it and get something different, trust your machine and tell me.

1.1 What you need

```
pip install numpy pandas scikit-learn matplotlib seaborn jupyter
```

Every number in these notes was produced with:

```
numpy 2.2.6 | sklearn 1.7.2 | pandas 2.3.3
```

1.2 Learning outcomes

By the end of this lecture you should be able to:

1. State what a loss function is and how it differs from a metric (CO3).
2. Compute MSE, MAE, Huber, cross-entropy and hinge loss by hand on small examples, and reproduce each with NumPy and scikit-learn (CO3).
3. Explain why the choice of loss changes what the fitted model *is*, not merely how well it scores (CO3).
4. Choose an appropriate loss for a stated problem and defend the choice (CO3).

2 What a loss function is

Every supervised model is fitted by the same three-step recipe:

1. Propose a family of candidate models, controlled by parameters.
2. Define a single number saying how badly a given candidate does.
3. Search the parameters for a candidate that makes that number small.

Step 2 is the **loss function**. This lecture is entirely about it. Step 3 is Unit III.

Key idea

A loss function converts *a set of mistakes* into *one number*. That compression is the whole trick, and it is also where all the judgement lives: by choosing how mistakes are converted, you are declaring what kind of mistake you consider serious.

2.1 Three words that are not synonyms

Students routinely mix these up, and the confusion causes real bugs.

Term	Meaning
Loss	Computed on <i>one</i> example. “How bad was this single prediction?”
Cost or objective	The loss averaged (or summed) over the training set, plus any regularization term. This is the quantity actually minimised.
Metric	What you <i>report</i> to a human. Need not be differentiable, need not be the thing you optimised.

In this course, and in most writing, “loss” is used loosely for both the per-example loss and the average. The distinction that matters is the third one.

The loss and the metric are allowed to differ, and often should

You may train a classifier with cross-entropy and report F1. You may train a regressor with Huber and report RMSE. That is not sloppiness — it reflects that the training objective must be smooth and differentiable, while the reported metric must be meaningful to a person.

What is *not* acceptable is failing to notice. If your loss and your metric disagree about which model is better, you need to be able to say which one you trust and why.

2.2 The two properties a loss must have

1. **It must be small when the model is right.** Obvious, but it rules out fewer things than you would think.
2. **It must be usable by the optimiser.** In practice: differentiable almost everywhere, with a gradient that actually points somewhere useful.

The second requirement is why we cannot simply optimise the thing we care about. We return to that in Section 4.

3 Regression losses

The prediction is a number, so the mistake is a number: the **residual**

$$r_i = y_i - \hat{y}_i.$$

Every regression loss is a rule for turning a residual into a non-negative penalty, then averaging.

3.1 Mean squared error

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n r_i^2$$

Squaring does two things: it removes the sign, and it makes large residuals disproportionately expensive. A residual twice as large costs four times as much.

3.2 Mean absolute error

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |r_i|$$

The absolute value also removes the sign, but charges in proportion. A residual twice as large costs twice as much — no more.

3.3 Root mean squared error

$$\text{RMSE} = \sqrt{\text{MSE}}$$

RMSE is not really a separate loss. Minimising RMSE and minimising MSE give the identical model, because the square root is increasing. RMSE exists for *reporting*: it is back in the units of y , so it can be read aloud.

MSE, MAE and RMSE by hand

Four houses. Actual prices y , predicted $\hat{y}$, in lakhs of rupees.

y	3.0	5.0	2.0	8.0
$\hat{y}$	2.5	5.5	4.0	7.0
$r = y - \hat{y}$	0.5	-0.5	-2.0	1.0
r^2	0.25	0.25	4.0	1.0
$ r $	0.5	0.5	2.0	1.0

$$\text{MSE} = \frac{0.25 + 0.25 + 4 + 1}{4} = \frac{5.5}{4} = \mathbf{1.375} \quad \text{MAE} = \frac{0.5 + 0.5 + 2 + 1}{4} = \frac{4}{4} = \mathbf{1.0}$$

$$\text{RMSE} = \sqrt{1.375} = \mathbf{1.1726}$$

Notice that the third house, which is off by 2, contributes 4 of the 5.5 squared-error total — **73% of the loss from 25% of the data**. Under MAE it contributes 2 of 4, exactly half. That single difference is the whole story of this section.

Now the same thing in code.

```
import numpy as np
from sklearn.metrics import mean_squared_error, mean_absolute_error
```

```

y = np.array([3.0, 5.0, 2.0, 8.0])
yhat = np.array([2.5, 5.5, 4.0, 7.0])
r = y - yhat

print("residuals ", r)
print("MSE np ", np.mean(r**2))
print("MAE np ", np.mean(np.abs(r)))
print("RMSE np ", np.sqrt(np.mean(r**2)))
print("MSE sklearn ", mean_squared_error(y, yhat))
print("MAE sklearn ", mean_absolute_error(y, yhat))

```

```

residuals [ 0.5 -0.5 -2.  1. ]
MSE np 1.375
MAE np 1.0
RMSE np 1.1726039399558574
MSE sklearn 1.375
MAE sklearn 1.0

```

The library agrees with the arithmetic, as it must. There is nothing inside `mean_squared_error` that you did not just do on paper.




Figure 1: What each loss charges for a single residual. The squared error is a bowl; the absolute error is a V with a kink at zero; Huber is the bowl near zero and the V far away.

3.4 The result that matters: MSE fits the mean, MAE fits the median

Here is the cleanest way to see the difference. Forget features entirely and ask the simplest question in statistics: *if I must predict one constant c for every observation, what constant should I choose?*

The answer depends entirely on the loss:

- Minimising $\sum (y_i - c)^2$ gives $c = \bar{y}$, the **mean**.
- Minimising $\sum |y_i - c|$ gives $c = \text{median}(y)$.

Both are worth proving to yourself; the first is one line of calculus, the second a counting argument (Question 2).

Now the consequence. Take six observations, and drag one of them far away.

```
clean = np.array([10., 11., 12., 13., 14., 14.])
outlier = np.array([10., 11., 12., 13., 14., 94.])

for name, d in [("clean", clean), ("with outlier", outlier)]:
    print("%-13s mean=%.4f median=%.4f" % (name, d.mean(), np.median(d)))

clean mean=12.3333 median=12.5000
with outlier mean=25.6667 median=12.5000
```

One corrupted reading moved the MSE-optimal prediction from 12.33 to **25.67** — past every single clean observation. The MAE-optimal prediction did not move at all.

We can confirm that these really are the minimisers rather than taking it on faith, by evaluating both losses on a fine grid of candidate constants:

```
grid = np.linspace(0, 100, 200001)
mse_curve = ((outlier[:, None] - grid[None, :])**2).mean(axis=0)
mae_curve = (np.abs(outlier[:, None] - grid[None, :])).mean(axis=0)
print("argmin MSE ->", round(grid[mse_curve.argmin()], 4), " mean =", outlier.mean())
print("argmin MAE ->", round(grid[mae_curve.argmin()], 4), " median =", np.median(
    outlier))

argmin MSE -> 25.6665 mean = 25.666666666666668
argmin MAE -> 12.0 median = 12.5
```

Why the MAE grid search returned 12.0 and not 12.5

This is not a bug, and it is worth understanding. With an *even* number of observations the absolute-error loss is not minimised at a single point — it is flat across the whole interval between the two central observations, here [12, 13]. Every constant in that interval gives exactly the same MAE.

`argmin` returns the *first* index achieving the minimum, so it reports the left end, 12.0. NumPy's `median` conventionally returns the midpoint, 12.5. Both are correct minimisers; neither is more right than the other.

The flat region is also why MAE is harder to optimise than MSE: a flat loss gives a zero gradient, and an optimiser standing on a plateau does not know which way to go.

The sentence to remember

Choosing a loss is not choosing how to *score* your model. It is choosing *what your model will become*. MSE builds a model of the average case; MAE builds a model of the typical case. On clean symmetric data they nearly agree. On data with a heavy tail they answer different questions.

3.5 Huber loss

We would like MSE's smooth, well-behaved gradient near zero and MAE's indifference to outliers far from zero. Huber loss simply uses each where it is wanted:

$$L_{\delta}(r) = \begin{cases} \frac{1}{2}r^2, & |r| \leq \delta, \\ \delta \left(|r| - \frac{1}{2}\delta \right), & |r| > \delta. \end{cases}$$




Figure 2: The best constant prediction as a single outlier is dragged further and further away. The MSE solution follows it without limit; the MAE solution is unmoved.

The second branch is exactly the tangent line to the first at $r = \delta$, so the two pieces meet with matching value *and* matching slope. The result is differentiable everywhere, including at the join.

Huber by hand, with delta equal to one

Same four residuals as before: 0.5, -0.5, -2.0, 1.0.

- $|0.5| \leq 1$, so quadratic: $\frac{1}{2}(0.5)^2 = 0.125$.
- $|-0.5| \leq 1$, so quadratic: 0.125.
- $|-2.0| > 1$, so linear: $1 \cdot (2 - 0.5) = 1.5$.
- $|1.0| \leq 1$, so quadratic: $\frac{1}{2}(1)^2 = 0.5$.

Mean = $(0.125 + 0.125 + 1.5 + 0.5)/4 = 2.25/4 = \mathbf{0.5625}$.

Compare with the squared error, where that third point alone contributed 4.0. Huber charged it 1.5.

```
def huber(r, d=1.0):
    a = np.abs(r)
    return np.where(a <= d, 0.5*r**2, d*(a - 0.5*d))

print("per point ", huber(r, 1.0))
print("mean ", huber(r, 1.0).mean())
```

```
per point [0.125 0.125 1.5 0.5 ]
mean 0.5625
```

δ is a hyperparameter with units

δ is measured in the units of y . A δ of 1 means something completely different for house prices in lakhs than for temperatures in degrees. Choose it by cross-validation, or set it from a robust scale estimate of the residuals — never copy it from a tutorial.




Figure 3: δ controls where the loss stops being quadratic. Small δ behaves almost like MAE; large δ is almost pure MSE over the range shown.

Try it yourself

Re-run the constant-prediction experiment from Section 3.4 using Huber loss with $\delta = 1$ and again with $\delta = 50$, and find the minimising constant on a grid. You should find the $\delta = 1$ answer sits near the median and the $\delta = 50$ answer near the mean, because when δ is larger than every residual Huber *is* scaled MSE.

4 Classification losses

Now the prediction is a class. The obvious loss is: charge 1 for a wrong answer and 0 for a right one. This is the **0–1 loss**, and it is exactly what accuracy measures.

It is also almost useless for training.

4.1 Why accuracy cannot be the loss

Consider a one-parameter classifier that predicts class 1 when $w x > 0$. Vary w and watch both accuracy and cross-entropy:

```
X = np.array([[1.0], [2.0], [-1.0], [-2.0]])
y = np.array([1, 1, 0, 0])

for w in [-2.0, -1.0, 0.0, 1.0, 2.0]:
    z = X.ravel() * w
    acc = ((z > 0).astype(int) == y).mean()
    p = 1 / (1 + np.exp(-z))
    ll = -np.mean(y*np.log(p+1e-12) + (1-y)*np.log(1-p+1e-12))
    print("w=%5.1f accuracy=%.2f log-loss=%.4f" % (w, acc, ll))
```

```
w= -2.0 accuracy=0.00 log-loss=3.0725
w= -1.0 accuracy=0.00 log-loss=1.7201
w=  0.0 accuracy=0.50 log-loss=0.6931
w=  1.0 accuracy=1.00 log-loss=0.2201
w=  2.0 accuracy=1.00 log-loss=0.0725
```

Look at the two columns.

Accuracy takes three values in total: 0, 0.5, 1. It is a step function — flat almost everywhere, with sudden jumps. Its derivative is zero wherever it is defined, so gradient descent standing anywhere on it receives the message “every direction is equally good”. And accuracy cannot distinguish $w = 1$ from $w = 2$ at all, though the second is clearly a more confident fit.

Cross-entropy decreases smoothly and strictly the whole way. Every value of w gets a gradient telling it which way to move and by how much.

Surrogate losses

We optimise a **surrogate**: a smooth loss that is a good stand-in for the thing we actually care about. Cross-entropy and hinge are the two standard surrogates for the 0–1 loss. This is why the number you optimise and the number you report are so often different — not carelessness, but necessity.

4.2 Binary cross-entropy

The model outputs p_i , its claimed probability that example i belongs to class 1. With true label $y_i \in \{0, 1\}$:

$$\text{BCE} = -\frac{1}{n} \sum_{i=1}^n \left[y_i \log p_i + (1 - y_i) \log(1 - p_i) \right]$$

The bracket looks like two terms but only ever contributes one: when $y_i = 1$ the second term vanishes, and when $y_i = 0$ the first does. So the formula says something very simple:

Key idea

Take the probability the model assigned to *the class that was actually correct*. Take its negative logarithm. Average over the data. That is all cross-entropy is.

Binary cross-entropy by hand

Three emails. $y = 1$ means spam. The model reports p , its probability of spam.

true y	1	0	1
model p (spam)	0.9	0.2	0.4
prob. of correct class	0.9	0.8	0.4
$-\log(\cdot)$	0.1054	0.2231	0.9163

Row three is the only step with any content: for the first and third emails the correct class was spam, so we take p ; for the second the correct class was not-spam, so we take $1 - p = 0.8$.

$$\text{BCE} = \frac{0.1054 + 0.2231 + 0.9163}{3} = \mathbf{0.4149}$$

The third email dominates. The model was not confidently wrong there — it said 0.4 for something that was spam, a near coin-flip — and it still cost more than the other two combined.

```
from sklearn.metrics import log_loss

y = np.array([1, 0, 1])
p = np.array([0.9, 0.2, 0.4])
pc = np.where(y == 1, p, 1 - p) # prob assigned to the correct class

print("prob of correct", pc)
print("-log each ", -np.log(pc))
print("BCE by hand ", (-np.log(pc)).mean())
print("BCE sklearn ", log_loss(y, p))
```

```
prob of correct [0.9 0.8 0.4]
-log each [0.1054 0.2231 0.9163]
BCE by hand 0.414931599615397
BCE sklearn 0.414931599615397
```

4.3 Confident and wrong

The logarithm is what gives cross-entropy its character:

p on correct class	loss	
0.99	0.0101	confident and right — nearly free
0.90	0.1054	
0.50	0.6931	a coin flip costs $\log 2$
0.10	2.3026	
0.01	4.6052	confident and wrong — $456\times$ the cost of being right
0.001	6.9078	and it keeps going

As $p \rightarrow 0$ the loss grows without bound. A model that is certain and wrong on even one training example pays an enormous price, which is precisely the behaviour we want: it forces models to be honest about their uncertainty rather than confidently guessing.




Figure 4: Cross-entropy against the probability given to the correct class. Being right is cheap; being confidently wrong is unbundably expensive.

log 0 is not a number

A model that outputs exactly 0 or exactly 1 produces an infinite loss and your training run dies. Every library therefore clips probabilities into $[\varepsilon, 1 - \varepsilon]$. This is why you should feed `log_loss` probabilities rather than hard 0/1 predictions — and why you should never write your own cross-entropy without an ε , as in the `+1e-12` above.

4.4 Categorical and sparse categorical cross-entropy

With K classes the model outputs a probability vector $\mathbf{p}_i$ summing to one. The loss is unchanged in spirit:

$$\text{CCE} = -\frac{1}{n} \sum_{i=1}^n \sum_{k=1}^K y_{ik} \log p_{ik}$$

where y_{ik} is 1 if example i is of class k and 0 otherwise. Because only one y_{ik} per row is non-zero, the inner sum again picks out a single term: the log-probability of the correct class.

Categorical and **sparse categorical** cross-entropy are *the same loss*. They differ only in how you hand over the labels:

	Label format	Example for 3 classes
Categorical	one-hot rows	<code>[[1,0,0], [0,1,0], [0,0,1]]</code>
Sparse categorical	integer class indices	<code>[0, 1, 2]</code>

```
P = np.array([[0.7, 0.2, 0.1],
              [0.1, 0.8, 0.1],
              [0.2, 0.2, 0.6]])
onehot = np.array([[1,0,0], [0,1,0], [0,0,1]])
sparse = np.array([0, 1, 2])
```

```
print("categorical", -np.sum(onehot * np.log(P), axis=1).mean())
print("sparse ", -np.log(P[np.arange(3), sparse]).mean())
print("sklearn ", log_loss(sparse, P))
```

```
categorical 0.3635480396729776
sparse 0.3635480396729776
sklearn 0.3635480396729776
```

Identical to sixteen decimal places, because it is the same arithmetic.

Why the sparse form exists

It is purely a memory and speed optimisation. With $K = 10,000$ classes — a language model's vocabulary, say — the one-hot matrix for a batch of 256 examples holds 2.56 million numbers of which 256 are non-zero. The sparse form stores 256 integers and indexes directly. Same loss, same gradients, a fraction of the memory.

Choosing the wrong one is one of the most common beginner errors in Keras and PyTorch, and the error message is rarely helpful. Look at the shape of your labels: 2-D means categorical, 1-D means sparse.

5 Margin losses

Cross-entropy needs a probability. Some classifiers do not produce one — a support vector machine outputs a raw score $f(x)$, positive for one class and negative for the other. For these we use **margin** losses.

Relabel the classes as $y \in \{-1, +1\}$ and define the **margin**

$$m_i = y_i f(x_i).$$

The margin is positive exactly when the prediction is correct, and its magnitude says how decisively. This single number is all a margin loss looks at.

5.1 Hinge loss

$$L_{\text{hinge}}(m) = \max(0, 1 - m)$$

Read it directly: if the margin is at least 1, the loss is zero. Otherwise you pay the shortfall.

Key idea

Hinge loss stops caring once you are right by a comfortable margin. Cross-entropy never entirely stops caring — it will always take a little more confidence if it can get it. That difference is why SVMs are determined by a handful of support vectors while logistic regression is influenced by every point.

Hinge loss by hand

Three points, true labels in $\{-1, +1\}$, raw scores f .

true y	+1	-1	+1
score $f(x)$	2.0	-0.5	0.3
margin $m = yf$	2.0	0.5	0.3
$\max(0, 1 - m)$	0.0	0.5	0.7
$L = \frac{0.0 + 0.5 + 0.7}{3} = \frac{1.2}{3} = \mathbf{0.4}$			

Every one of these three points is classified *correctly* — all three margins are positive. Yet two of them still incur a loss, because they sit inside the margin band. This is the behaviour that produces a wide separating boundary rather than one that merely separates.

```
from sklearn.metrics import hinge_loss

ys = np.array([1, -1, 1])
f = np.array([2.0, -0.5, 0.3])
m = ys * f

print("margins ", m)
print("hinge each ", np.maximum(0, 1 - m))
print("hinge by hand", np.maximum(0, 1 - m).mean())
print("hinge sklearn", hinge_loss(ys, f))
```

```
margins [2. 0.5 0.3]
hinge each [0. 0.5 0.7]
hinge by hand 0.39999999999999997
hinge sklearn 0.39999999999999997
```

Caution

That is 0.4, printed as 0.39999999999999997. Binary floating point cannot represent 0.4 exactly. Never test loss values with `==`; use `np.isclose` or `math.isclose`.

5.2 Squared hinge, and the family portrait

Squaring the hinge, $\max(0, 1 - m)^2$, gives a smooth loss (no kink at $m = 1$) that punishes large violations harder. It is more sensitive to outliers than plain hinge, for exactly the reason MSE is more sensitive than MAE.




Figure 5: The 0–1 loss we would like to minimise, and the three smooth surrogates that we minimise instead. All three are upper bounds on the 0–1 loss, which is what makes them legitimate stand-ins.

Read Figure 5 carefully; it is the summary of this whole section.

- The grey step is what we care about: wrong ($m \leq 0$) costs 1, right costs 0.
- Each smooth curve sits *above* the step everywhere. Driving the surrogate down therefore drives the 0–1 loss down too.
- Hinge reaches exactly zero at $m = 1$ and stays there. Logistic approaches zero but never arrives.
- For very negative margins, hinge grows linearly and squared hinge quadratically — so one badly mislabelled point disturbs a squared-hinge model far more.

5.3 Triplet loss

The last loss in this lecture answers a different question. Sometimes we do not want to classify at all — we want to *learn a representation* in which similar things are close together. Face recognition is the standard example: you cannot train a classifier with one output per person when new people keep arriving.

Triplet loss works on three examples at a time:

- an **anchor** a ,
- a **positive** p — same identity as the anchor,
- a **negative** n — a different identity.

$$L = \max(0, d(a, p) - d(a, n) + \alpha)$$

In words: the anchor should be closer to the positive than to the negative, by at least a margin α . If it already is, the loss is zero and nothing is learned from this triplet.




Figure 6: A triplet that still has something to teach. The anchor is closer to the positive than to the negative, but not by the required margin, so the loss is non-zero and the embedding will be adjusted.

For the configuration in Figure 6, with $\alpha = 1$:

$$L = \max(0, 0.97 - 1.54 + 1.0) = \mathbf{0.42}$$

The ordering is already correct — the positive is nearer — but only by 0.58, short of the margin of 1, so a penalty remains.

Most triplets teach you nothing

Pick three examples at random and the constraint is usually already satisfied, giving zero loss and zero gradient. Practical systems therefore spend most of their engineering effort on *triplet mining*: actively searching for the hard cases where $d(a, p)$ is nearly as large as $d(a, n)$. Without mining, training stalls almost immediately.

6 Choosing a loss

Situation	Use	Because
Regression, clean data, errors roughly symmetric	MSE	Smooth, fast, and the mean is the sensible target
Regression, outliers you do not trust	MAE	Fits the typical case; one bad reading cannot drag it
Regression, outliers you cannot exclude but must not dominate	Huber	Quadratic where it matters, linear where it would distort
Two classes, want probabilities	Binary cross-entropy	Calibrated outputs; punishes confident errors
K classes, labels one-hot	Categorical cross-entropy	Same loss, one-hot API
K classes, labels are integers	Sparse categorical cross-entropy	Same loss, far less memory
Two classes, want a wide boundary, no probabilities needed	Hinge	Ignores points already safely correct
Learning an embedding where "similar" must mean "close"	Triplet	Trains on relative distance, handles unseen identities

Do not choose by asking which gives the smaller number

Losses are on different scales and are not comparable to each other. An MAE of 1.0 is not "better" than an MSE of 1.375 — they are different quantities in different units, computed from the identical predictions, as the worked example in Section 3 shows.

Compare *models* using a single fixed metric on held-out data. Use the loss to train; use the metric to decide.

7 Summary

1. A loss compresses a set of mistakes into one number, and that choice declares what kind of mistake you consider serious.
2. Loss, cost and metric are three different things. Loss and metric are allowed to differ, and often must.
3. MSE fits the mean; MAE fits the median. One outlier moved our MSE-optimal constant from 12.33 to 25.67 and left the MAE-optimal constant untouched.
4. Huber is MSE near zero and MAE far away, joined smoothly at δ . δ carries the units of y .
5. Accuracy cannot be a training loss: it is a step function with zero gradient. We minimise smooth surrogates instead.
6. Cross-entropy is the negative log of the probability given to the correct class. Confident and wrong is unboundedly expensive.
7. Categorical and sparse categorical cross-entropy are the same loss with different label formats.

8. Hinge loss charges nothing once the margin exceeds 1, which is what produces a wide margin and a small set of support vectors.
9. Triplet loss trains representations by relative distance, and needs hard-triplet mining to make progress.

8 Practice questions

Attempt these before reading Section 9.

1. Predictions $\hat{y} = [4, 7, 10]$ for true values $y = [5, 5, 5]$. Compute MSE, MAE and RMSE by hand.
2. Prove that the constant c minimising $\sum_i |y_i - c|$ is a median of the y_i . *Hint:* think about how the total changes as c moves slightly right.
3. A model predicts probability 0.7 for the correct class on every one of 50 examples. What is the cross-entropy? Now one example changes to 0.001 on the correct class. What is the new value?
4. You are predicting delivery times. Being 10 minutes late is more than twice as bad as being 5 minutes late for your users, and occasional sensor glitches produce impossible readings. Which loss, and why?
5. For $y = +1$ and score $f(x) = 3$, give the hinge loss and the logistic loss. Which is exactly zero, and what does that imply about whether this point affects training?
6. Huber with $\delta = 100$ is applied to residuals that never exceed 10. What loss have you effectively chosen? What does this tell you about setting δ without looking at your data?
7. A classmate reports "my model achieved a loss of 0.31, which beats the published baseline of 0.44." Give two distinct reasons this comparison may be meaningless.
8. Your labels are the integers $[2, 0, 1]$ and your model outputs a 3×3 matrix of probabilities. Which cross-entropy variant do you use, and what breaks if you choose the other one?
9. For a triplet with $d(a, p) = 0.4$, $d(a, n) = 1.9$ and $\alpha = 1$, compute the loss. What does your answer imply about how much this triplet contributes to training?
10. MSE is minimised by the mean and MAE by the median. What summary statistic is minimised by the 0–1 loss, if we apply it to a discrete prediction problem? Justify your answer.

9 Solutions

1. Residuals $r = y - \hat{y} = [1, -2, -5]$. $\text{MSE} = (1+4+25)/3 = 30/3 = 10$. $\text{MAE} = (1+2+5)/3 = 8/3 = 2.667$. $\text{RMSE} = \sqrt{10} = 3.162$. Note $\text{RMSE} > \text{MAE}$, which is always true and is a direct consequence of squaring inflating the largest residual.

2. Order the observations. Suppose c sits so that k observations lie strictly below it and $n - k$ strictly above. Move c right by a small ϵ (small enough to cross no observation). Each of the k points below becomes ϵ further away; each of the $n - k$ above becomes ϵ nearer. The total changes by $\epsilon(k - (n - k)) = \epsilon(2k - n)$.

This is negative while $k < n/2$, so moving right helps; positive once $k > n/2$, so moving right hurts. The minimum is therefore where $k = n/2$ — a median. When n is even any point between the two central observations gives $2k - n = 0$, so the whole interval is optimal, which is exactly the plateau we observed numerically in Section 3.4.

3. All fifty at 0.7: the loss is just $-\log 0.7 = 0.3567$ (the mean of identical values). Changing one to 0.001:

$$\frac{49 \times 0.3567 + 6.9078}{50} = \frac{17.478 + 6.908}{50} = 0.4877.$$

A single confidently wrong example raised the loss by 37%. That sensitivity is the design intent, but it also means one mislabelled training row can noticeably distort your objective — a reason to inspect the highest-loss examples after training.

4. The two requirements conflict, which is the point of the question. "More than twice as bad" argues for a super-linear penalty, i.e. MSE. "Impossible sensor readings" argues against MSE, which those readings would dominate. **Huber** resolves it: quadratic over the range of genuine lateness, linear beyond, with δ set at the boundary of physically plausible delays. Cleaning the impossible readings first is better still — choosing a robust loss is not a substitute for fixing your data.

5. Margin $m = yf(x) = 3$. Hinge = $\max(0, 1 - 3) = 0$. Logistic = $\log(1 + e^{-3}) = 0.0486$. The hinge loss is exactly zero, so this point contributes no gradient and could be deleted from the training set without changing the SVM at all — it is not a support vector. Under logistic loss it still contributes a small pull, so it is never entirely irrelevant.

6. With $\delta = 100$ and $|r| \leq 10$ always, the condition $|r| \leq \delta$ holds for every point, so only the quadratic branch is ever used. You have chosen $\frac{1}{2} \times \text{MSE}$, and gained none of Huber's robustness while adding a hyperparameter. The lesson: δ must sit inside the range your residuals actually occupy, so it has to be set after looking at the residual distribution.

7. Any two of: (i) the two numbers may be different losses entirely — 0.31 MAE against 0.44 MSE compares nothing; (ii) they may be computed on different datasets, or different train/test splits; (iii) one may be a training loss and the other a test loss; (iv) the target variable may be scaled differently, and every regression loss inherits the units of y ; (v) with a different class balance, cross-entropy values are not comparable.

8. Labels are 1-D integers, so use **sparse categorical cross-entropy**. If you pass integer labels to the categorical form, the library will either raise a shape error (expecting $n \times K$, receiving n) or, in the worst case, silently broadcast and compute a meaningless number. This is why you check label shape before choosing, rather than after debugging.

9. $L = \max(0, 0.4 - 1.9 + 1) = \max(0, -0.5) = \mathbf{0}$. The constraint is already satisfied with room to spare, so the gradient is zero and this triplet contributes *nothing* to training. It is exactly the kind of easy triplet that mining exists to avoid.

10. The **mode**. The 0–1 loss charges 1 for every prediction that is not exactly right, so the expected loss of predicting class c is $1 - P(Y = c)$, minimised by choosing the most probable class. This completes a pleasing pattern: squared error gives the mean, absolute error the median, 0–1 error the mode — the three classical measures of central tendency each fall out of a different loss.

10 References and further reading

All links were checked on 5 August 2026. Free resources are marked [free].

Textbooks

- Bishop, *Pattern Recognition and Machine Learning*, Springer 2006. §1.5 (decision theory and loss), §4.3 (logistic regression and cross-entropy), §7.1 (maximum margin and hinge). A free PDF is hosted by Microsoft Research; their site blocks automated link checks, so no direct URL is printed here — search for the title plus “Microsoft Research”.
- Müller and Guido, *Introduction to Machine Learning with Python*, O’Reilly 2016. Chapter 2 for where these losses sit inside the estimators you will actually call.
- Shalev-Shwartz and Ben-David, *Understanding Machine Learning: From Theory to Algorithms*, Cambridge 2014. [free] Chapter 15 on surrogate losses and why upper-bounding the 0–1 loss is the right move. <https://www.cs.huji.ac.il/~shais/UnderstandingMachineLearning/>
- Murphy, *Probabilistic Machine Learning: An Introduction*, MIT Press 2022. [free] Chapter 5 for the decision-theoretic view. <https://probml.github.io/pml-book/book1.html>

Online courses

- **NPTEL, Introduction to Machine Learning** — Prof. Balaraman Ravindran, IIT Madras. The closest match to this syllabus. [free] <https://nptel.ac.in/courses/106106139>
- **SWAYAM** — national portal hosting NPTEL and other Indian university courses. [free] <https://swayam.gov.in/>
- **Google Machine Learning Crash Course** — the "Descending into ML" and "Logistic Regression" modules cover squared loss and log loss with interactive exercises. [free] <https://developers.google.com/machine-learning/crash-course>

Videos

- StatQuest, *Entropy (for data science) Clearly Explained*. Builds the logarithm in cross-entropy from first principles. [free] <https://youtu.be/YtebGVx-Fxw>
- StatQuest, *Support Vector Machines Part 1 (of 3): Main Ideas*. Where the margin comes from and why points beyond it stop mattering. [free] <https://youtu.be/efR1C6CvhmE>
- StatQuest video index — short, clear explanations for most topics in this syllabus. [free] <https://statquest.org/video-index/>

Documentation

- scikit-learn, *Metrics and scoring*. The catalogue of losses and metrics, and the scorer interface. [free] https://scikit-learn.org/stable/modules/model_evaluation.html
- scikit-learn, *Common pitfalls and recommended practices*. [free] https://scikit-learn.org/stable/common_pitfalls.html
- *Dive into Deep Learning*, §3.4 (softmax regression and cross-entropy). Runnable, with the derivations laid out. [free] https://d2l.ai/chapter_linear-classification/softmax-regression.html

Articles

- Huber, “Robust Estimation of a Location Parameter”, *Annals of Mathematical Statistics* 35(1), 1964. The original paper; the first two pages are readable and set up exactly the mean-versus-median tension in Section 3.4.
- Schroff, Kalenichenko and Philbin, “FaceNet: A Unified Embedding for Face Recognition and Clustering”, CVPR 2015. Where triplet loss and triplet mining were popularised. [free] <https://arxiv.org/abs/1503.03832>

Sources: Bishop, *Pattern Recognition and Machine Learning*, §1.5, §4.3, §7.1.; Zhang, Lipton, Li and Smola, *Dive into Deep Learning*, §3.4.; Shalev-Shwartz and Ben-David, *Understanding Machine Learning*, Ch. 15.; Murphy, *Probabilistic Machine Learning: An Introduction*, Ch. 5.