



**"RANKED 45
IN INDIA**

#University Category



**WORLD
UNIVERSITY
RANKINGS**

NO.1 PVT. UNIVERSITY IN
ACADEMIC REPUTATION IN INDIA



ACCREDITED GRADE 'A'
BY NAAC



PERFECT SCORE OF **150/150** AS A TESTAMENT
TO EXCEPTIONAL E-LEARNING METHODS

1



Applied Machine Learning

Unit 06 – Lecture 16: Ensembles – Bagging, Boosting and Random Forests

Tofik Ali

School of Computer Science, UPES Dehradun

Autumn 2026

Why a committee of mediocre models beats a good one

Topic focus: **diversity is the resource**

Repository: <https://github.com/tali7c/Applied-Machine-Learning>

Sources: Han, Kamber and Pei, *Data Mining: Concepts and Techniques* (3e), Ch. 8–9.; James et al., *An Introduction to Statistical Learning with Python*, Ch. 8–9.; Bishop, *Pattern Recognition and Machine Learning*, Ch. 4–5, 7, 14.; Zhang et al., *Dive into Deep Learning*, Ch. 5.



Quick Links

[Why ensembles work](#)

[Bagging](#)

[Variance, not bias](#)

[Random forests](#)

[Feature importance](#)

[Boosting](#)

[Bagging vs boosting](#)

[Summary](#)

Agenda

Why an ensemble beats its members

Bagging and the bootstrap

Bagging reduces variance, not bias

Random forests

Feature importance, honestly

Boosting: AdaBoost and gradient boosting

Bagging against boosting

Summary

How this lecture works

Every idea appears twice

1. **By hand** — twenty-five voters, five rows in a hat, ten points on a line. Small enough that every number can be checked with a pen.
2. **In code** — the same arithmetic, then the same arithmetic again inside `scikit-learn`, so you can see the library is doing exactly this and nothing magic.

Commit before you look

Five slides ask you to predict a number before it is revealed. Write your guess down. Being wrong on paper is how the idea sticks.

The centrepiece is **three rounds of AdaBoost worked on ten points**: the weighted error, the α , the reweighting, and the combined sign. That arithmetic is the most examinable thing in the lecture.

Where we are

The decision-tree lecture left a problem

An unrestricted tree hits training accuracy 1.0000 every time and then falls over on new rows. Pruning bought some of that back, but a single tree stayed **unstable**: move a few rows and the root split changes.

This lecture turns that into an advantage

Instability means the trees you grow on

slig

Tofik Ali

Applied Machine Learning

Average many different unstable models

Today's one question

I have a mediocre model. *Can I build a good one out of many copies of it?*

And the answer has a condition

Yes — provided the copies make their mistakes on **different rows**. Everything in this lecture is machinery for forcing them to disagree: resample the rows (**bagging**), restrict the features (**random forests**), or reweight the rows towards what is still wrong (**boosting**).



UNIVERSITY OF TOMORROW

School of Computer Science, UPES Dehradun

One seed — and the two things that are not it

The one seed

Every arbitrary random choice in this deck uses **seed** 0: `np.random.default_rng(0)` or `random_state=0`. Splits, bootstrap draws, flipped labels. Every listing that shuffles or samples shows it.

A seed that defines data

The three planted noise columns use `default_rng(42)`. Change it and you have *different data*, not a different experiment — so it stays.

A sweep

Twenty splits for out-of-bag, ten per noise level, with `random_state=s` across the loop. Here the run-to-run variation *is* the demonstration — but the sweep starts at 0, so it reproduces exactly too.

Why this matters to you

Run the code as written and you get the numbers on the slide. That is the only reason to believe a number on a slide.

An ensemble is a committee with a voting rule

Train m models. Ask all of them. Take the majority vote.

The vocabulary

The individual models are the **base learners** or **members**. The rule that combines them is the **aggregation**: majority vote for classification, mean for regression, weighted versions of both.

The claim to be justified

Twenty-five models that are each right 60% of the time can be combined into one that is right about 85% of the time — the difference between a model you would throw away and one you would deploy.

The claim is a piece of eighteenth-century probability, not machine learning: **Condorcet's jury theorem**, 1785.

Predict first #1 — twenty-five mediocre voters

Commit to a number

Twenty-five classifiers. Each one is right 60% of the time. Each one is **independent** of the others — they get their answers wrong on unrelated rows. You take the majority vote.

What accuracy does the majority vote have?

Predict first #1 — twenty-five mediocre voters

Commit to a number

Twenty-five classifiers. Each one is right 60% of the time. Each one is **independent** of the others — they get their answers wrong on unrelated rows. You take the majority vote.

What accuracy does the majority vote have?

$$\Pr[\text{at least 13 of 25 correct}] = \mathbf{0.846232}$$

A gain of +**0.2462** over any single member — and every member is still only 60% accurate. Nothing was improved. Only combined.

By hand: the binomial sum

Let K be the number of members that get a given row right. Members are independent, so

$$\Pr[K = k] = \binom{25}{k} (0.6)^k (0.4)^{25-k}, \quad \Pr[\text{vote correct}] = \sum_{k=13}^{25} \Pr[K = k].$$

The first few terms

k	$\binom{25}{k}$	term
13	5 200 300	0.11395006
14	4 457 400	0.14650722
15	3 268 760	0.16115794
16	2 042 975	0.15108557

The sum

$\sum_{k=13}^{25} = \mathbf{0.846232}$. The largest term is at $k = 15 = 0.6 \times 25$: the mode sits on the right of the majority threshold, so most of the mass lies above it.



The condition that does all the work

The same thing in code

```
from scipy.special import comb
from scipy import stats
n, p = 25, 0.60
maj = sum(comb(n, k, exact=True) * p**k * (1-p)**(n-k) for k in range(13,
    26))
print("majority (k >= 13) =", round(maj, 6))
print("scipy check =", round(stats.binom.sf(12, n, p), 6))
```

```
majority (k >= 13) = 0.846232
scipy check        = 0.846232
single member      = 0.600000
the ensemble gains +0.2462
```

`stats.binom.sf(12, 25, 0.6)` is the survival function — $\Pr[K > 12]$ — and it agrees to every digit. The sum is the definition; the survival function is the same sum with a better name.

The vote amplifies whatever the member already is

p	1 voter	11 voters	25 voters	101 voters
0.40	0.4000	0.2465	0.1538	0.0209
0.45	0.4500	0.3669	0.3063	0.1562
0.49	0.4900	0.4729	0.4598	0.4201
0.50	0.5000	0.5000	0.5000	0.5000
0.51	0.5100	0.5271	0.5402	0.5799
0.55	0.5500	0.6331	0.6937	0.8438
0.60	0.6000	0.7535	0.8462	0.9791
0.70	0.7000	0.9218	0.9825	1.0000

Above 0.5: runs to 1



Even $p = 0.51$ climbs. It just climbs slowly — 101 voters buy you 0.5799.

Below 0.5: runs to 0

$p = 0.40$ gets worse: $0.4000 \rightarrow 0.0209$.
An ensemble of members worse than chance is confidently, reliably wrong.

$p = 0.50$ is a fixed point. This is why the word is *weak* learner, not *bad* learner: better than chance, however slightly.

Drawn: the jury theorem, and what breaks it



Left: independent voters, the curves of the table above. Right: the same twenty-five voters at $p = 0.60$, but now correlated. **The left-hand panel is the advertisement. The right-hand panel is the contract.**

Predict first #2 — what correlation costs

Commit to a number

Same twenty-five voters, each still exactly 60% accurate. But now their *correctness* is correlated: the mean pairwise correlation between “member i was right” and “member j was right” is **0.13** — a correlation most people would call small.

The independent answer was 0.846. What is it now?

Predict first #2 — what correlation costs

Commit to a number

Same twenty-five voters, each still exactly 60% accurate. But now their *correctness* is correlated: the mean pairwise correlation between “member i was right” and “member j was right” is **0.13** — a correlation most people would call small.

The independent answer was 0.846. What is it now?

0.6932 — more than half of the gain, gone.

A correlation of 0.13 throws away 0.153 of the 0.246 the ensemble had won.

Independence is the whole game

Give each voter a latent score $Z_i = \sqrt{\rho} W + \sqrt{1 - \rho} E_i$ and call the voter correct when $Z_i < \Phi^{-1}(0.6)$. The shared W is the common cause; ρ dials how much the voters agree beyond chance. Every voter is still right exactly 60% of the time.

rho_latent	mean pairwise corr	ensemble accuracy	members
0.00	+0.0001	0.8458	0.6002
0.05	+0.0310	0.7798	0.6006
0.10	+0.0627	0.7381	0.5993
0.20	+0.1260	0.6932	0.5995
0.40	+0.2595	0.6476	0.5997
0.70	+0.4883	0.6179	0.6002
1.00	+1.0000	0.5993	0.5993

The last column never moves. **Correlation destroys the ensemble without touching a single member's accuracy.** At $\rho = 1$ all twenty-five voters are the same voter and the committee is worth exactly one member.

Ensembles of classifiers: three different learners, one vo

Diversity does not have to come from resampling. It can come from using *different kinds of model*.

```
members = [("knn", make_pipeline(StandardScaler(), KNeighborsClassifier(5))),  
            ("tree", DecisionTreeClassifier(max_depth=3, random_state=0)),  
            ("lr", make_pipeline(StandardScaler(),  
                                LogisticRegression(max_iter=5000)))]  
vote = VotingClassifier(members, voting="hard").fit(Xtr, ytr)
```

```
knn    test accuracy 0.9474  
tree   test accuracy 0.9006  
lr     test accuracy 0.9591  
vote   test accuracy 0.9474  
mean pairwise disagreement between members 0.0702  
rows where the members were not unanimous 18 of 171
```

Report the result you got. The vote did *not* beat the best member. Three members is far too few, one of them is much worse than the others, and on 153 of 171 rows all three already agree — so the vote can only act on 18 rows. Voting is not free magic.

Bootstrap aggregating, in one paragraph

Draw n rows **with replacement** from a training set of n rows. Fit a model. Repeat B times. Vote.

Why with replacement

Without replacement, drawing n from n gives you the training set back. With replacement, some rows appear twice or three times and some do not appear at all — and *that* is what makes the bags differ.

Who it works for

Bagging helps **unstable** learners — ones whose output changes a lot when the data changes a little. Deep decision trees are the canonical case. Bagging a linear regression is close to a waste of electricity.

Breiman, 1996. The name is a contraction of **bootstrap aggregating**.

By hand: five rows, four bags

Training set $\{A, B, C, D, E\}$. Each bag is five draws with replacement.

bag 1: E D C B B	in-bag unique 4/5	out-of-bag A
bag 2: A A A A E	in-bag unique 2/5	out-of-bag B, C, D
bag 3: D E C D E	in-bag unique 3/5	out-of-bag A, B
bag 4: D D C C E	in-bag unique 3/5	out-of-bag A, B

Read the third column

4, 2, 3, 3 distinct rows out of five — an average of $3/5 = 0.60$. Every bag is a different training set; bag 2 has never heard of three of the five rows.

Read the fourth column

Every bag leaves somebody out. The model trained on bag 1 **never saw A or B** — so A and B are honest test rows for it, for free. Hold that thought.

Predict first #3 — how much of the data is in a bag?

Commit to a number

A training set has $n = 1000$ rows. You draw 1000 of them with replacement.

What fraction of the original rows appears at least once?

Predict first #3 — how much of the data is in a bag?

Commit to a number

A training set has $n = 1000$ rows. You draw 1000 of them with replacement.

What fraction of the original rows appears at least once?

0.6323, and it is essentially the same number for every large n .

Most students guess something near 1.0 (“I drew as many as I had”) or exactly 0.5. The true answer is a constant of nature: $1 - 1/e = 0.632121$.

By hand: deriving 63.2%

Fix one row. On a single draw, the chance it is *not* chosen is $1 - \frac{1}{n}$. The n draws are independent, so

$$\Pr[\text{row never chosen}] = \left(1 - \frac{1}{n}\right)^n, \quad \Pr[\text{row in the bag}] = 1 - \left(1 - \frac{1}{n}\right)^n.$$

Take the limit

$$\lim_{n \rightarrow \infty} \left(1 - \frac{1}{n}\right)^n = e^{-1} = 0.367879,$$

so the in-bag fraction tends to

Check it at $n = 5$

$(1 - 1/5)^5 = 0.8^5 = 0.32768$, so the in-bag fraction is 0.67232 — and the four bags on the previous slide averaged 0.60, which is what four draws of a mean-0.672 quantity look like.



The consequence

The same thing in code, and measured



```
rng = np.random.default_rng(0) # the one seed for this lecture
for n in (5, 10, 50, 100, 1000, 10000):
    out = (1 - 1/n) ** n
    emp = np.mean([len(np.unique(rng.integers(0, n, n))) / n
                   for _ in range(2000)])
    print("%6d %.6f %.6f %.6f" % (n, out, 1 - out, emp))
```

n	$(1-1/n)^n$	$1-(1-1/n)^n$	empirical unique fraction
5	0.327680	0.672320	0.674900
10	0.348678	0.651322	0.655600
50	0.364170	0.635830	0.637460
100	0.366032	0.633968	0.634460
1000	0.367695	0.632305	0.632152
10000	0.367861	0.632139	0.632156

limit $1 - 1/e = 0.632121$

Formula and measurement agree to four decimals from $n = 50$ upwards. The convergence is fast: by $n = 100$ you are within 0.002 of the limit.

Drawn: the bootstrap fraction



Right: twenty thousand bootstrap samples of the 398-row training split. Mean 0.6325, standard deviation 0.0158, range 0.5704 to 0.6910. The 63.2% is an average, not a guarantee — individual bags vary by several percentage points.

Bagging in code: one tree against two hundred

```
tree = DecisionTreeClassifier(random_state=0).fit(Xtr, ytr)
bag = BaggingClassifier(DecisionTreeClassifier(random_state=0),
                        n_estimators=200, oob_score=True,
                        random_state=0).fit(Xtr, ytr)
```

```
single unpruned tree    test accuracy 0.9064
bag of 200 such trees   test accuracy 0.9298
bag out-of-bag score    0.9598
OOB minus test          +0.0300
5-fold CV on the training rows 0.9497
```

breast_cancer, 398 training rows, 171 test rows. Two hundred copies of a model that scores 0.9064 give 0.9298 — a gain of +0.0234 with **no change to the learner at all**. Same class, same hyperparameters, same code.

Out-of-bag error: a validation set you did not have to pay for

For each training row, average only the members that *did not* see it. Score those predictions against the truth.

Why it is honest

A row is out-of-bag for about 37% of the members. Those members never used it to choose a split. Their prediction for it is a genuine prediction on unseen data.

Why it is free

No rows are held back and no model is refitted. Training is done with 5-fold cross-validation.

The caveat you must state

Each OOB prediction comes from a sub-ensemble of roughly $0.37B$ members, not all B . For small B that makes OOB *pessimistic*. For B in the hundreds the effect is negligible.


In scikit-learn

`oob_score=True` on

Is the free estimate honest? Twenty splits say yes

On the single split above, OOB read 0.9598 against a test score of 0.9298 — optimistic by 0.0300. One split proves nothing. Repeat it.

```
20 random splits of breast_cancer, forest of 300 trees
mean OOB score      0.9584 (sd 0.0054)
mean held-out score 0.9576 (sd 0.0137)
mean OOB - held-out +0.0008
paired t = 0.209, p = 0.836
```



Bias +0.0008, and $p = 0.836$ — no detectable systematic optimism. Note the standard deviations: OOB 0.0054 against held-out 0.0137. The free estimate is not only unbiased here, it is **less noisy**, because it uses all 398 rows instead of 171.

What averaging can and cannot fix

For squared-error regression the expected error at a point splits exactly:

$$\mathbb{E}[(y - \hat{f}(x))^2] = \underbrace{(\mathbb{E}[\hat{f}(x)] - f(x))^2}_{\text{bias}^2} + \underbrace{\text{Var}[\hat{f}(x)]}_{\text{variance}} + \underbrace{\sigma^2}_{\text{noise}}.$$

Averaging attacks the middle term




The mean of B things has smaller variance than one of them. It has *the same expectation*. So bagging can only move the variance.

Which is why the members should be deep

Use **unpruned** trees inside a bag. High variance is the thing bagging is good at removing; bias is the thing it cannot touch. Prune the members and you have thrown away the wrong half.

Measure this rather than assert it: 300 independent training sets of 120 points from $f(x) = \sin(4x) + x$ with noise $\sigma = 0.35$.

Drawn: three hundred training sets, two learners



Left and centre: forty of the three hundred fitted curves. The single trees scatter wildly around the truth; the bagged curves hug it. Both **averages** (the thick lines) sit on the dashed truth — that is the bias being unchanged.

The numbers behind that picture

```
single deep tree  bias^2 0.0005  variance 0.1230  bias^2+var 0.1235
bag of 50 trees   bias^2 0.0002  variance 0.0601  bias^2+var 0.0603
variance fell to 48.8% of the single tree's; bias^2 moved -0.0003
irreducible noise variance 0.1225
```

Variance: halved

0.1230 $\rightarrow$ 0.0601. That is the entire benefit, and it is worth 0.0632 of squared error per point.

Bias: did not move

0.0005 $\rightarrow$ 0.0002. Both are noise next to the variance term. A deep tree is nearly unbiased to begin with, so there was nothing to fix.

The noise floor is $\sigma^2 = 0.35^2 = 0.1225$. No ensemble of anything can go below it.

The other half: bag a stump and watch nothing happen

A depth-1 stump on a sine wave is badly **biased** — it can only produce two values. If bagging could fix bias, it would show here.

```
single stump      bias^2 0.1856   variance 0.0281   bias^2+var 0.2137
bag of 50 stumps  bias^2 0.1842   variance 0.0147   bias^2+var 0.1989
bagging removed 48% of the stump's variance and left the bias where it was
```

Variance halved again — $0.0281 \rightarrow 0.0147$, the same 48%. And the bias moved by **0.0014**, less than one percent of itself. Fifty stumps are still a stump.

The one-line rule

Bagging is a variance reducer. If your model is underfitting, bagging it will not help. That is boosting's job, and it is why boosting exists.

Random forest = bagging + a random subset of features

At **every split of every tree**, choose `max_features` columns at random and search only those for the best split.

Why bagging alone is not enough

If one feature is strongly predictive, every bagged tree will put it at the root. The trees then look almost the same, and by the right-hand panel of the jury figure, an ensemble of near-identical members is worth barely more than one member.

What the restriction buys

Forbid the dominant feature at some splits and those trees are forced to find a different structure. Each individual tree gets *worse*. The correlation between trees falls further, and the ensemble gets **better**.

Defaults: $\lfloor \sqrt{d} \rfloor$ for classification, $d/3$ for regression. With $d = 30$, $\sqrt{30} = 5.4772$, so scikit-learn's 'sqrt' uses **5**.

By hand: the variance of a correlated average

If B members each have variance σ^2 and every pair has correlation ρ , the variance of their average is

$$\text{Var}\left(\frac{1}{B} \sum_{b=1}^B T_b\right) = \rho \sigma^2 + \frac{1-\rho}{B} \sigma^2.$$

As a fraction of σ^2 , $B = 100$

ρ	$\rho + (1-\rho)/B$
0.00	0.0100
0.10	0.1090
0.30	0.3070
0.50	0.5050
1.00	1.0000

Read the first term

$\rho\sigma^2$ does **not** contain B . Adding trees drives the second term to zero and leaves the first one exactly where it was.

At $\rho = 0.5$, a thousand trees and a hundred trees are both stuck at half the variance of one tree.

So the only lever left

Predict first #4 — worse trees, better forest?

Commit to an ordering

Three ensembles of 100 trees on `breast_cancer` ($d = 30$): plain bagging (`max_features=30`), the default forest (`=5`), and one random feature per split (`=1`). **Rank them by ensemble accuracy, then by average member accuracy.** The two rankings are not the same.

Predict first #4 — worse trees, better forest?

Commit to an ordering

Three ensembles of 100 trees on breast_cancer ($d = 30$): plain bagging ($\text{max_features}=30$), the default forest ($=5$), and one random feature per split ($=1$). **Rank them by ensemble accuracy, then by average member accuracy.** The two rankings are not the same.

bagging ($\text{max_features}=30$)	tree corr 0.8261	mean single tree 0.9130	ensemble 0.9240
forest ($\text{max_features}=5$)	tree corr 0.7945	mean single tree 0.9110	ensemble 0.9532
forest ($\text{max_features}=1$)	tree corr 0.7010	mean single tree 0.8861	ensemble 0.9415

Members: monotone down



0.9130 $\rightarrow$ 0.9110 $\rightarrow$ 0.8861. Restricting the split search can only make each tree worse. It does.

Ensemble: peaks in the middle

0.9240 $\rightarrow$ **0.9532** $\rightarrow$ 0.9415. Correlation falls the whole way (0.8261 $\rightarrow$ 0.7010), so the two effects trade off and the optimum is interior.

$\text{max_features}=1$ decorrelates the most but the trees have become too weak to carry the vote. This trade-off is the entire design of a random forest, and it is why the default is a *small* number, not 1 and not d .

Drawn and tabulated: the whole max_features sweep



max_features	tree corr	mean single tree	OOB	test
1	0.7046	0.8864	0.9673	0.9357
2	0.7582	0.9030	0.9673	0.9532
5	0.7935	0.9110	0.9698	0.9532
12	0.8132	0.9141	0.9598	0.9532
30	0.8282	0.9135	0.9548	0.9240

OOB error is minimised at `max_features=5` (0.0302) against 0.0452 at 30 — the default is right here. And note the OOB column was computed **without touching the test set**: this is how you tune `max_features` honestly.

A forest is worse than every one of its trees, and better than

```
a forest is 300 trees; the first three have 13, 16, 18 leaves
their individual test accuracies 0.8889 0.9123 0.8772
mean over all 300 trees          0.9120
worst single tree                0.8480
the forest                      0.9532
```

No individual tree is any good

The average member scores 0.9120; the worst scores 0.8480. Every one of them is a small overfitted tree grown on 63% of the rows with five features per split.

The forest scores 0.9532

+0.0412 above the average member, and above the *best* of the three shown too. This is the jury theorem happening on real data.

Say the sentence out loud: **the forest is better than any tree in it**. You cannot get that by improving a tree. You get it by having many that fail differently.

What feature importances actually is

Mean decrease in impurity: for each feature, add up the impurity reduction of every split that used it, weighted by the rows reaching that split, and average over the trees.

Why people love it

It is free — the numbers fall out of the fitting. It sums to 1. It needs no extra data.

Why it lies

A column with many distinct values offers many candidate thresholds. More candidates means a better *best* candidate, purely by chance. So it **systematically favours high-cardinality features** — exactly the bias information gain had in the decision-tree lecture, now inherited by the forest.

Two demonstrations follow, deliberately of different strengths. Report both.

Demonstration 1: plant three useless columns in breast

A coin (2 values), a die (6 values) and a row identifier (569 values). All three are drawn from a random number generator and carry **no information whatsoever**.

```
forest test accuracy with the three noise columns added 0.9357
the three planted noise columns:
  noise_uid(569)          0.0028    rank 30 of 33
  noise_dice(6)           0.0012    rank 32 of 33
  noise_coin(2)           0.0005    rank 33 of 33
the useless 569-level column scores 5.9 times the useless 2-level column
it outranks 1 of the 30 real clinical measurements
```

Be honest about the strength of this. The ordering among the three is exactly cardinality order, $5.9\times$ from 2 levels to 569 — the bias is real and measurable. But the thirty clinical measurements are genuinely strong, so the noise still lands at ranks 30, 32, 33. On this dataset the bias distorts; it does not mislead.

Demonstration 2: a table where it does mislead

Five columns, 800 rows. `signal_binary` agrees with `y` 75% of the time. **The other four are pure noise.** No model can beat 0.7500; the forest gets 0.7333.

```
              impurity  permutation
signal_binary(2)    0.2715    +0.2136 (sd 0.0294)
noise_coin(2)       0.0237    +0.0083 (sd 0.0115)
noise_dice(6)       0.0829    +0.0189 (sd 0.0131)
noise_cont(800)     0.3110    +0.0081 (sd 0.0145)
noise_uid(800)      0.3109    +0.0114 (sd 0.0127)
impurity importance puts TWO useless columns above the only informative one
useless 800-level column / useless 2-level column = 13.1 times
```

Impurity importance ranks a **random uniform draw** and a **row identifier** above the only column that carries signal. A student who read the top of that list would delete the one useful feature and keep the two useless ones.

The fix: permutation importance

Score the fitted model on **held-out** rows. Shuffle one column. Score again. The drop is that column's importance. Repeat and average.

Why it is not fooled

It never looks at the tree structure. It asks the only question that matters: *if this column were destroyed, would the predictions get worse?* Destroying a random identifier changes nothing, so it scores about zero.




And it gets the answer right

On the designed table: **+0.2136** for the signal against +0.0081 and +0.0114 for the two 800-level columns. A $19\times$ separation, in the correct direction.

Its own two caveats

Compute it on **held-out** rows, not training rows. And with correlated

Drawn: the two measures side by side



Left: on breast_cancer the three planted columns rise with cardinality even though all three are noise. Centre and right: on the designed table, impurity puts two red bars above the purple one and permutation does not. **Red is noise in every panel.**

Boosting inverts every design decision in bagging

Bagging

Members are trained **in parallel** on **resampled** data, are **deep** (low bias, high variance), and get an **equal** vote.

Boosting

Members are trained **in sequence** on **reweighted** data, are **shallow** (high bias, low variance), and get a **weighted** vote.

The idea in one sentence

Each new member is trained on a version of the data in which the rows the committee currently gets *wrong* have been made heavier.

Because each member attacks the current errors, boosting reduces **bias**. That is the half of the decomposition bagging could not reach — and it is why a boosted ensemble of *stumps* works at all.

AdaBoost, written out

Labels $y_i \in \{-1, +1\}$. Start with $w_i = 1/n$ for all n rows. For $t = 1, \dots, T$:

1. Fit a weak learner h_t to the data weighted by w .
2. Weighted error $\varepsilon_t = \sum_{i: h_t(x_i) \neq y_i} w_i$.
3. Member weight $\alpha_t = \frac{1}{2} \ln\left(\frac{1 - \varepsilon_t}{\varepsilon_t}\right)$.
4. Reweight: $w_i \leftarrow w_i e^{-\alpha_t y_i h_t(x_i)}$, then divide by the sum so the weights add to 1.

Final prediction: $F(x) = \sum_{t=1}^T \alpha_t h_t(x)$, classify by sign $F(x)$.

Read step 4

$y_i h_t(x_i) = +1$ when the member is right, so the weight is multiplied by $e^{-\alpha_t} < 1$ and **shrinks**. It is -1 when the member is wrong, so the weight is multiplied by $e^{+\alpha_t} > 1$ and **grows**. One formula does both.

By hand, round 1: ten points on a line

$x = 0, 1, \dots, 9$ with labels $y = + + + - - - + + + -$. Weak learner: a decision stump $h(x) = \pm 1$ according to $x < \theta$. All weights start at 0.1.

```
round 1: stump  h(x) = +1 if x < 2.5 else -1
               weighted error  eps = 0.3000
               alpha = 0.5*ln((1-eps)/eps) = 0.4236
               misclassified points x = 6, 7, 8
               new weights 0.0714 0.0714 0.0714 0.0714 0.0714 0.0714 0.1667 0.1667 0.1667 0.0714
```

Check ε_1 by hand

Three points wrong, each of weight 0.1:
 $\varepsilon_1 = 0.3$. Then $\alpha_1 = \frac{1}{2} \ln(0.7/0.3) =$
 $\frac{1}{2} \ln 2.3333 = \frac{1}{2}(0.8473) = \mathbf{0.4236}$.

Check the new weights by hand

Wrong: $0.1 e^{0.4236} = 0.15275$, three of them.
 Right: $0.1 e^{-0.4236} = 0.06547$, seven of them.
 Sum = $3(0.15275) + 7(0.06547) = 0.91654$.
 Divide: $0.15275/0.91654 = \mathbf{0.1667}$ and
 $0.06547/0.91654 = \mathbf{0.0714}$.

By hand, rounds 2 and 3

```
round 2: stump  h(x) = +1 if x < 8.5 else -1
              weighted error  eps = 0.2143
              alpha = 0.5*ln((1-eps)/eps) = 0.6496
              misclassified points x = 3, 4, 5
              new weights 0.0455 0.0455 0.0455 0.1667 0.1667 0.1667 0.1061 0.1061 0.1061 0.0455

round 3: stump  h(x) = -1 if x < 5.5 else +1
              weighted error  eps = 0.1818
              alpha = 0.5*ln((1-eps)/eps) = 0.7520
              misclassified points x = 0, 1, 2, 9
              new weights 0.1250 0.1250 0.1250 0.1019 0.1019 0.1019 0.0648 0.0648 0.0648 0.1250
```

Follow the story. Round 1 fails on 6, 7, 8, so round 2 arrives with those three heavy — and picks a stump that gets all three right, at the price of 3, 4, 5. Round 3 then arrives with 3, 4, 5 heavy and *flips polarity* to fix them.

Watch ϵ fall, $0.3000 \rightarrow 0.2143 \rightarrow 0.1818$, and α rise, $0.4236 \rightarrow 0.6496 \rightarrow 0.7520$. Each member is trusted more than the last, because each one solved a harder problem.

Predict first #5 — can three stumps beat the best stump

Commit

The best single stump on these ten points misclassifies **three** of them, and no stump can do better — the labels are + + + - - - + + + -, which no single threshold can separate.

How many of the ten does the weighted vote of the three stumps get right?

Predict first #5 — can three stumps beat the best stump

Commit

The best single stump on these ten points misclassifies **three** of them, and no stump can do better — the labels are + + + - - - + + + -, which no single threshold can separate.

How many of the ten does the weighted vote of the three stumps get right?

10 of 10. Training error zero.

Three models, each of which *cannot* solve the problem, combine into one that does. This is the sense in which boosting builds a strong learner out of weak ones.

The reveal: $F(x) = \alpha_1 h_1 + \alpha_2 h_2 + \alpha_3 h_3$


x	f1	f2	f3	F(x)	sign	y
0	+1.0000	+1.0000	-1.0000	+0.3213	+1	+1
1	+1.0000	+1.0000	-1.0000	+0.3213	+1	+1
2	+1.0000	+1.0000	-1.0000	+0.3213	+1	+1
3	-1.0000	+1.0000	-1.0000	-0.5260	-1	-1
4	-1.0000	+1.0000	-1.0000	-0.5260	-1	-1
5	-1.0000	+1.0000	-1.0000	-0.5260	-1	-1
6	-1.0000	+1.0000	+1.0000	+0.9780	+1	+1
7	-1.0000	+1.0000	+1.0000	+0.9780	+1	+1
8	-1.0000	+1.0000	+1.0000	+0.9780	+1	+1
9	-1.0000	-1.0000	+1.0000	-0.3213	-1	-1

training errors of the 3-round ensemble: 0

training errors of the best single stump: 3

Check row 0 with a pen: $+0.4236 + 0.6496 - 0.7520 = +0.3212$. And row 6:
 $-0.4236 + 0.6496 + 0.7520 = +0.9780$. **Not one column of f_1, f_2, f_3 matches y — but their weighted sum does.**

Drawn: the weights growing and the margin appearing



Marker area is the row's weight. Red circles are the points that round's stump got wrong — and they are exactly the points that grow before the next panel. The last panel is $F(x)$: green where the sign is right.

And now the identical result from AdaBoostClassifier

```
ada = AdaBoostClassifier(DecisionTreeClassifier(max_depth=1, random_state=0),  
                          n_estimators=3, random_state=0).fit(x[:, None], y)
```

```
sklearn estimator_errors_ 0.3000 0.2143 0.1818  
hand-computed eps        0.3000 0.2143 0.1818  
sklearn estimator_weights_ 0.8473 1.2993 1.5041  
hand-computed alpha       0.4236 0.6496 0.7520  
ratio sklearn/hand        2.0000 2.0000 2.0000  
sklearn training errors 0, stump thresholds [2.5, 8.5, 5.5]
```

Every ε matches to four decimals and every threshold matches exactly. The α 's differ by **exactly a factor of two**, every time: scikit-learn implements SAMME, whose weight is $\ln \frac{1-\varepsilon}{\varepsilon} + \ln(K-1)$, and at $K=2$ the second term is $\ln 1 = 0$, leaving 2α . A constant factor on every member cannot change sign F , so the predictions are identical.


Gradient boosting: the same idea, with residuals instead

For squared-error regression there is nothing to reweight — so fit the next member to **what is left over**. Eight points, $y = 5, 5, 6, 6, 9, 9, 10, 10$.

```
stage 0: F0 = mean(y) = 7.5000
  residuals -2.500 -2.500 -1.500 -1.500 +1.500 +1.500 +2.500 +2.500
stage 1: stump on the residuals, split at x < 4.5
  stump output -2.000 -2.000 -2.000 -2.000 +2.000 +2.000 +2.000 +2.000
  F1           +5.500 +5.500 +5.500 +5.500 +9.500 +9.500 +9.500 +9.500
  SSE 2.0000
stage 2: stump on the residuals, split at x < 2.5
  F2           +5.000 +5.000 +5.667 +5.667 +9.667 +9.667 +9.667 +9.667
  SSE 1.3333
stage 3: stump on the residuals, split at x < 6.5
  F3           +4.889 +4.889 +5.556 +5.556 +9.556 +9.556 +10.000 +10.000
  SSE 1.0370
```

F_0 is the mean. Every stage after it fits a stump to $y - F_{m-1}$ and adds it on. SSE
 $34 \rightarrow 2 \rightarrow 1.333 \rightarrow 1.037$.

Drawn: fitting the residuals



Top row: the staircase F_m closing on the points. Bottom row: the residuals it is chasing, shrinking towards zero. **Boosting is gradient descent in function space** — for squared error the negative gradient *is* the residual.

The library does exactly that

```
gbm = GradientBoostingRegressor(n_estimators=3, learning_rate=1.0,  
                                max_depth=1,  
                                random_state=0).fit(xr[:, None], yr)
```

```
sklearn init_.constant_ 7.5000 (hand: mean = 7.5000)  
sklearn predictions +4.889 +4.889 +5.556 +5.556 +9.556 +9.556 +10.000 +10.000  
hand F3             +4.889 +4.889 +5.556 +5.556 +9.556 +9.556 +10.000 +10.000  
max abs difference  0.00e+00
```

One knob that matters more than any other

`learning_rate` was set to 1.0 here so the hand computation and the library agree exactly. In practice use 0.05 to 0.1: add $\eta \times$ the stump instead of the whole stump. Small steps and many of them generalise better than a few large ones — and, unlike a forest, `n_estimators` then becomes a parameter you must actually tune.

The comparison, sharply

	Bagging / forests	Boosting
Members trained	in parallel, independently	in sequence, each on the last's errors
Data each sees	a bootstrap resample	the same rows, reweighted
Member depth	deep, unpruned	shallow — stumps or depth 2–3
Vote	equal weight	weight α_t , rises as ε_t falls
Attacks	variance	bias
More members	plateaus, safe	can overfit, must be watched
Label noise	robust — averages it away	fragile — chases it
Parallelisable	yes, trivially	no, inherently sequential
Tuning needed	almost none	learning rate, depth, T

The row that costs people real money


A mislabelled row is a row every member keeps getting wrong. Bagging averages over it.

Boosting makes it heavier and heavier until the ensemble is built around it.

Demonstrated: flip the training labels and watch

Flip a fraction of the *training* labels; test on clean labels; average over ten splits.

noise	RandomForest(200)	AdaBoost(200 stumps)	single tree
0%	0.9596	0.9684	0.9322
10%	0.9503	0.9129	0.8298
20%	0.9368	0.8789	0.7415
30%	0.8871	0.8240	0.6661
40%	0.7327	0.6725	0.5637



At 0% noise **AdaBoost wins** — by 0.0088. Say that before you say anything else. Then follow the columns down: by 10% the forest is ahead by 0.0374, and by 40% by 0.0602. Error bars are the standard deviation over the ten splits; both ensembles beat the single tree everywhere, by a wide margin.

How many estimators? Forests plateau; boosting can t

25% of the training labels flipped

trees	forest test acc	AdaBoost test acc
1	0.6959	0.9006
10	0.8129	0.9006
50	0.9298	0.8480
100	0.9298	0.8480
200	0.9298	0.8129
400	0.9298	0.8070
800	0.9240	0.7719

forest: best 0.9298 at $T=50$, final 0.9240, drop from peak -0.0058
AdaBoost: best 0.9064 at $T=19$, final 0.7719, drop from peak -0.1345
AdaBoost training accuracy at $T=800$: 0.9648


The forest climbs and then flattens: from $T = 50$ to $T = 800$ it moves by 0.0058, which is one test row. AdaBoost peaks at $T = 19$ and then loses **0.1345** — while its *training* accuracy climbs to 0.9648 on labels that are 25% wrong. That is the definition of overfitting.

And on clean labels, so the comparison is fair

clean labels, same split

AdaBoost: best 0.9766 at $T=324$, final 0.9591, drop from peak -0.0175

forest: best 0.9532 at $T=100$, final 0.9415, drop from peak -0.0117



On clean data AdaBoost reaches 0.9766, above the forest's 0.9532, and gives back 0.0175 by $T = 800$; the forest gives back 0.0117 from its own peak. **Two or three test rows each — jitter, not a trend.** The difference is the shape: the forest is flat from $T = 50$ on both panels, AdaBoost falls by 0.1345 on the noisy one and is still falling.

All five, on three bundled datasets, five-fold cross-validation

	wine	breast	digits
single tree	0.9273	0.9262	0.8592
bagging 200	0.9608	0.9579	0.9533
forest 200	0.9832	0.9667	0.9766
AdaBoost 200	0.9665	0.9754	0.8458
grad boosting	0.9494	0.9649	0.9655

Every ensemble beats the tree

On every dataset, with one exception — and the exception is the interesting part.

AdaBoost on digits: 0.8458

Below the single tree. Ten classes, and SAMME needs $\varepsilon_t < 1 - 1/K = 0.9$ every round; a stump on a ten-class problem is barely above chance. **Boosted stumps are a two-class instrument.**

A random forest with default settings is top or near-top on all three and needed no tuning at all. That is why it is the standard first thing to try on tabular data.

Summary — twelve lines

1. An ensemble trains many models and combines them. It helps only when the members are **better than chance** and **make different mistakes**.
2. Twenty-five *independent* voters at $p = 0.60$ vote correctly with probability **0.846232**, a gain of +0.2462. Below $p = 0.5$ the same machinery makes things worse: $0.40 \rightarrow 0.0209$ at 101 voters.
3. Correlation destroys that gain without touching any member's accuracy: mean pairwise correlation 0.126 drops the vote from 0.8458 to **0.6932**; at correlation 1 it is 0.5993.
4. A bootstrap sample of n from n contains $1 - (1 - 1/n)^n \rightarrow 1 - 1/e = \mathbf{0.632121}$ of the rows. Measured at $n = 10\,000$: 0.632156.
5. The other 37% is the **out-of-bag** set — free validation. Over 20 splits, OOB minus held-out was **+0.0008** ($p = 0.836$), with *smaller* standard deviation than the held-out score.
6. Bagging attacks variance only. 300 resamples: variance $0.1230 \rightarrow \mathbf{0.0601}$, $\text{bias}^2 \ 0.0005 \rightarrow 0.0002$. Bagging fifty *stumps* halved the variance and left bias^2 at $0.1856 \rightarrow 0.1842$.
7. $\text{Var}(\bar{T}) = \rho\sigma^2 + \frac{1-\rho}{B}\sigma^2$. The first term has no B in it, so once B is large the only remaining lever is ρ .
8. A forest pulls ρ down with `max_features`: $0.8261 \rightarrow 0.7945 \rightarrow 0.7010$ for 30, 5, 1 features. Members get worse ($0.9130 \rightarrow 0.8861$) and the ensemble peaks in the middle at **0.9532**. Mean member 0.9120, worst member 0.8480, forest 0.9532.
9. Impurity importance is biased towards high-cardinality columns. On a designed table, two **pure-noise** 800-level columns (0.3110, 0.3109) outranked the only informative one (0.2715). Permutation importance on held-out rows got it right: **+0.2136** against +0.0081 and +0.0114.
10. AdaBoost on ten points: $\varepsilon = 0.3000, 0.2143, 0.1818$ and $\alpha = 0.4236, 0.6496, 0.7520$. Three stumps that each get three points wrong combine to get **10** of 10 right. scikit-learn reproduces every ε and every threshold; its weights are exactly 2α because SAMME uses $\ln \frac{1-\varepsilon}{\varepsilon}$.
11. Gradient boosting fits the residuals: $F_0 = 7.5$, then stumps on $y - F_{m-1}$, $\text{SSE } 34 \rightarrow 2 \rightarrow 1.333 \rightarrow \mathbf{1.037}$, matching GradientBoostingRegressor to 0.00e+00.
12. Label noise separates them. At 0% AdaBoost wins by 0.0088; at 40% the forest wins by **0.0602**. With 25% flipped labels AdaBoost peaks at $T = 19$ then loses 0.1345 while its training accuracy climbs to 0.9648. The forest's peak-to-final drop is 0.0058, one test row.

Before the next lecture

Read

notes.pdf — the jury theorem derived, the bootstrap derived, three rounds of AdaBoost with every multiplication shown, and **practice questions with full worked solutions**.

Do

Reproduce the three AdaBoost rounds with a pen: ε_t , α_t , the reweighting, the normalisation, and sign $F(x)$ for all ten points. Then compute $1 - (1 - 1/n)^n$ for $n = 5$ and $n = 10$ and say why the two answers are so close.

The habit to take away

On tabular data, fit a random forest with default settings **first**. It needs no scaling, no tuning and no held-out set (use `oob_score=True`), and it tells you immediately what accuracy is achievable. Everything else is measured against it.

And the habit to unlearn

Do not read `feature_importances_` and believe it. Compute permutation importance on held-out rows before you delete a column or write a sentence about which variables matter.

Materials: <https://github.com/tali7c/Applied-Machine-Learning>

Sources: Han, Kamber and Pei, *Data Mining: Concepts and Techniques* (3e), Ch. 8–9.; James et al., *An Introduction to Statistical Learning with Python*, Ch. 8–9.; Bishop, *Pattern Recognition and Machine Learning*, Ch. 4–5, 7, 14.; Zhang et al., *Dive into Deep Learning*, Ch. 5.