



**"RANKED 45
IN INDIA**

#University Category



**WORLD
UNIVERSITY
RANKINGS**

NO.1 PVT. UNIVERSITY IN
ACADEMIC REPUTATION IN INDIA



ACCREDITED GRADE 'A'
BY NAAC



PERFECT SCORE OF **150/150** AS A TESTAMENT
TO EXCEPTIONAL E-LEARNING METHODS

1



Applied Machine Learning

Unit 06 – Lecture 19: Classification Model Evaluation: ROC and AUC

Tofik Ali

School of Computer Science, UPES Dehradun

Autumn 2026

The last lecture of the unit, and the one that decides which of the others you use

Topic focus: **a classifier outputs a score, not a class**

Repository: <https://github.com/tali7c/Applied-Machine-Learning>

Sources: Han, Kamber and Pei, *Data Mining: Concepts and Techniques* (3e), Ch. 8–9.; James et al., *An Introduction to Statistical Learning with Python*, Ch. 8–9.; Bishop, *Pattern Recognition and Machine Learning*, Ch. 4–5, 7, 14.; Zhang et al., *Dive into Deep Learning*, Ch. 5.



Accuracy
000

Confusion
00000

Scores
0000

ROC
000000

AUC
0000000

ROC/PR
000

Threshold
00000

Multiclass
000000

Selection
000000

Summary
00

Quick Links

[Why not accuracy](#)

[Confusion matrix](#)

[Scores](#)

[ROC by hand](#)

[AUC](#)

[ROC vs PR](#)

[Thresholds](#)

[Multiclass](#)

[Selection](#)

[Summary](#)

Agenda

Accuracy, and the one place it lies

The confusion matrix

Scores, not labels

Building an ROC curve by hand

AUC, computed two ways

ROC against precision–recall

Choosing the threshold is a business decision

More than two classes

Model selection, done properly

Summary

How this lecture works

Every idea appears twice

1. **By hand** — twenty patients, ten scored rows, a hundred rows in three classes. Small enough to check on paper, and small enough to appear on a test.
2. **In code** — the same arithmetic through `sklearn.metrics`, so you see that `roc_curve` is bookkeeping on a sorted list and nothing magic.

Commit before you look

Six frames ask you to predict a number before it is revealed. Write your guess down. Being wrong on paper is how the idea sticks.

Three numbers to carry out: at 1% positives a useless model scores 0.9900 accuracy and a useful one 0.9897; the same scores give **ROC AUC** 0.9200 and **average precision** 0.1702; and the best two of seven classifiers differ by 0.0003 against a fold-to-fold spread of 0.0058.

Where we are: six models, no way to choose

The unit built the models

Nearest neighbours, decision trees, Bayesian classifiers, ensembles, neural networks, support vector machines. Every one of those lectures ended with *one accuracy on one split*, and we treated that number as though it settled something.

It settles nothing. Two questions are open

- (1) **Is this model any good?** An accuracy is meaningless until you know what a *useless* model scores on your problem.
- (2) **Is it better than that one?** A single number from a single split has a standard error you have not measured.

The whole lecture in one line: **a classifier does not output a class — it outputs a score, and a threshold turns that score into a class.** Accuracy, precision, recall and F_1 describe *one* threshold. ROC describes all of them at once.

One seed — and the two things that are not it

The one seed

Everything arbitrary in this deck uses **seed 0**: every `train_test_split`, every subsample of the negatives, the coin-flip scores, every estimator's `random_state`. Every listing that samples shows it.

A seed that defines data

`make_classification(random_state=0)` fixes the 99:1 problem and `random_state=1` the three-class one. Change either and you have *different data*. They never move.

A sweep

The model-selection section scores every model over $5 \times 4 = 20$ folds. The *spread* between those folds is the argument — it is what tells you a gap of 0.002 means nothing. So it stays a sweep, generated from seed 0.

Why this matters to you

Run the code as written and you get the numbers on the slide, *including* the ones whose whole purpose is to vary.

Predict first #1

Accuracy is the fraction of predictions that were right — but the weights in that average are the **class frequencies**. A synthetic table: 20,000 rows, six informative features, and a 99:1 class ratio, so a stratified 70/30 split leaves 6000 held-out rows with 60 positives. Two models: a `DummyClassifier` that answers “negative” for every row, and a scaled `LogisticRegression` that has looked at all 25 features.

What accuracy does each score — and which of the two is higher?

Predict first #1

Accuracy is the fraction of predictions that were right — but the weights in that average are the **class frequencies**. A synthetic table: 20,000 rows, six informative features, and a 99:1 class ratio, so a stratified 70/30 split leaves 6000 held-out rows with 60 positives. Two models: a `DummyClassifier` that answers “negative” for every row, and a scaled `LogisticRegression` that has looked at all 25 features.

What accuracy does each score — and which of the two is higher?

```
rows 20000    positives 200    negatives 19800    prevalence 0.0100

A model that always answers 'negative', on the 6000 held-out rows:
accuracy 0.9900
confusion matrix
[[5940    0]
 [  60    0]]
precision 0.0000    recall 0.0000    F1 0.0000
it found 0 of the 60 positives
```

Predict first #1

Accuracy is the fraction of predictions that were right — but the weights in that average are the **class frequencies**. A synthetic table: 20,000 rows, six informative features, and a 99:1 class ratio, so a stratified 70/30 split leaves 6000 held-out rows with 60 positives. Two models: a `DummyClassifier` that answers “negative” for every row, and a scaled `LogisticRegression` that has looked at all 25 features.

What accuracy does each score — and which of the two is higher?

```
rows 20000    positives 200    negatives 19800    prevalence 0.0100

A model that always answers 'negative', on the 6000 held-out rows:
accuracy 0.9900
confusion matrix
[[5940    0]
 [  60    0]]
precision 0.0000    recall 0.0000    F1 0.0000
it found 0 of the 60 positives
```

A model that has not looked at a single feature scores 99%. Write that number on a slide and it reads as a triumph. **Now fit something real.**

The result that should stop you

An actual logistic regression on the same split:

accuracy 0.9897 (the constant model got 0.9900)

accuracy gained over the constant model: -0.0003

precision 0.3333 recall 0.0333 F1 0.0606

it predicted 'positive' 6 times and was right 2 times

ROC AUC 0.9200 average precision 0.1702

The useful model scores LOWER accuracy than the useless one.

The result that should stop you

An actual logistic regression on the same split:

```
accuracy 0.9897 (the constant model got 0.9900)
accuracy gained over the constant model: -0.0003
precision 0.3333 recall 0.0333 F1 0.0606
it predicted 'positive' 6 times and was right 2 times
ROC AUC 0.9200 average precision 0.1702
```

The useful model scores LOWER accuracy than the useless one.

Accuracy and AUC rank the two models in *opposite* orders

The logistic regression is **worse than the constant model by accuracy** (0.9897 against 0.9900) and its ROC AUC is **0.9200** against the constant model's 0.5000. These are not two metrics that disagree about the size of a difference; they disagree about its *sign*. **Report accuracy on an imbalanced problem and you will throw away the model that works.**

The result that should stop you

An actual logistic regression on the same split:

```
accuracy 0.9897 (the constant model got 0.9900)
accuracy gained over the constant model: -0.0003
precision 0.3333 recall 0.0333 F1 0.0606
it predicted 'positive' 6 times and was right 2 times
ROC AUC 0.9200 average precision 0.1702
```

The useful model scores LOWER accuracy than the useless one.


Accuracy and AUC rank the two models in *opposite* orders

The logistic regression is **worse than the constant model by accuracy** (0.9897 against 0.9900) and its ROC AUC is **0.9200** against the constant model's 0.5000. These are not two metrics that disagree about the size of a difference; they disagree about its *sign*. **Report accuracy on an imbalanced problem and you will throw away the model that works.**


Two rescues, and they are the rest of the lecture: stop looking at one number and look at the **four underneath it**; then stop looking at one **threshold**.

Drawn: the same two models, four ways

60 positives in 6000 rows (1%): accuracy says nothing



The two threshold-free summaries disagree sharply



Left: four metrics for the constant model and for logistic regression on the 99:1 problem. **The accuracy bars are indistinguishable**; precision and recall are not. Right: two threshold-free summaries of the *same* logistic regression, each against its own baseline — ROC AUC 0.9200 against 0.5, average precision 0.1702 against the prevalence 0.0100. They disagree, and we will see exactly why.

Four boxes, and the naming rule

	predicted negative	predicted positive	row total
actually negative	TN true negative	FP false positive	$N = FP + TN$
actually positive	FN false negative	TP true positive	$P = TP + FN$

The second word is what the model said; the first word is whether it was right. A false positive is a row the model called positive and was wrong about.

By hand. Twenty people are screened. Ten have the disease (label 1), ten do not. The model says “disease” for six of the ten who have it, and for two of the ten who do not.

patient	1	2	3	4	5	6	7	8	9	10
actual	1	1	1	1	1	1	1	1	1	1
predicted	1	1	1	1	1	1	0	0	0	0
cell	TP	TP	TP	TP	TP	TP	FN	FN	FN	FN

patient	11	12	13	14	15	16	17	18	19	20
actual	0	0	0	0	0	0	0	0	0	0
predicted	1	1	0	0	0	0	0	0	0	0
cell	FP	FP	TN	TN	TN	TN	TN	TN	TN	TN

Count them: $TP = 6$, $FN = 4$, $FP = 2$, $TN = 8$, so $P = 10$ and $N = 10$.

By hand: six metrics from four counts

$$\text{accuracy} = \frac{TP + TN}{P + N} = \frac{6 + 8}{20} = 0.7000$$

$$\text{precision} = \frac{TP}{TP + FP} = \frac{6}{8} = 0.7500$$

$$\text{recall} = \text{TPR} = \frac{TP}{P} = \frac{6}{10} = 0.6000$$

$$\text{specificity} = \frac{TN}{N} = \frac{8}{10} = 0.8000$$

$$\text{FPR} = \frac{FP}{N} = \frac{2}{10} = 0.2000$$

$$F_1 = \frac{2pr}{p + r} = \frac{2(0.75)(0.60)}{1.35} = \frac{0.9}{1.35} = 0.6667$$

Six numbers, one source

Every metric in this lecture is a ratio of two of {TP, FP, FN, TN}. Nothing else is going on. Learn the four counts and you can rebuild the rest under examination conditions.

Two habits

- (1) Write P and N down first — most arithmetic errors are a wrong denominator.
- (2) FPR and specificity always sum to 1. Compute one, subtract for the other, and you have a free check.

The same twenty rows, in code

```
y_true = np.array([1]*10 + [0]*10)
y_pred = np.array([1]*6 + [0]*4 + [1]*2 + [0]*8)
print(confusion_matrix(y_true, y_pred))
print(classification_report(y_true, y_pred, digits=4,
                             target_names=["healthy (0)", "disease (1)"]))
```

```
confusion_matrix(y_true, y_pred) =
[[8 2]
 [4 6]]
rows = actual, cols = predicted, label order [0, 1]
so cm[1,1]=TP=6  cm[0,1]=FP=2  cm[1,0]=FN=4  cm[0,0]=TN=8
```

	precision	recall	f1-score	support
healthy (0)	0.6667	0.8000	0.7273	10
disease (1)	0.7500	0.6000	0.6667	10
accuracy			0.7000	20
macro avg	0.7083	0.7000	0.6970	20
weighted avg	0.7083	0.7000	0.6970	20

The same twenty rows, in code

```
y_true = np.array([1]*10 + [0]*10)
y_pred = np.array([1]*6 + [0]*4 + [1]*2 + [0]*8)
print(confusion_matrix(y_true, y_pred))
print(classification_report(y_true, y_pred, digits=4,
                             target_names=["healthy (0)", "disease (1)"]))
```


```
confusion_matrix(y_true, y_pred) =
[[8 2]
 [4 6]]
rows = actual, cols = predicted, label order [0, 1]
so cm[1,1]=TP=6  cm[0,1]=FP=2  cm[1,0]=FN=4  cm[0,0]=TN=8
```

	precision	recall	f1-score	support
healthy (0)	0.6667	0.8000	0.7273	10
disease (1)	0.7500	0.6000	0.6667	10
accuracy			0.7000	20
macro avg	0.7083	0.7000	0.6970	20
weighted avg	0.7083	0.7000	0.6970	20

The layout that catches everybody

`confusion_matrix` puts **actual on the rows, predicted on the columns**, labels in sorted order — so the top-left cell is TN and the bottom-right is TP. That is the *transpose* of Han, Kamber and Pei §8.5.1 and of most textbook diagrams. Print it and check one cell against a count you did yourself.

Drawn: four counts, six rates



name	formula	denominator is	answers
recall / sensitivity / TPR	TP/P	the actual positives	of the ill, how many did we find?
specificity / TNR	TN/N	the actual negatives	of the well, how many did we clear?
FPR	FP/N	the actual negatives	of the well, how many did we alarm?
precision / PPV	$TP/(TP+FP)$	the <i>predicted</i> positives	of our alarms, how many were real?

Precision is the odd one out

Recall and FPR divide by a *class size*, so they do not move when you change how many negatives you happened to collect. Precision divides by TP + FP, which the **prevalence** controls. **That one difference is why the ROC curve is blind to imbalance and the precision–recall curve is not.**

Why the harmonic mean

F_1 combines precision and recall, and the harmonic mean is *dominated by the smaller of the two*. A model that fires almost never but is always right when it does: $p = 1.0$, $r = 0.02$. Arithmetic mean 0.51; harmonic mean **0.0392**. The first flatters, the second reports.

F_β , when a miss and an alarm differ

$$F_\beta = (1 + \beta^2) \frac{p \times r}{\beta^2 p + r}$$

$\beta = 1$ is F_1 . F_2 weights recall four times as heavily and is the choice when a miss hurts more. **Choosing β is a business decision that has walked into the arithmetic.**

predict applies a threshold you never chose

model	score	default rule
LogisticRegression	<code>predict_proba(X)[:, 1]</code>	positive if ≥ 0.5
RandomForestClassifier	<code>predict_proba(X)[:, 1]</code>	positive if ≥ 0.5
GaussianNB	<code>predict_proba(X)[:, 1]</code>	positive if ≥ 0.5
KNeighborsClassifier	fraction of the k neighbours	positive if ≥ 0.5
SVC	<code>decision_function(X)</code>	positive if ≥ 0
MLPClassifier	<code>predict_proba(X)[:, 1]</code>	positive if ≥ 0.5

Everything in the previous section describes *one* confusion matrix, and a confusion matrix needs hard labels. **But a classifier does not naturally produce hard labels.** It produces a score, and `predict` silently compares it with 0.5 — a number nobody in the room chose.

The model we use for the rest of the lecture

Deliberately *imperfect*: logistic regression on three of the thirty breast_cancer columns, so that the choice of threshold has visible consequences.

```
COLS = [0, 1, 4] # mean radius, mean texture, mean smoothness
Xb_tr, Xb_te, yb_tr, yb_te = train_test_split(
    Xb[:, COLS], yb, test_size=0.3, stratify=yb, random_state=0)
clf = make_pipeline(StandardScaler(),
                    LogisticRegression(max_iter=5000)).fit(Xb_tr, yb_tr)
sb = clf.predict_proba(Xb_te)[: , 1]
```

```
features used: ['mean radius', 'mean texture', 'mean smoothness']
held-out rows 171  benign (positive) 107  malignant 64
ROC AUC 0.9661  accuracy at 0.5 0.9064
for comparison, the same split with all 30 features: ROC AUC 0.9956
we use the three-feature model from here on: a near-perfect model
makes every threshold look equally good, which teaches nothing.
```

The model we use for the rest of the lecture

Deliberately *imperfect*: logistic regression on three of the thirty `breast_cancer` columns, so that the choice of threshold has visible consequences.

```
COLS = [0, 1, 4] # mean radius, mean texture, mean smoothness
Xb_tr, Xb_te, yb_tr, yb_te = train_test_split(
    Xb[:, COLS], yb, test_size=0.3, stratify=yb, random_state=0)
clf = make_pipeline(StandardScaler(),
                    LogisticRegression(max_iter=5000)).fit(Xb_tr, yb_tr)
sb = clf.predict_proba(Xb_te)[: , 1]
```

```
features used: ['mean radius', 'mean texture', 'mean smoothness']
held-out rows 171   benign (positive) 107   malignant 64
ROC AUC 0.9661   accuracy at 0.5 0.9064
for comparison, the same split with all 30 features: ROC AUC 0.9956
we use the three-feature model from here on: a near-perfect model
makes every threshold look equally good, which teaches nothing.
```

Note which class is positive: in `load_breast_cancer` the label 1 is **benign**, the harmless one.
Every metric today is asymmetric between the classes, so decide which class is positive and say so out loud before you compute anything.

One set of scores, seven confusion matrices

```
for t in [0.05, 0.20, 0.35, 0.50, 0.65, 0.80, 0.95]:  
    p = (sb >= t).astype(int)  
    print(t, accuracy_score(yb_te, p), precision_score(yb_te, p),  
          recall_score(yb_te, p), f1_score(yb_te, p))  
print("ROC AUC", roc_auc_score(yb_te, sb))
```

one set of scores, one AUC, many thresholds

threshold	TP	FP	FN	TN	accuracy	precision	recall	F1
0.05	107	28	0	36	0.8363	0.7926	1.0000	0.8843
0.20	107	19	0	45	0.8889	0.8492	1.0000	0.9185
0.35	106	13	1	51	0.9181	0.8908	0.9907	0.9381
0.50	100	9	7	55	0.9064	0.9174	0.9346	0.9259
0.65	94	4	13	60	0.9006	0.9592	0.8785	0.9171
0.80	86	4	21	60	0.8538	0.9556	0.8037	0.8731
0.95	58	0	49	64	0.7135	1.0000	0.5421	0.7030

ROC AUC, for every row above: 0.9661

AUC is a property of the scores. The threshold is a separate decision.

One set of scores, seven confusion matrices



```
for t in [0.05, 0.20, 0.35, 0.50, 0.65, 0.80, 0.95]:
    p = (sb >= t).astype(int)
    print(t, accuracy_score(yb_te, p), precision_score(yb_te, p),
          recall_score(yb_te, p), f1_score(yb_te, p))
print("ROC AUC", roc_auc_score(yb_te, sb))
```

```
one set of scores, one AUC, many thresholds
threshold TP FP FN TN accuracy precision recall F1
0.05      107 28 0 36 0.8363 0.7926 1.0000 0.8843
0.20      107 19 0 45 0.8889 0.8492 1.0000 0.9185
0.35      106 13 1 51 0.9181 0.8908 0.9907 0.9381
0.50      100 9 7 55 0.9064 0.9174 0.9346 0.9259
0.65      94 4 13 60 0.9006 0.9592 0.8785 0.9171
0.80      86 4 21 60 0.8538 0.9556 0.8037 0.8731
0.95      58 0 49 64 0.7135 1.0000 0.5421 0.7030
ROC AUC, for every row above: 0.9661
AUC is a property of the scores. The threshold is a separate decision.
```

Accuracy 0.7135 to 0.9181, from one fitted model

Lower the threshold and recall climbs to 1.0000 while precision falls to 0.7926; raise it and the reverse. There is no threshold at which both improve, because raising it can only move rows from “predicted positive” to “predicted negative”. **Quoting a single accuracy is quoting a model and an unstated threshold as if they were one thing.**

Drawn: the whole threshold sweep, two ways



Left: the four threshold-dependent metrics as continuous functions of the threshold. Right: **the same information as a single curve**. Each labelled point is one row of the table you just saw; the curve is every threshold at once. **That curve is the object this lecture is about**, and we are going to build it by hand before we let a library draw it.

Ten scored examples, by hand

Five positives, five negatives, each with a score between 0 and 1. Sorted by score, highest first.

rank	1	2	3	4	5	6	7	8	9	10
label	POS	POS	neg	POS	POS	neg	neg	POS	neg	neg
score	0.95	0.90	0.85	0.75	0.60	0.55	0.45	0.35	0.25	0.10

$$P = 5, N = 5$$

The model is not perfect — a negative sits at rank 3, above two positives — but it is clearly better than nothing. Those two facts are what the curve will show.

Get a pen out now

We are going to sweep the threshold down this list and compute (FPR, TPR) at every step. **Do it with me.** The library will not appear for four frames, and when it does its output must match yours row for row.

The recipe: up for POS, right for neg

Start with the threshold *above every score*: nothing is predicted positive, $TP = FP = 0$, and the point is $(0, 0)$. Now lower the threshold to each score in turn.

Cross a POSITIVE

One row moves from FN to TP. So TPR rises by $1/P = 0.2$ and FPR does not move.

A step UP.

Cross a NEGATIVE

One row moves from TN to FP. So FPR rises by $1/N = 0.2$ and TPR does not move.

A step RIGHT.

Ten rows, ten steps, eleven points — and the two endpoints are guaranteed: $(0, 0)$ when nothing is called positive, $(1, 1)$ when everything is. **Every ROC curve you will ever see runs between those two corners.**

By hand: the first five steps, one at a time

The sorted labels are **POS, POS, neg, POS, POS, ...**

- **Threshold above everything.** $TP = 0$, $FP = 0$, $FN = 5$, $TN = 5$. Point (0.0, 0.0). Start here, always.

By hand: the first five steps, one at a time

The sorted labels are **POS, POS, neg, POS, POS, ...**

- **Threshold above everything.** $TP = 0$, $FP = 0$, $FN = 5$, $TN = 5$. Point $(0.0, 0.0)$. Start here, always.
- $t = 0.95$, a **POSITIVE**. $TP = 1$, $FP = 0$. $FPR = 0/5 = 0.0$, $TPR = 1/5 = 0.2$. Point $(0.0, 0.2)$ — *up*.

By hand: the first five steps, one at a time

The sorted labels are **POS, POS, neg, POS, POS, ...**

- **Threshold above everything.** $TP = 0$, $FP = 0$, $FN = 5$, $TN = 5$. Point $(0.0, 0.0)$. Start here, always.
- $t = 0.95$, a **POSITIVE**. $TP = 1$, $FP = 0$. $FPR = 0/5 = 0.0$, $TPR = 1/5 = 0.2$. Point $(0.0, 0.2)$ — *up*.
- $t = 0.90$, a **POSITIVE**. $TP = 2$, $FP = 0$. Point $(0.0, 0.4)$ — *up* again. We are climbing the left-hand edge, which is exactly what a good model does.

By hand: the first five steps, one at a time

The sorted labels are **POS, POS, neg, POS, POS, ...**

- **Threshold above everything.** $TP = 0, FP = 0, FN = 5, TN = 5$. Point $(0.0, 0.0)$. Start here, always.
- $t = 0.95$, a **POSITIVE**. $TP = 1, FP = 0$. $FPR = 0/5 = 0.0$, $TPR = 1/5 = 0.2$. Point $(0.0, 0.2)$ — *up*.
- $t = 0.90$, a **POSITIVE**. $TP = 2, FP = 0$. Point $(0.0, 0.4)$ — *up* again. We are climbing the left-hand edge, which is exactly what a good model does.
- $t = 0.85$, a **negative**. $TP = 2, FP = 1$. $FPR = 1/5 = 0.2$, TPR unchanged at 0.4. Point $(0.2, 0.4)$ — *right*. **This is the model's one mistake, and it is the only reason the area will not be 1.**

By hand: the first five steps, one at a time

The sorted labels are **POS, POS, neg, POS, POS, ...**

- **Threshold above everything.** $TP = 0, FP = 0, FN = 5, TN = 5$. Point $(0.0, 0.0)$. Start here, always.
- $t = 0.95$, a **POSITIVE**. $TP = 1, FP = 0$. $FPR = 0/5 = 0.0$, $TPR = 1/5 = 0.2$. Point $(0.0, 0.2)$ — *up*.
- $t = 0.90$, a **POSITIVE**. $TP = 2, FP = 0$. Point $(0.0, 0.4)$ — *up* again. We are climbing the left-hand edge, which is exactly what a good model does.
- $t = 0.85$, a **negative**. $TP = 2, FP = 1$. $FPR = 1/5 = 0.2$, TPR unchanged at 0.4. Point $(0.2, 0.4)$ — *right*. **This is the model's one mistake, and it is the only reason the area will not be 1.**
- $t = 0.75$, a **POSITIVE**. $TP = 3, FP = 1$. Point $(0.2, 0.6)$ — *up*.

By hand: the first five steps, one at a time

The sorted labels are **POS, POS, neg, POS, POS, ...**

- **Threshold above everything.** $TP = 0, FP = 0, FN = 5, TN = 5$. Point $(0.0, 0.0)$. Start here, always.
- $t = 0.95$, a **POSITIVE**. $TP = 1, FP = 0$. $FPR = 0/5 = 0.0$, $TPR = 1/5 = 0.2$. Point $(0.0, 0.2)$ — *up*.
- $t = 0.90$, a **POSITIVE**. $TP = 2, FP = 0$. Point $(0.0, 0.4)$ — *up* again. We are climbing the left-hand edge, which is exactly what a good model does.
- $t = 0.85$, a **negative**. $TP = 2, FP = 1$. $FPR = 1/5 = 0.2$, TPR unchanged at 0.4. Point $(0.2, 0.4)$ — *right*. **This is the model's one mistake, and it is the only reason the area will not be 1.**
- $t = 0.75$, a **POSITIVE**. $TP = 3, FP = 1$. Point $(0.2, 0.6)$ — *up*.
- **Nothing has been fitted, smoothed or estimated.** This is bookkeeping on a sorted list, and you could finish it in the margin of a question paper.

By hand: the whole sweep, eleven points

threshold	TP	FP	FN	TN	FPR	TPR
(nothing)	0	0	5	5	0.0	0.0
0.95 (POS)	1	0	4	5	0.0	0.2
0.90 (POS)	2	0	3	5	0.0	0.4
0.85 (neg)	2	1	3	4	0.2	0.4
0.75 (POS)	3	1	2	4	0.2	0.6
0.60 (POS)	4	1	1	4	0.2	0.8
0.55 (neg)	4	2	1	3	0.4	0.8
0.45 (neg)	4	3	1	2	0.6	0.8
0.35 (POS)	5	3	0	2	0.6	1.0
0.25 (neg)	5	4	0	1	0.8	1.0
0.10 (neg)	5	5	0	0	1.0	1.0

That table *is* the ROC curve

Plot the eleven (FPR, TPR) pairs, join them with horizontal and vertical segments, and the staircase is finished.

Read the pattern in the label column

POS POS *neg* POS POS *neg neg* POS
neg neg becomes
 up up **right** up up **right right** up **right**
right.

The curve depends only on the order of the scores, not on their values — so any monotone rescaling leaves it untouched.

Predict first #2

```
y10 = np.array([1, 1, 0, 1, 1, 0, 0, 1, 0, 0])  
s10 = np.array([0.95, 0.90, 0.85, 0.75, 0.60, 0.55, 0.45, 0.35, 0.25, 0.10])  
fpr, tpr, thr = roc_curve(y10, s10, drop_intermediate=False)
```

How many rows will this print, what will the first threshold be, and will any row disagree with your table?

Predict first #2

```
y10 = np.array([1, 1, 0, 1, 1, 0, 0, 1, 0, 0])  
s10 = np.array([0.95, 0.90, 0.85, 0.75, 0.60, 0.55, 0.45, 0.35, 0.25, 0.10])  
fpr, tpr, thr = roc_curve(y10, s10, drop_intermediate=False)
```

How many rows will this print, what will the first threshold be, and will any row disagree with your table?

```
sklearn roc_curve(y, s, drop_intermediate=False)  
threshold      FPR      TPR  
inf  0.0000  0.0000  
0.95 0.0000  0.2000  
0.90 0.0000  0.4000  
0.85 0.2000  0.4000  
0.75 0.2000  0.6000  
0.60 0.2000  0.8000  
0.55 0.4000  0.8000  
0.45 0.6000  0.8000  
0.35 0.6000  1.0000  
0.25 0.8000  1.0000  
0.10 1.0000  1.0000  
the first threshold is np.inf: the 'predict nothing' corner (0,0).  
every other row matches the hand sweep exactly.
```


Predict first #2

```
y10 = np.array([1, 1, 0, 1, 1, 0, 0, 1, 0, 0])
s10 = np.array([0.95, 0.90, 0.85, 0.75, 0.60, 0.55, 0.45, 0.35, 0.25, 0.10])
fpr, tpr, thr = roc_curve(y10, s10, drop_intermediate=False)
```

How many rows will this print, what will the first threshold be, and will any row disagree with your table?

```
sklearn roc_curve(y, s, drop_intermediate=False)
threshold      FPR      TPR
inf  0.0000  0.0000
0.95  0.0000  0.2000
0.90  0.0000  0.4000
0.85  0.2000  0.4000
0.75  0.2000  0.6000
0.60  0.2000  0.8000
0.55  0.4000  0.8000
0.45  0.6000  0.8000
0.35  0.6000  1.0000
0.25  0.8000  1.0000
0.10  1.0000  1.0000
the first threshold is np.inf: the 'predict nothing' corner (0,0).
every other row matches the hand sweep exactly.
with the default drop_intermediate=True sklearn keeps only the corners:
8 points instead of 11, same curve, same area 0.8000
```

Drawn: your staircase, and how to read one



Left: your staircase with its threshold labels, and `roc_curve`'s crosses on top. **The largest disagreement is 0.00.** Right: the five strips whose areas we are about to add up.

Four things to look for. (0,1) is the perfect classifier, so closer is better; the diagonal is a score with no information; *steepness at the left* is what matters in practice; and a dip below the diagonal means the score is anti-correlated over that band.

By hand: the area, as five trapezoids

The curve is piecewise linear, so the trapezoidal rule is *exact*. Cut the region into vertical strips, one per horizontal move; each strip is 0.2 wide and as tall as the TPR during that move.

from	to	width	TPR	area
0.0	0.2	0.2	0.4	0.08
0.2	0.4	0.2	0.8	0.16
0.4	0.6	0.2	0.8	0.16
0.6	0.8	0.2	1.0	0.20
0.8	1.0	0.2	1.0	0.20
total				0.80

```

from    to    width  height (TPR)  strip
0.00    0.20    0.20    0.40          0.0800
0.20    0.40    0.20    0.80          0.1600
0.40    0.60    0.20    0.80          0.1600
0.60    0.80    0.20    1.00          0.2000
0.80    1.00    0.20    1.00          0.2000
trapezoidal area
0.8000
sklearn auc(fpr, tpr)
0.8000
sklearn roc_auc_score(y, s)
0.8000

```

By hand: the area, as five trapezoids

The curve is piecewise linear, so the trapezoidal rule is *exact*. Cut the region into vertical strips, one per horizontal move; each strip is 0.2 wide and as tall as the TPR during that move.

from	to	width	TPR	area
0.0	0.2	0.2	0.4	0.08
0.2	0.4	0.2	0.8	0.16
0.4	0.6	0.2	0.8	0.16
0.6	0.8	0.2	1.0	0.20
0.8	1.0	0.2	1.0	0.20
total				0.80

```

from    to    width  height (TPR)  strip
0.00    0.20    0.20      0.40      0.0800
0.20    0.40    0.20      0.80      0.1600
0.40    0.60    0.20      0.80      0.1600
0.60    0.80    0.20      1.00      0.2000
0.80    1.00    0.20      1.00      0.2000
trapezoidal area
0.8000
sklearn auc(fpr, tpr)
0.8000
sklearn roc_auc_score(y, s)
0.8000

```

`auc(fpr, tpr)` integrates the curve you were given; `roc_auc_score(y, s)` builds the curve itself and integrates that. **Same 0.8000, and the hand sum $0.08 + 0.16 + 0.16 + 0.20 + 0.20$ is the same arithmetic.**

By hand: the same number, without a drawing

There are $P \times N = 5 \times 5 = 25$ ways to pick one positive and one negative. Count the pairs the model ordered *correctly*.

negative scored	positives above it
0.85	2
0.55	4
0.45	4
0.25	5
0.10	5
total	20 of 25

Only 0.95 and 0.90 beat the negative at 0.85; every positive beats the two lowest negatives.

$$20/25 = 0.80$$

The same number the trapezoids gave. **It is not a coincidence and it is not an approximation.** It is an identity:

$$\text{AUC} = \Pr(s(X^+) > s(X^-)) + \frac{1}{2} \Pr(s(X^+) = s(X^-))$$

with X^+, X^- drawn independently from the positives and the negatives. **Ties count a half** — which is exactly what the diagonal segments of a staircase contribute.

The pair count, and what it turns out to be

```
for each negative: how many positives outrank it?
negative scored 0.85 -> 2 of the 5 positives above it
negative scored 0.55 -> 4 of the 5 positives above it
negative scored 0.45 -> 4 of the 5 positives above it
negative scored 0.25 -> 5 of the 5 positives above it
negative scored 0.10 -> 5 of the 5 positives above it
concordant pairs 20   ties 0   total pairs 25
P(random positive outranks random negative) = 20/25 = 0.8000
roc_auc_score                               = 0.8000
```

The pair count, and what it turns out to be

```
for each negative: how many positives outrank it?
negative scored 0.85 -> 2 of the 5 positives above it
negative scored 0.55 -> 4 of the 5 positives above it
negative scored 0.45 -> 4 of the 5 positives above it
negative scored 0.25 -> 5 of the 5 positives above it
negative scored 0.10 -> 5 of the 5 positives above it
concordant pairs 20   ties 0   total pairs 25
P(random positive outranks random negative) = 20/25 = 0.8000
roc_auc_score                               = 0.8000
```

```
U = stats.mannwhitneyu(pos, neg, alternative="greater").statistic
print(U, U / (len(pos) * len(neg)), roc_auc_score(y10, s10))
```

```
Mann-Whitney U statistic  20.0
U / (n_pos * n_neg)      0.8000
roc_auc_score            0.8000
AUC and the Mann-Whitney U statistic are the same quantity, rescaled.
```

The pair count, and what it turns out to be


```
for each negative: how many positives outrank it?
negative scored 0.85 -> 2 of the 5 positives above it
negative scored 0.55 -> 4 of the 5 positives above it
negative scored 0.45 -> 4 of the 5 positives above it
negative scored 0.25 -> 5 of the 5 positives above it
negative scored 0.10 -> 5 of the 5 positives above it
concordant pairs 20   ties 0   total pairs 25
P(random positive outranks random negative) = 20/25 = 0.8000
roc_auc_score                               = 0.8000
```

```
U = stats.mannwhitneyu(pos, neg, alternative="greater").statistic
print(U, U / (len(pos) * len(neg)), roc_auc_score(y10, s10))
```

```
Mann-Whitney U statistic  20.0
U / (n_pos * n_neg)      0.8000
roc_auc_score            0.8000
AUC and the Mann-Whitney U statistic are the same quantity, rescaled.
```

So the AUC of a set of scores and the rank-sum test of “do the positives score higher than the negatives?” are **the same computation**, and Hanley and McNeil’s 1982 paper is where that was first written down.

Drawn: twenty-five pairs, and a million



Left: all 25 positive-negative pairs from the ten hand rows; a mark means the model ordered that pair correctly. **Twenty of twenty-five, AUC 0.80** — you can literally count the picture. Right: the sampled estimate on `breast_cancer` converging on the analytic 0.966121, inside the ± 2 standard-error band of a Bernoulli proportion — 100 pairs give 0.980000, ten million give 0.966230, and counting all **6848** pairs exactly reproduces `roc_auc_score` to a difference of **0.00e+00**. **AUC is not like that probability; it is that probability.**

Predict first #3

```
rng = np.random.default_rng(0) # the one seed for this lecture
yr, sr = rng.integers(0, 2, 2000), rng.random(2000)
print(roc_auc_score(yr, sr))
```

Two thousand coin-flip labels, two thousand uniform random scores. There is no model here at all.

What AUC comes out — and, harder, how far from it can one run stray?

Predict first #3

```
rng = np.random.default_rng(0) # the one seed for this lecture
yr, sr = rng.integers(0, 2, 2000), rng.random(2000)
print(roc_auc_score(yr, sr))
```

Two thousand coin-flip labels, two thousand uniform random scores. There is no model here at all.

What AUC comes out — and, harder, how far from it can one run stray?

```
2000 coin-flip labels, 2000 uniform random scores
roc_auc_score 0.5023
accuracy at threshold 0.5 0.5090
200 repetitions: mean 0.4994 sd 0.0136 min 0.4645 max 0.5342
the diagonal is not a curve you draw, it is the answer you get
when the score carries no information about the label.
```

Predict first #3

```
rng = np.random.default_rng(0) # the one seed for this lecture
yr, sr = rng.integers(0, 2, 2000), rng.random(2000)
print(roc_auc_score(yr, sr))
```

Two thousand coin-flip labels, two thousand uniform random scores. There is no model here at all.

What AUC comes out — and, harder, how far from it can one run stray?


```
2000 coin-flip labels, 2000 uniform random scores
roc_auc_score 0.5023
accuracy at threshold 0.5 0.5090
200 repetitions: mean 0.4994 sd 0.0136 min 0.4645 max 0.5342
the diagonal is not a curve you draw, it is the answer you get
when the score carries no information about the label.
```

That spread is the useful part. **On 2000 rows an AUC of 0.53 is not evidence of anything** — the noise runs to 0.5342 on its own.

An AUC below 0.5 is a different animal

```
print(roc_auc_score(yb_te, sb),  
      roc_auc_score(yb_te, -sb),  
      roc_auc_score(1 - yb_te, sb))
```


```
the fitted breast_cancer model      AUC  
0.9661  
the same model with its scores negated  AUC  
0.0339  
the same model with the labels swapped  AUC  
0.0339  
1 - AUC  
0.0339  
An AUC of 0.0339 is not a bad model. It is this model  
wired backwards:  
flip the sign of the score and you get 0.9661 back.
```



An AUC below 0.5 is a different animal

```
print(roc_auc_score(yb_te, sb),  
      roc_auc_score(yb_te, -sb),  
      roc_auc_score(1 - yb_te, sb))
```

```
the fitted breast_cancer model      AUC  
0.9661  
the same model with its scores negated  AUC  
0.0339  
the same model with the labels swapped  AUC  
0.0339  
1 - AUC  
0.0339  
An AUC of 0.0339 is not a bad model. It is this model  
wired backwards:  
flip the sign of the score and you get 0.9661 back.
```



$AUC(-s) = 1 - AUC(s)$, exactly

Negating the score reverses every pairwise comparison, so 0.5 — not 0 — is the floor for a model you are allowed to relabel. **An AUC well below 0.5 almost always means you passed the wrong column of predict_proba, or defined the positive class one way and scored it the other.** In the figure the "inverted" curve at 0.192 becomes 0.808 the moment you flip a sign.

The two invariances

Threshold. Look again at the sweep: the threshold is a *row label*, never an input. Accuracy moved $0.7135 \rightarrow 0.9181$ across seven thresholds while AUC sat at 0.9661 throughout.

Class balance. Keep every positive, throw negatives away at random. TPR lives inside the positives and FPR inside the negatives.

same model, same scores; only the negatives are subsampled

negatives kept	prevalence	ROC AUC	avg precision
5940	0.0100	0.9200	0.1702
2000	0.0291	0.9197	0.3380
500	0.1071	0.9172	0.6699
120	0.3333	0.9243	0.8839
60	0.5000	0.9158	0.9202

ROC AUC is invariant to class balance. Average precision is not:
it rose from 0.1702 to 0.9202 -- a factor of 5.4 -- without the model
changing at all.

The two invariances

Threshold. Look again at the sweep: the threshold is a *row label*, never an input. Accuracy moved $0.7135 \rightarrow 0.9181$ across seven thresholds while AUC sat at 0.9661 throughout.

Class balance. Keep every positive, throw negatives away at random. TPR lives inside the positives and FPR inside the negatives.

same model, same scores; only the negatives are subsampled

negatives kept	prevalence	ROC AUC	avg precision
5940	0.0100	0.9200	0.1702
2000	0.0291	0.9197	0.3380
500	0.1071	0.9172	0.6699
120	0.3333	0.9243	0.8839
60	0.5000	0.9158	0.9202

ROC AUC is invariant to class balance. Average precision is not:
it rose from 0.1702 to 0.9202 -- a factor of 5.4 -- without the model
changing at all.

ROC AUC moved 0.0092 over a **fiftyfold** change in prevalence — sampling noise. Average precision moved by a factor of **5.4**. **One fixed set of scores. Nothing about the model changed.** That single table is the whole of the next section.

Where ROC flatters: the 99:1 problem again

```
the 99:1 problem, logistic regression, 6000 held-out rows
prevalence          0.0100 (60 positives, 5940 negatives)
ROC AUC             0.9200 (baseline 0.5000)
average precision    0.1702 (baseline = prevalence = 0.0100)
```

recall	precision	threshold	TP	FP
0.10	0.2400	0.3634	6	19
0.30	0.2647	0.2004	18	50
0.50	0.1648	0.0719	30	152
0.70	0.0955	0.0275	42	398
0.90	0.0439	0.0069	54	1177
1.00	0.0100	0.0000	60	5940

ROC reads: at TPR 0.9000 the FPR is only 0.1946 -- encouraging.

PR reads: 0.1946 x 5940 = 1156 false alarms for 54 true ones,
precision 0.0446. One alert in 22 is real.

Where ROC flatters: the 99:1 problem again

```
the 99:1 problem, logistic regression, 6000 held-out rows
prevalence          0.0100 (60 positives, 5940 negatives)
ROC AUC             0.9200 (baseline 0.5000)
average precision    0.1702 (baseline = prevalence = 0.0100)
```

recall	precision	threshold	TP	FP
0.10	0.2400	0.3634	6	19
0.30	0.2647	0.2004	18	50
0.50	0.1648	0.0719	30	152
0.70	0.0955	0.0275	42	398
0.90	0.0439	0.0069	54	1177
1.00	0.0100	0.0000	60	5940



ROC reads: at TPR 0.9000 the FPR is only 0.1946 -- encouraging.

PR reads: 0.1946 x 5940 = 1156 false alarms for 54 true ones,
precision 0.0446. One alert in 22 is real.

Both readings are true. Only one is the client's question

AUC 0.9200 *is* a strong model in the precise sense AUC measures: pick a random positive and a random negative and it orders them correctly 92% of the time. But to catch 90% of the positives you must accept 1156 **false alarms alongside 54 real ones**. The FPR of 0.1946 that looked small on the ROC axis is 19.46% of 5940 negatives — and 5940 is a large number.


Drawn: one model, two pictures, two verdicts



	axes	baseline of a useless model	moves with prevalence?
ROC	TPR against FPR	the diagonal, AUC = 0.5	no
PR	precision against recall	the horizontal line at the prevalence	yes

FPR divides by the number of negatives, which is huge when positives are rare, so a large absolute number of false alarms shows up as a small FPR. Precision divides by TP + FP instead, and when FP swamps TP it falls off a cliff.

Drawn: five prevalences, one set of scores



Report ROC AUC when...

the classes are comparably common; or you want a summary that will not shift because next month's data has a different mix; or a false positive and a false negative cost about the same.

Report average precision when...

positives are rare and somebody is going to act on the alerts. Fraud review, medical triage, moderation — a human has finite capacity, and precision is exactly the fraction of their time that is not wasted. **Always quote the prevalence beside it.**

And the invariance runs both ways: an AUC measured on a balanced subsample *transfers* to the imbalanced population; **an average precision does not**. Final evaluation goes on data with the production prevalence.

0.5 is an assumption — and Youden's J by hand

predict thresholds at 0.5. That is optimal under exactly one condition: **a false positive and a false negative cost the same.**

threshold	FPR	TPR	J
∞	0.0	0.0	0.0
0.95	0.0	0.2	0.2
0.90	0.0	0.4	0.4
0.85	0.2	0.4	0.2
0.75	0.2	0.6	0.4
0.60	0.2	0.8	0.6
0.55	0.4	0.8	0.4
0.45	0.6	0.8	0.2
0.35	0.6	1.0	0.4
0.25	0.8	1.0	0.2
0.10	1.0	1.0	0.0

$$J = \text{TPR} - \text{FPR} = \text{sens.} + \text{spec.} - 1$$

Maximising it picks the point on the curve **furthest above the diagonal, measured vertically**. On the ten hand rows the maximum is $J = 0.6$ at threshold **0.60**, giving $(\text{FPR}, \text{TPR}) = (0.2, 0.8)$.

And here is what it quietly assumes

J weights a unit of sensitivity and a unit of specificity **equally**, so it is the equal-cost answer wearing different clothes. A reasonable default when nobody will tell you the costs, and **the wrong answer the moment somebody does.**

Youden on real data

On the three-feature breast_cancer model, with *malignant* as the positive class this time.

```
the same rule on breast_cancer, malignant as the positive class:
Youden threshold 0.3612  J 0.8160  FPR 0.1215  TPR 0.9375
at t = 0.5000: recall 0.8594  precision 0.8871  accuracy 0.9064
at Youden      : recall 0.9375  precision 0.8219  accuracy 0.9006
```

Youden on real data

On the three-feature breast_cancer model, with *malignant* as the positive class this time.

```
the same rule on breast_cancer, malignant as the positive class:
Youden threshold 0.3612   J 0.8160   FPR 0.1215   TPR 0.9375
at t = 0.5000: recall 0.8594   precision 0.8871   accuracy 0.9064
at Youden      : recall 0.9375   precision 0.8219   accuracy 0.9006
```

What the move bought

Recall on the class you must not miss went
0.8594 → **0.9375**: eight points.

What it cost

Precision 0.8871 → 0.8219 and accuracy
0.9064 → 0.9006. **Whether that is a good trade
is not a statistical question.**

Predict first #4

Give a miss a cost C_{FN} and a false alarm a cost C_{FP} , and minimise $\text{cost}(t) = C_{FN}FN(t) + C_{FP}FP(t)$. Only the *ratio* matters, so fix $C_{FP} = 1$. Youden picked 0.60 on the ten hand rows.

Which threshold is cheapest at 1:1, and which at 10:1?

Predict first #4

Give a miss a cost C_{FN} and a false alarm a cost C_{FP} , and minimise $\text{cost}(t) = C_{FN}FN(t) + C_{FP}FP(t)$. Only the *ratio* matters, so fix $C_{FP} = 1$. Youden picked 0.60 on the ten hand rows.

Which threshold is cheapest at 1:1, and which at 10:1?

```
ten hand rows, cost of a miss C_FN, cost of a false alarm C_FP = 1
threshold TP FP FN TN cost(1:1) cost(10:1)
0.95 1 0 4 5 4 40
0.90 2 0 3 5 3 30
0.85 2 1 3 4 4 31
0.75 3 1 2 4 3 21
0.60 4 1 1 4 2 11
0.55 4 2 1 3 3 12
0.45 4 3 1 2 4 13
0.35 5 3 0 2 3 3
0.25 5 4 0 1 4 4
0.10 5 5 0 0 5 5
cheapest at 1:1 threshold 0.60 cost 2 (TPR 0.80, FPR 0.20)
cheapest at 10:1 threshold 0.35 cost 3 (TPR 1.00, FPR 0.60)
Youden picked 0.60. Youden is the 1:1 answer in disguise, on a
balanced set. Change the price of a miss and the answer changes.
```

Predict first #4

Give a miss a cost C_{FN} and a false alarm a cost C_{FP} , and minimise $\text{cost}(t) = C_{FN}FN(t) + C_{FP}FP(t)$. Only the *ratio* matters, so fix $C_{FP} = 1$. Youden picked 0.60 on the ten hand rows.

Which threshold is cheapest at 1:1, and which at 10:1?

```
ten hand rows, cost of a miss C_FN, cost of a false alarm C_FP = 1
threshold TP FP FN TN    cost(1:1)    cost(10:1)
0.95      1  0  4  5         4         40
0.90      2  0  3  5         3         30
0.85      2  1  3  4         4         31
0.75      3  1  2  4         3         21
0.60      4  1  1  4         2         11
0.55      4  2  1  3         3         12
0.45      4  3  1  2         4         13
0.35      5  3  0  2         3          3
0.25      5  4  0  1         4          4
0.10      5  5  0  0         5          5
cheapest at 1:1 threshold 0.60 cost 2 (TPR 0.80, FPR 0.20)
cheapest at 10:1 threshold 0.35 cost 3 (TPR 1.00, FPR 0.60)
Youden picked 0.60. Youden is the 1:1 answer in disguise, on a
balanced set. Change the price of a miss and the answer changes.
```

Making a miss ten times as expensive moved the threshold from 0.60 to 0.35, took recall from 0.80 to 1.00, and tripled the false alarm rate. Read down the 10:1 column and watch the minimum walk left as the price of a miss rises.

The same costing on breast_cancer

```
grid = np.unique(sm) # sm = 1 - sb, malignant score
for ratio in [1, 2, 5, 10, 20, 50]:
    costs = [ratio * FN(t) + FP(t) for t in grid]
    t = grid[int(np.argmin(costs))]
```

```
the same costing on breast_cancer; malignant is now the positive
malignant 64   benign 107   ROC AUC 0.9661
```

C_FN:C_FP	threshold	cost	FN	FP	recall	precision
1:1	0.6855	14	13	1	0.7969	0.9808
2:1	0.3612	21	4	13	0.9375	0.8219
5:1	0.3612	33	4	13	0.9375	0.8219
10:1	0.0670	42	0	42	1.0000	0.6038
20:1	0.0670	42	0	42	1.0000	0.6038
50:1	0.0670	42	0	42	1.0000	0.6038

Raising the price of a miss from 1 to 10 moved the threshold from 0.6855 to 0.0670, recall from 0.7969 to 1.0000 and precision from 0.9808 to 0.6038.

Nothing about the model changed. Only the price list.

The same costing on breast_cancer

```
grid = np.unique(sm) # sm = 1 - sb, malignant score
for ratio in [1, 2, 5, 10, 20, 50]:
    costs = [ratio * FN(t) + FP(t) for t in grid]
    t = grid[int(np.argmin(costs))]
```

```
the same costing on breast_cancer; malignant is now the positive
malignant 64   benign 107   ROC AUC 0.9661
```

C_FN:C_FP	threshold	cost	FN	FP	recall	precision
1:1	0.6855	14	13	1	0.7969	0.9808
2:1	0.3612	21	4	13	0.9375	0.8219
5:1	0.3612	33	4	13	0.9375	0.8219
10:1	0.0670	42	0	42	1.0000	0.6038
20:1	0.0670	42	0	42	1.0000	0.6038
50:1	0.0670	42	0	42	1.0000	0.6038


Raising the price of a miss from 1 to 10 moved the threshold from 0.6855 to 0.0670, recall from 0.7969 to 1.0000 and precision from 0.9808 to 0.6038.

Nothing about the model changed. Only the price list.

13 missed cancers and one false alarm, or none and forty-two

Same model, same scores, same day — a completely different clinical instrument. Report the threshold *with* the model, always, and say where it came from. Note too that the cost curve is a step function: 2:1 and 5:1 share a threshold, and once $FN = 0$ you cannot buy fewer than zero misses.

Drawn: the cost curve, three operating points



Left: total cost against threshold on the ten hand rows at 1:1 and 10:1; the minima are at 0.60 and 0.35. Right: three operating points on one breast_cancer ROC curve. **The curve is the model; the points are decisions.**

The threshold is a parameter, and parameters leak

Maximising J , or minimising cost, **on the test set** is fitting to the test set — the same error as choosing features on it, just smaller. Choose the threshold on a validation split or inside the folds; report once, at the threshold you already chose. TunedThresholdClassifierCV does exactly this inside a Pipeline.

By hand: a hundred rows in three classes

For class i : TP_i is the diagonal entry, FP_i is the rest of *column* i , FN_i is the rest of *row* i .

	pred A	pred B	pred C	total
actual A	54	4	2	60
actual B	6	21	3	30
actual C	4	2	4	10
col. total	64	27	9	100

The off-diagonal pattern is the informative part

A model that confuses 3 with 8 has a different problem from one that confuses 3 with 5. **No single average will tell you which.** Accuracy is the diagonal over the total: $(54 + 21 + 4)/100 = \mathbf{0.79}$.

class	TP	FP	FN	support	precision	recall	F_1
A	54	10	6	60	$54/64 = 0.843750$	0.900000	0.870968
B	21	6	9	30	$21/27 = 0.777778$	0.700000	0.736842
C	4	5	6	10	$4/9 = 0.444444$	0.400000	0.421053

By hand: micro, macro and weighted

Micro pools the counts before dividing:

$\sum TP = 79$, $\sum FP = 21$, $\sum FN = 21$, so
micro precision = recall = $F_1 = 79/100 = 0.79$.

$$\text{macro } p = \frac{1}{3}(0.843750 + 0.777778 + 0.444444) = 0.688657$$

$$\text{macro } r = \frac{1}{3}(0.9 + 0.7 + 0.4) = 0.666667$$

$$\begin{aligned}\text{macro } F_1 &= \frac{1}{3}(0.870968 + 0.736842 + 0.421053) \\ &= 0.676287\end{aligned}$$

Weighted uses the supports 60, 30, 10:

$$\begin{aligned}\text{wtd } F_1 &= \frac{60(0.870968) + 30(0.736842) + 10(0.421053)}{100} \\ &= 0.785739\end{aligned}$$

Two facts that fall out

(1) In single-label multiclass, micro precision, micro recall and micro F_1 all equal *accuracy* — always, because every row contributes exactly one TP or one FP and one FN.

(2) **Weighted recall equals accuracy**, always, because $\sum_i n_i(TP_i/n_i) = \sum_i TP_i$.

And the gap is the warning

Macro F_1 is **0.1095** *below* weighted F_1 , because class C is a tenth of the rows, a third of the macro vote, and the model is bad at it.

The same hundred rows, in code

```
print(classification_report(y_true3, y_pred3, digits=6,  
                             target_names=["A", "B", "C"]))
```

	precision	recall	f1-score	support
A	0.843750	0.900000	0.870968	60
B	0.777778	0.700000	0.736842	30
C	0.444444	0.400000	0.421053	10
accuracy			0.790000	100
macro avg	0.688657	0.666667	0.676287	100
weighted avg	0.784028	0.790000	0.785739	100

The same hundred rows, in code

```
print(classification_report(y_true3, y_pred3, digits=6,
                           target_names=["A", "B", "C"]))
```

	precision	recall	f1-score	support
A	0.843750	0.900000	0.870968	60
B	0.777778	0.700000	0.736842	30
C	0.444444	0.400000	0.421053	10
accuracy			0.790000	100
macro avg	0.688657	0.666667	0.676287	100
weighted avg	0.784028	0.790000	0.785739	100

Micro asks “how often is the model right?” **Macro** asks “how well does it do on a typical *class*?” **Weighted** asks “how well does it do on a typical *row*?” Choose by deciding who you are answerable to: a per-class regulator wants macro, a throughput target wants micro.

On real data, where the rare classes hide

load_digits with two classes made rare on purpose — only 20 examples each of the digits 3 and 8.

```
imbalanced digits: 1480 rows, class counts [178 182 177 20 181 182 181 179 20 180]
accuracy 0.9662
  micro      precision 0.9662  recall 0.9662  F1 0.9662
  macro      precision 0.9730  recall 0.9131  F1 0.9339
  weighted   precision 0.9669  recall 0.9662  F1 0.9652
per-class recall: [1.      0.9818 0.9811 0.8333 0.963  0.9273 0.9815 0.9815 0.5    0.9815]
classes 3 and 8 have 6 test rows each and recall 0.8333 and 0.5000.
micro hides them completely; macro is the only average that shows them.
```

On real data, where the rare classes hide

load_digits with two classes made rare on purpose — only 20 examples each of the digits 3 and 8.

```
imbalanced digits: 1480 rows, class counts [178 182 177 20 181 182 181 179 20 180]
accuracy 0.9662
  micro      precision 0.9662  recall 0.9662  F1 0.9662
  macro      precision 0.9730  recall 0.9131  F1 0.9339
  weighted   precision 0.9669  recall 0.9662  F1 0.9652
per-class recall: [1.      0.9818 0.9811 0.8333 0.963  0.9273 0.9815 0.9815 0.5    0.9815]
classes 3 and 8 have 6 test rows each and recall 0.8333 and 0.5000.
micro hides them completely; macro is the only average that shows them.
```

Micro and weighted both read 0.97. Macro reads 0.93, and macro *recall* reads **0.9131** — because the model finds half of the eights. **If the eights are the cases you care about, only the macro row and the per-class array told you anything.**

Predict first #5

ROC needs a positive class, so multiclass ROC means k one-vs-rest curves. A three-class problem at roughly 80/15/5; the model **never once predicts class 2**, so its recall and F_1 are 0.0000.

What is class 2's one-vs-rest AUC?

Predict first #5

ROC needs a positive class, so multiclass ROC means k one-vs-rest curves. A three-class problem at roughly 80/15/5; the model **never once predicts class 2**, so its recall and F_1 are 0.0000.

What is class 2's one-vs-rest AUC?

```
three classes, [709 140 51] in the test set
accuracy 0.8600

class  support  one-vs-rest AUC  recall    F1
0         709         0.8920   0.9633  0.9236
1         140         0.9090   0.6500  0.6766
2          51         0.7793   0.0000  0.0000

macro   OvR AUC 0.8601   F1 0.5334
weighted OvR AUC 0.8883   F1 0.8328
micro                                F1 0.8600 (= accuracy 0.8600)
macro Ov0 AUC 0.8151

macro F1 0.5334 against weighted F1 0.8328: a gap of 0.2994.
Report the one that matches who you will disappoint.
```

Predict first #5

ROC needs a positive class, so multiclass ROC means k one-vs-rest curves. A three-class problem at roughly 80/15/5; the model **never once predicts class 2**, so its recall and F_1 are 0.0000.

What is class 2's one-vs-rest AUC?

```
three classes, [709 140 51] in the test set
accuracy 0.8600

class  support  one-vs-rest AUC  recall    F1
0         709         0.8920   0.9633  0.9236
1         140         0.9090   0.6500  0.6766
2          51         0.7793   0.0000  0.0000


macro   OvR AUC 0.8601   F1 0.5334
weighted OvR AUC 0.8883   F1 0.8328
micro                                F1 0.8600 (= accuracy 0.8600)
macro Ov0 AUC 0.8151

macro F1 0.5334 against weighted F1 0.8328: a gap of 0.2994.
Report the one that matches who you will disappoint.
```

Recall 0.0000, F_1 0.0000, **AUC 0.7793**

There is real information in the scores for class 2; `argmax` simply never lets it win, because class 0 is sixteen times as common and the prior buries it. **AUC sees the ranking; F_1 sees only the decision.** Report weighted $F_1 = 0.8328$ and stop, and you ship a model that cannot detect one of its three classes.

Drawn: three curves, and the averaging choice



Left: three one-vs-rest ROC curves and their macro average. Right: per-class F_1 against class share, with macro and weighted as horizontal lines — **they differ by 0.2994**. There is also `multi_class="ovo"`, averaging all $\binom{k}{2}$ pairwise AUCs (0.8151 here), which is the more natural choice when the class sizes are wildly different. **Report the average that matches who you will disappoint.**

Seven classifiers, twenty folds, one metric

```
cv = RepeatedStratifiedKFold(n_splits=5, n_repeats=4, random_state=0)
res = {n: cross_val_score(m, Xb, yb, cv=cv, scoring="roc_auc")
      for n, m in models.items()}
```

breast_cancer, all 30 features, 5-fold x 4 repeats = 20 folds

model	mean AUC	sd	2 s.e.	worst	best
SVM (RBF)	0.9950	0.0058	0.0026	0.9771	1.0000
Logistic reg.	0.9948	0.0056	0.0025	0.9803	1.0000
MLP (100)	0.9944	0.0060	0.0027	0.9800	1.0000
Random forest	0.9905	0.0090	0.0040	0.9705	1.0000
Gaussian NB	0.9882	0.0069	0.0031	0.9705	0.9993
KNN (k=5)	0.9849	0.0107	0.0048	0.9633	1.0000
Decision tree	0.9227	0.0205	0.0092	0.8809	0.9554

Seven classifiers, twenty folds, one metric

```
cv = RepeatedStratifiedKFold(n_splits=5, n_repeats=4, random_state=0)
res = {n: cross_val_score(m, Xb, yb, cv=cv, scoring="roc_auc")
       for n, m in models.items()}
```

```
breast_cancer, all 30 features, 5-fold x 4 repeats = 20 folds
model      mean AUC      sd    2 s.e.    worst    best
SVM (RBF)   0.9950  0.0058  0.0026  0.9771  1.0000
Logistic reg. 0.9948  0.0056  0.0025  0.9803  1.0000
MLP (100)   0.9944  0.0060  0.0027  0.9800  1.0000
Random forest 0.9905  0.0090  0.0040  0.9705  1.0000
Gaussian NB  0.9882  0.0069  0.0031  0.9705  0.9993
KNN (k=5)   0.9849  0.0107  0.0048  0.9633  1.0000
Decision tree 0.9227  0.0205  0.0092  0.8809  0.9554
```

Three notes on the machinery. `RepeatedStratifiedKFold` gives twenty estimates rather than five, which is what makes the standard error worth quoting. **Every scaler sits inside a Pipeline**, so it is refitted in each fold. And `scoring="roc_auc"` was chosen *before* the run, not after looking at the results.

Accuracy
○○○

Confusion
○○○○○

Scores
○○○○

ROC
○○○○○○

AUC
○○○○○○○○

ROC/PR
○○○

Threshold
○○○○○

Multiclass
○○○○○○

Selection
○●○○○○

Summary
○○

Predict first #6

The table above is easy to read as a ranking: SVM (RBF) 0.9950 beats logistic regression 0.9948.

Is that a result? How many of the twenty folds did the SVM actually win?

Predict first #6

The table above is easy to read as a ranking: SVM (RBF) 0.9950 beats logistic regression 0.9948.

Is that a result? How many of the twenty folds did the SVM actually win?

```
best    SVM (RBF)          0.9950
second  Logistic reg.      0.9948
gap 0.0003; fold-to-fold sd of the gap 0.0022
paired t = 0.598, p = 0.5569 over 20 folds
SVM (RBF) beat Logistic reg. on 7 of the 20 folds
```

Predict first #6

The table above is easy to read as a ranking: SVM (RBF) 0.9950 beats logistic regression 0.9948.

Is that a result? How many of the twenty folds did the SVM actually win?

```
best    SVM (RBF)          0.9950
second  Logistic reg.      0.9948
gap 0.0003; fold-to-fold sd of the gap 0.0022
paired t = 0.598, p = 0.5569 over 20 folds
SVM (RBF) beat Logistic reg. on 7 of the 20 folds
```

0.9950 against 0.9948 is not a result

The lead is 0.0003. The fold-to-fold standard deviation of that lead is 0.0022, **seven times larger**, and the nominal winner *lost thirteen of the twenty folds*. A paired t -test gives $p = 0.5569$. **When the gap is smaller than the noise, choose on something else:** training time, interpretability, behaviour on the tail, whether you can explain it to a clinician.

The ranking flips from fold to fold

fold-by-fold, the top four models, first eight folds

fold	SVM (RBF)	Logistic reg.	MLP (100)	Random forest	winner
0	0.9869	0.9846	0.9840	0.9813	SVM (RBF)
1	0.9987	0.9990	0.9954	0.9990	Logistic reg.
2	0.9967	0.9980	0.9974	0.9866	Logistic reg.
3	1.0000	1.0000	1.0000	0.9964	SVM (RBF)
4	0.9963	0.9956	0.9946	0.9980	Random forest
5	0.9925	0.9905	0.9889	0.9880	SVM (RBF)
6	0.9977	0.9987	0.9984	0.9838	Logistic reg.
7	0.9983	0.9950	0.9997	0.9960	MLP (100)


The ranking flips from fold to fold

fold-by-fold, the top four models, first eight folds

fold	SVM (RBF)	Logistic reg.	MLP (100)	Random forest	winner
0	0.9869	0.9846	0.9840	0.9813	SVM (RBF)
1	0.9987	0.9990	0.9954	0.9990	Logistic reg.
2	0.9967	0.9980	0.9974	0.9866	Logistic reg.
3	1.0000	1.0000	1.0000	0.9964	SVM (RBF)
4	0.9963	0.9956	0.9946	0.9980	Random forest
5	0.9925	0.9905	0.9889	0.9880	SVM (RBF)
6	0.9977	0.9987	0.9984	0.9838	Logistic reg.
7	0.9983	0.9950	0.9997	0.9960	MLP (100)

Four different winners in eight folds. Had you run a single 80/20 split you would have crowned one of them, and which one would have been decided by `random_state`.

Drawn: error bars, and every fold



fold-to-fold sd of the leader: 0.0058

models within one sd of the leader: Logistic reg., MLP (100), Random forest, SVM (RBF)

A gap smaller than the fold-to-fold spread is not a difference.

Left: mean cross-validated AUC with ± 2 standard-error bars; the ticks are the twenty individual folds and the band is one fold-to-fold standard deviation of the leader. Right: the top four models fold by fold — **the lines cross constantly, and six of the seven models won at least one fold.**

And the metric decides the ranking

The same seven models, the same folds, scored by accuracy as well as AUC.

the same seven models scored by accuracy and by AUC, same folds

model	accuracy	AUC	acc rank	AUC rank
SVM (RBF)	0.9771	0.9957	3	1
Logistic reg.	0.9789	0.9955	1	2
MLP (100)	0.9789	0.9943	2	3
Random forest	0.9578	0.9923	5	4
Gaussian NB	0.9385	0.9885	6	5
KNN (k=5)	0.9649	0.9806	4	6
Decision tree	0.9262	0.9210	7	7

models whose rank differs between the two metrics: 6 of 7

Pick the metric before you pick the model, and say why.

And the metric decides the ranking

The same seven models, the same folds, scored by accuracy as well as AUC.

the same seven models scored by accuracy and by AUC, same folds

model	accuracy	AUC	acc rank	AUC rank
SVM (RBF)	0.9771	0.9957	3	1
Logistic reg.	0.9789	0.9955	1	2
MLP (100)	0.9789	0.9943	2	3
Random forest	0.9578	0.9923	5	4
Gaussian NB	0.9385	0.9885	6	5
KNN (k=5)	0.9649	0.9806	4	6
Decision tree	0.9262	0.9210	7	7

models whose rank differs between the two metrics: 6 of 7

Pick the metric before you pick the model, and say why.

Six of seven models change rank

k-NN is fourth by accuracy and sixth by AUC; the random forest is fifth by accuracy and fourth by AUC.

There is no metric-free notion of “the best model” — which closes the loop with the first section, where accuracy and AUC ranked two models in opposite orders.

The checklist for your classifier comparison

1. State the **positive class** and the **prevalence**.
2. State the metric *before* you run anything, and say why that metric.
3. Put every preprocessing step inside a Pipeline.
4. Use **repeated stratified cross-validation**, not one split.
5. Report mean, standard deviation and the number of folds — never a mean alone.
6. Include a baseline: `DummyClassifier`, and one simple model.
7. Plot the ROC curves together, and the PR curves too if positives are rare.
8. Choose the threshold from a stated cost or capacity constraint, **on validation data**.
9. Say plainly when two models are indistinguishable, and give the *non-statistical* reason you picked one.

Nine lines. **Every one of them is a mark on the assignment, and eight of them are one sentence in your report.**

Summary — eleven lines

1. Accuracy is a weighted summary and the weights are the class frequencies. At 1% prevalence a constant predictor scored **0.9900** and a useful logistic regression scored **0.9897** — **accuracy ranked them backwards** while AUC ranked them 0.5000 against 0.9200.
2. Four counts, six metrics. On twenty patients with TP = 6, FP = 2, FN = 4, TN = 8: accuracy 0.7000, precision 0.7500, recall 0.6000, specificity 0.8000, FPR = 0.2000, $F_1 = 0.6667$ — and `confusion_matrix` is the *transpose* of the textbook layout.
3. A classifier outputs a score and predict silently thresholds it. One fixed set of scores gave accuracies from 0.7135 to 0.9181 across seven thresholds, **with AUC fixed at 0.9661 throughout**.
4. An ROC curve is a sorted list and some bookkeeping: up $1/P$ for a positive, right $1/N$ for a negative. The eleven-point hand sweep matched `roc_curve(drop_intermediate=False)` exactly; the default keeps 8 of the 11 points and the same area.
5. AUC two ways, same answer. Trapezoids $0.08 + 0.16 + 0.16 + 0.20 + 0.20 = 0.80$; concordant pairs $20/25 = 0.80$. It is the probability a random positive outranks a random negative, and the Mann–Whitney U rescaled. All 6848 pairs on `breast_cancer` matched `roc_auc_score` to $0.00e+00$.
6. AUC 0.5 is chance — 200 coin-flip runs gave mean 0.4994, sd 0.0136, max 0.5342. Below 0.5 is an *inverted* classifier: 0.9661 becomes 0.0339, and $AUC(-s) = 1 - AUC(s)$ exactly.
7. AUC is invariant to the threshold by construction and to the class balance by measurement: over a fiftyfold change in prevalence it moved inside a band of 0.0085 while average precision moved $0.1702 \rightarrow 0.9202$.
8. **Under heavy imbalance ROC flatters**. ROC AUC 0.9200, average precision 0.1702; at 90% recall that is 1156 false alarms for 54 true ones — one alert in twenty-two. **The PR baseline is the prevalence, not 0.5**.
9. The threshold is a business decision. Ten hand rows: 1:1 picks 0.60 (recall 0.80), 10:1 picks 0.35 (recall 1.00). On `breast_cancer`, 1:1 $\rightarrow$ 10:1 moved the threshold 0.6855 $\rightarrow$ 0.0670, misses 13 $\rightarrow$ 0, false alarms 1 $\rightarrow$ 42. Youden's J is the equal-cost answer in disguise.
10. Multiclass: micro always equals accuracy, weighted recall always equals accuracy, macro gives the rarest class a full vote. Hand example: macro F_1 0.676287 against weighted 0.785739. And a class with $F_1 = 0.0000$ had a one-vs-rest **AUC of 0.7793**.
11. Model selection: the best two of seven differed by **0.0003** with a fold-to-fold spread of **0.0058**, $p = 0.5569$, the leader losing 13 of 20 folds, and **six of seven models changing rank** when the metric changed. **A difference smaller than the fold-to-fold spread is not a difference.**

Before the assessment session

Read

notes.pdf — the full development, the sampling check on the AUC identity, the leakage argument for thresholds, and **10 practice questions with worked solutions**.

Do

Questions 1, 2, 3, 5 and 7, **with a pen**. Question 3 is an eight-row ROC sweep with the AUC computed both ways; question 5 costs the same eight rows at 8:1; question 7 is the multiclass averaging.

This closes the classification unit

The next test covers neural networks, support vector machines and **this lecture**, and the ROC sweep is the numeric question most likely to appear. The assignment that follows is a **classifier comparison with ROC and AUC** — the nine-line checklist two frames back is its marking scheme in disguise.

One habit to start now

Never quote an accuracy without saying **what the baseline scores and what threshold you used**. If you cannot answer both, you have not evaluated anything.

Materials: <https://github.com/tali7c/Applied-Machine-Learning>

Sources: Han, Kamber and Pei, *Data Mining: Concepts and Techniques* (3e), Ch. 8–9.; James et al., *An Introduction to Statistical Learning with Python*, Ch. 8–9.; Bishop, *Pattern Recognition and Machine Learning*, Ch. 4–5, 7, 14.; Zhang et al., *Dive into Deep Learning*, Ch. 5.