

Applied Machine Learning

Unit 07 — Lecture 24 Notes

Course Revision and Consolidation

Tofik Ali

Autumn 2026

Read this first

These notes are written so that you can learn the material *without* having been in the lecture. Every number you see was produced by running the code shown beside it, and every figure was generated from real computation rather than drawn by hand.

This is the last lecture of the course, and it is not a summary. A summary would list the topics; you already have twenty-three sets of notes that do that. The job here is different: to turn seven units into **one argument**, so that when someone hands you a problem you have never seen, you can say which method to try, why, and — the part that separates a competent answer from a good one — *what would tell you it was the wrong choice*.

The argument runs like this. Units II and III answered two separate questions: *what quantity are we minimising*, and *how is the minimisation carried out*. Unit IV established that what you feed the optimiser decides what it can possibly find, and that anything learned from the data belongs inside the fold. Units V and VI are the *same machinery* with different hypothesis classes and different losses: least squares, logistic regression, trees, ensembles, networks and support vector machines are not seven unrelated inventions. Unit VII removed the target column and showed what changes when there is nothing to score against.

Three results carry this set of notes. The first is Section 4: seven models on one dataset through one honest pipeline, where the best and second-best differ in AUC by **0.0013** and the fold-to-fold standard deviation is **0.0046** — so the nominal winner *lost 6 of the 20 folds*. The second is Section 6: five mistakes, each with a measured price, the largest of which is +0.2900 of pure fiction from fitting a feature selector outside a pipeline. The third is Section 7: one complete run from raw arrays with holes in them to a result you could defend in the project quiz.

Section 3 is the revision map. It is deliberately *short* about what is examinable, because a list of everything is the same as no list at all.

Contents

1	How to use these notes	2
1.1	What you need	3
1.2	Learning outcomes	4
2	One argument, not seven units	4
2.1	Unit II — what quantity are we minimising	5
2.2	Unit III — how the minimisation is carried out	5
2.3	Unit IV — what you feed the optimiser	5

2.4	Units V and VI — the same machinery, different classes	6
2.5	Unit VII — what changes when there is no target	6
3	The revision map	7
4	One dataset, every model	7
4.1	The protocol	7
4.2	What it produced	8
4.3	Is the winner the winner?	9
4.4	The rank depends on the metric	10
4.5	And what the seven cost	10
5	Choosing a method, and defending the choice	10
5.1	How to phrase a defensible choice	12
6	The five mistakes that cost the most	12
6.1	Mistake 1: fitting a transformer outside the pipeline	12
6.2	Mistake 2: tuning and evaluating on the same folds	14
6.3	Mistake 3: quoting accuracy under imbalance	14
6.4	Mistake 4: unscaled distance-based or gradient-based methods	15
6.5	Mistake 5: reading a coefficient as a cause	15
7	End to end: raw arrays to a defended result	16
7.1	Steps 0 to 3: look, split, baseline, one pipeline	17
7.2	Steps 4 to 6: tune inside, score outside, then score once	17
7.3	Step 7: the threshold is a separate decision	18
7.4	Steps 8 and 9: how sure are you, and what is the model	18
7.5	What a defended result reads like	19
8	The arithmetic that recurs across the units	19
8.1	Least squares on five points	19
8.2	Ridge, which is the same calculation with one extra term	20
8.3	Entropy and information gain	20
8.4	One forward and one backward pass	21
8.5	Bayes, reversing a conditional	21
8.6	One k -means iteration	22
8.7	A confusion matrix and one ROC point	22
8.8	One logistic gradient step, and one linkage divergence	22
9	The remaining internal assessment	23
10	Summary	24
11	Practice questions	24
12	Solutions	26
13	References and further reading	28

1 How to use these notes

Read Section 2 first and slowly. It is four pages and it is the whole course; everything after it is evidence for it. Then read Section 3, the revision map, and use it to decide what to reread

elsewhere. Sections 4 to 7 are the measurements: one dataset through every model, a table for choosing between them, the five expensive mistakes, and a complete worked pipeline. Section 8 collects the by-hand arithmetic that recurs across units, and it is the section to work with a pen rather than to read. Section 9 is short and factual: what is left to be assessed.

There are four kinds of box. **Key idea** — the one sentence to remember a month from now. **Worked example** — a calculation done by hand, with small numbers, before any library is used. **Caution** — the specific mistake that section exists to prevent. **Try it yourself** — something to run, with the expected output given so you can check.

Code appears in a framed block, and directly beneath it a shaded block showing what that code *actually printed*. If you run it and get something different, trust your machine and tell me. Timings are the one exception: those depend on your hardware, and only their *ratios* are meant to reproduce.

1.1 What you need

```
pip install numpy scipy scikit-learn matplotlib
```

Every number in these notes was produced with:

```
import numpy, scipy, sklearn, matplotlib
print("numpy", numpy.__version__, "| scipy", scipy.__version__,
      "| sklearn", sklearn.__version__,
      "| matplotlib", matplotlib.__version__)
```

```
numpy 2.2.6 | scipy 1.15.3 | sklearn 1.7.2 | matplotlib 3.10.9
```

Four bundled datasets recur and **none of them downloads anything**: `load_breast_cancer` (569×30), `load_wine` (178×13), `load_diabetes` (442×10) and `load_iris` (150×4). Everything else is a fixed-seed synthetic problem built with `make_classification` or `make_moons`. The random seed throughout is 24.

Do not reach for `fetch_*`

Every `fetch_*` loader downloads over HTTPS, and on the campus network those downloads fail with a certificate error you cannot fix from your machine. Nothing in this course depends on one.

One seed for the whole revision, and two sweeps

This lecture re-measures every claim the course makes, so it has to do it under *one* protocol rather than a dozen.

The one seed is 24. Every estimator that takes a `random_state`, every `train_test_split`, every shuffled `StratifiedKFold` and every synthetic set is built from it, and every listing below shows it rather than hiding it in a helper. (Where a second, independent split is needed — the outer loop of a nested cross-validation — it is written `random_state=25`, i.e. seed + 1, so that inner and outer folds are not the same folds.)

Two sweeps, where the variation is the point. Mistakes 1 and 2 are each measured on **twenty fresh noise problems**, not one, because a single trial of a noise experiment proves nothing — and they draw from *different* streams, so that the two results are independent rather than two views of the same twenty point sets.

Every number in these notes therefore reproduces exactly, with one honest exception: the millisecond columns of the model table, which no seed can fix and which are labelled where they appear (Section 4).

1.2 Learning outcomes

These notes serve all four course outcomes, which is why they are the last package: **CO1** — *explain the concepts and types of machine learning and the workflow of a learning problem*; **CO2** — *apply appropriate representations, similarity measures and pre-processing to a data set*; **CO3** — *evaluate models and interpret their results*; **CO4** — *apply and evaluate learning techniques and select an appropriate model for a given problem*. By the end you should be able to:

1. State the four questions that organise the entire course — what is being minimised, how, over what class of functions, and how you would know it worked — and place any method from the syllabus inside that frame. **CO1**
2. Say precisely what changes when the target column is removed, and why the burden of justification moves onto you. **CO1**
3. Choose and defend a pre-processing chain for a described data set, and state which steps must sit inside the cross-validation fold and why. **CO2**
4. Compare several models on one data set under a protocol that is honest about its own noise, and refuse to declare a winner when the gap is smaller than the fold-to-fold spread. **CO3**
5. Choose a metric from the properties of the problem — prevalence, cost asymmetry, the need to rank rather than to decide — and set a decision threshold from a stated cost ratio rather than from the default. **CO3**
6. Given a described situation — few rows and many columns, a requirement to explain, a non-convex boundary, heavy imbalance, no labels at all — name a method, justify it with a measured property rather than folklore, and name the observation that would refute the choice. **CO4**
7. Recognise and avoid the five mistakes of Section 6, and state the measured cost of each. **CO4**
8. Carry out, with a pen, the calculations that recur across the units: least squares on a small table, entropy and information gain, one forward and backward pass, a Bayes inversion, one k -means iteration, and a confusion matrix with its ROC point. **CO2, CO3, CO4**

2 One argument, not seven units

Every supervised method in this course is the same three decisions with different contents. Write them down as a template:

$$\hat{f} = \arg \min_{f \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n L(y_i, f(x_i)) + \lambda \Omega(f) \quad (1)$$

Four things in that line are choices, and the whole syllabus is an investigation of them. L is the **loss** — Unit II. The arg min is an **algorithm** — Unit III. $\mathcal{H}$ is the **hypothesis class** — Units V and VI. And Ω with its dial λ is **regularization**, which is the answer Unit V gave to three different pathologies at once. What is *not* in the line, and matters more than any of it, is the x_i : those columns were manufactured by you, and Unit IV is about that manufacture. Finally, nothing in the line tells you whether $\hat{f}$ is any good — that requires data the fit never saw, which is Unit IV again and then the whole of classification evaluation.




Figure 1: The course on one page. Read left to right along the top: choose a loss, minimise it, over a class of functions, and then prove it on data the fit never saw. The bottom-left box feeds all four. The bottom-right box is what happens when the target column is deleted.

2.1 Unit II — what quantity are we minimising

A loss is not a scoring rule applied after the fact. It is the specification of what the model will become. The demonstration that makes this concrete took six sensor readings, corrupted one of them, and asked for the single constant that best summarises them. Under squared error the answer moved from 12.33 to 25.67 — past every clean observation — while under absolute error it did not move at all. The reason is a one-line calculus fact: the constant minimising $\sum (y_i - c)^2$ is the *mean*, and the constant minimising $\sum |y_i - c|$ is the *median*.

Key idea

Choosing a loss is choosing what the model becomes, not how it is graded. Every property you later complain about — sensitivity to outliers, refusal to output a probability, indifference to a rare class — was decided here.

2.2 Unit III — how the minimisation is carried out

Given a loss, descent is three questions: how big a step, computed from how much data, and with how much memory of the last step. Those are the learning rate, the batch size and the momentum coefficient. The payoff of having separated Units II and III is that *the answer does not depend on which loss you picked*: the same gradient-descent loop fits a squared-error line, a cross-entropy classifier and a ten-thousand-parameter network.

Two measurements from that unit are worth keeping. There is a hard stability cliff: on one quadratic, $\eta = 1.00$ made the iterates cycle $0 \rightarrow 6 \rightarrow 0$ forever and $\eta = 1.05$ sent them to -17.18 in twenty steps. And the adaptive optimisers, tuned, all land within two per cent of each other; what they really buy is a *wider target* — Adagrad worked over 2.24 decades of learning rate against plain SGD's 1.31.

2.3 Unit IV — what you feed the optimiser

This is the unit that decides your ceiling. One engineered product column took a linear model from cross-validated $R^2 = 0.9065$ to 0.9978; a random forest reached the same place unaided, which is the useful form of the lesson — *feature engineering is worth most to the simplest models*.

Method	Hypothesis class $\mathcal{H}$	Loss L	Minimised by
Least squares	linear in x	squared error	closed form, $(X^\top X)^{-1}X^\top y$
Ridge, lasso	linear in x	squared + Ω	closed form; coordinate descent
Logistic	linear in the log-odds	cross-entropy	gradient descent, Newton/IRLS
Naive Bayes	product of conditionals	(counting)	closed-form counts
Decision tree	axis-aligned boxes	impurity reduction	greedy recursive search
Random forest	average of trees	impurity reduction	greedy + bootstrap + column sampling
Boosting	additive sum of trees	any differentiable	stagewise gradient steps
Neural network	composed affine + nonlinearity	any differentiable	backpropagation and SGD
SVM	linear in a feature map	hinge + $\ w\ ^2$	convex quadratic programme
k -NN	the training table itself	—	— (there is no fit)

Table 1: Nine supervised methods as three choices each. The differences between rows are real and they matter, but they are differences *within one template*, not between unrelated inventions.

The rule that runs through the entire unit, and that Section 6 prices, is one sentence.

Key idea

Anything you learn from the data — a mean, a median, a category list, a rotation, a feature subset, a resampled row, a hyper-parameter — is part of the model, and belongs inside the cross-validation fold.

2.4 Units V and VI — the same machinery, different classes

Table 1 is the point of these two units. Read it down the columns rather than across the rows: the rows differ in $\mathcal{H}$ and in L , and in almost nothing else.

Three connections in that table are worth stating explicitly, because they are the ones examinations ask about.

- **Least squares is not an arbitrary choice.** Assume Gaussian noise, write the likelihood, take logarithms, and the log-likelihood is a constant minus $\text{SSE}/(2\sigma^2)$. Maximising the likelihood *is* minimising the sum of squared errors. Change the noise to Laplace and the maximum-likelihood estimator becomes least absolute deviations.
- **Logistic regression is least squares' template with two substitutions.** $\mathcal{H}$ becomes linear in the log-odds rather than in y , and L becomes the Bernoulli negative log-likelihood. The gradient is then almost embarrassingly familiar: $\nabla L = X^\top(p - y)$, prediction minus truth, exactly as in the linear case.
- **An ensemble is not a better model; it is a bet on disagreement.** Bagging perturbs rows, forests additionally perturb columns, boosting changes the problem after each round. As the members become more correlated the benefit collapses; independence is the whole game.

2.5 Unit VII — what changes when there is no target

Delete the y column from equation (1) and two of its four choices lose their meaning. There is no $L(y_i, f(x_i))$ because there is no y_i ; and there is no held-out score because there is nothing to be right about. What replaces the loss is a **distance**, chosen by you; what replaces the score is an **assumption**, stated by you.

That is not a gap in the theory. It is the defining property of the subject: clustering *proposes* a structure rather than discovering a fact, and the proposal is a joint consequence of the data,

the distance, the algorithm and the settings. The measured consequence is stark. On *wine*, k -means agrees with the cultivar at ARI 0.3711 on the raw columns and 0.8975 after standardising, because one column carries 99.77% of the raw variance and Euclidean distance is therefore almost entirely a comparison of that one column. Nothing about the algorithm changed.

The commonest Unit VII error in scripts

“The clustering was wrong because it did not match the labels.” If you have labels, you are not clustering. On *iris*, two internal criteria agree on $k = 2$ and the species say $k = 3$; both are correct answers to different questions. The number of clusters in the data and the number of classes in the label column are different quantities.

3 The revision map

This section is deliberately short, and it is short on purpose. A revision list that includes everything is worth nothing, because it cannot help you allocate an evening. What follows is what is *genuinely* examinable: calculations you can be asked to produce with a pen, and derivations you can be asked to reproduce. Everything else you should be able to recognise and discuss, which is a much lower bar.

Which notes to reread, and in what order. If you have one evening: the notes on *loss functions* for the mean-versus-median demonstration, the notes on *cross-validation and model selection* for what belongs inside the fold, the notes on *decision trees* for entropy and gain, the notes on *neural networks* for the backward pass, and the notes on *classification evaluation* for the ROC staircase. Those five carry most of the computable marks in the course. If you have a week, add the *regularization* notes for the ridge invertibility proof and soft thresholding, the *Bayesian* notes for the base-rate inversion, and the *hierarchical clustering* notes for the two linkage runs.

What is not worth memorising. Library argument names. Coefficient tables you can look up. Any number you can recompute in ten seconds. Spend the time instead on the arithmetic in Section 8, all of which you can be asked to *produce*.

4 One dataset, every model

The most useful revision exercise in the whole course is also the simplest: take one dataset and run everything on it under one protocol. It settles, with numbers, a question that folklore gets wrong.

4.1 The protocol

```
X, y = load_breast_cancer(return_X_y=True) # 569 x 30, 357 benign
cv = RepeatedStratifiedKFold(n_splits=5, n_repeats=4, random_state=24)

for name, est in seven_models.items():
    for tr, te in cv.split(X, y):
        pipe = Pipeline([("sc", StandardScaler()), ("m", clone(est))])
        pipe.fit(X[tr], y[tr]) # scaler refitted every fold
        acc.append(accuracy_score(y[te], pipe.predict(X[te])))
        auc.append(roc_auc_score(y[te], score_of(pipe, X[te])))
```

Two choices in those six lines are the difference between a measurement and an anecdote. *Twenty* folds rather than five, so that the fold-to-fold spread is itself estimated. And the scaler

Unit	The one idea	What you can be asked to compute or derive
I	Learning is fitting a function from examples; three kinds of task	classifying a described problem as supervised, unsupervised or reinforcement; MAE, MSE and RMSE on a handful of numbers; a 1-NN prediction from a distance table
II	Choosing a loss chooses what the model becomes	minimising a loss over a constant and showing MSE gives the mean and MAE the median; writing cross-entropy for a Bernoulli label; explaining why a squared error is not robust
III	Step downhill: how big, from how much data, with how much memory	one gradient-descent step by hand on a quadratic; the divergence condition on η ; one momentum update and one Adam update including the bias correction
IV	What you feed it decides what it can find	the survival probability $(1-p)^d$ under listwise deletion; the z , IQR and modified- z outlier rules on a small column; one standardisation by hand; deciding what must sit inside the fold, and saying why
V	Least squares has an answer you can write down, and one dial controls it	$b_1 = S_{xy}/S_{xx}$, $b_0 = \bar{y} - b_1\bar{x}$, SSE and R^2 ; the normal equations on a 2×2 or 3×3 ; the Gaussian-MLE derivation; ridge $\hat{\beta} = (X^\top X + \lambda I)^{-1} X^\top y$ and the invertibility argument; soft thresholding
VI	One recipe, several hypothesis classes	entropy, information gain and gain ratio on a small table; a k -NN prediction and the effect of k ; Bayes with a prior, and naive Bayes with Laplace smoothing; one AdaBoost round (ε , α , reweighting); one forward and backward pass on a small network; the margin and α_i on a few separable points; a confusion matrix, precision, recall, F_1 , and a ROC point with its AUC
VII	Delete the target and the distance becomes the whole decision	building a proximity matrix; a complete agglomerative run under single and under complete linkage, matrix rewritten after every merge; one Lloyd iteration of k -means; the silhouette of one point; deciding whether a described dissimilarity is a metric


Table 2: The revision map. If a topic is not in the right-hand column it is still on the syllabus — but it will be tested by asking you to explain, not to compute.

inside the pipeline, so that the mean and standard deviation used to transform a validation row were computed without ever looking at it.

4.2 What it produced

```
X shape (569, 30)  positives 357  negatives 212  majority rate 0.6274
folds = 20

model      acc      sd      AUC      sd      fit ms  pred ms
LogisticRegression  0.9767  0.0096  0.9943  0.0059    2.3    0.24
KNN (k=11)  0.9675  0.0179  0.9898  0.0077    0.6    1.39
DecisionTree  0.9324  0.0191  0.9290  0.0198    6.1    0.35
RandomForest  0.9626  0.0163  0.9899  0.0073  341.3   7.08
GradientBoosting  0.9605  0.0212  0.9908  0.0085  366.9   0.62
MLP (100,)  0.9767  0.0122  0.9939  0.0067  225.9   0.44
SVM (RBF)  0.9745  0.0099  0.9956  0.0046    2.5    0.71
```



Two of those nine columns are measurements, not results

The accuracies, the standard deviations and the AUCs above reproduce exactly: fixed data, one seed, same answer on your machine as on mine. **The two millisecond columns do not.** They depend on your processor, your BLAS build and what else is running, and repeated runs of this identical script moved the MLP's fit time between 225 and 253 milliseconds without a line of code changing.

So the seconds are reported here as one measurement and are never the basis of a claim. When this lecture says a model is expensive, it says so in something countable — how many parameters were fitted, how many trees, how many support vectors — and those are the same everywhere:

model	fitted quantities
LogisticRegression	31 coefficients
KNN (k=11)	0 fitted, 13650 numbers stored
DecisionTree	31 nodes
RandomForest	300 trees, 10534 nodes
GradientBoosting	100 trees, 1406 nodes
MLP (100,)	3201 weights and biases
SVM (RBF)	106 support vectors

That table explains the millisecond column without depending on any machine. The MLP fits 3201 weights against logistic regression's 31 — a factor of **103**, exactly — for two accuracies that agree to four decimal places. k -NN fits nothing at all and stores 13,650 numbers, which is why it is the cheapest to fit and the dearest to predict.

Six of the seven sit inside a band of 0.017 in accuracy. The seventh, the single unpruned tree, is the only one clearly worse — and it is also the only model in the list with no averaging, no penalty and no margin anywhere in it. A majority-class baseline would have scored 0.6274, so **the distance from any real model to the baseline is roughly twenty times the distance between the real models.**

4.3 Is the winner the winner?

```
best by AUC      : SVM (RBF) 0.9956
second by AUC   : LogisticRegression 0.9943
gap              : 0.0013
fold-to-fold sd : 0.0046 (winner) 0.0059 (runner-up)
paired t-test   : t = 2.4697 p = 0.0232
folds the nominal winner LOST : 6 of 20
```

The paired test is significant at the conventional level and the result is *immaterial*. A gap of 0.0013 against a fold-to-fold spread of 0.0046 is a quarter of the noise; the nominal winner lost six folds out of twenty. On accuracy the same two models go the other way by +0.0022, with




Figure 3: Left: the per-fold AUC difference between the two best models. Six bars point the wrong way. Right: the same twenty differences as a histogram, with the mean gap and one standard deviation of the gap marked. The distribution straddles zero comfortably.

$p = 0.3965$, seven wins, four losses and **nine ties**.

Key idea

A difference smaller than the fold-to-fold spread is not a difference. Report the spread every time you report a score, and when the intervals overlap, choose on cost, interpretability or maintainability instead — because those you can actually measure.

4.4 The rank depends on the metric

rank	by accuracy	by AUC
1	LogisticRegression	SVM (RBF)
2	MLP (100,)	LogisticRegression
3	SVM (RBF)	MLP (100,)
4	KNN (k=11)	GradientBoosting
5	RandomForest	RandomForest
6	GradientBoosting	KNN (k=11)
7	DecisionTree	DecisionTree

models whose rank changes with the metric : 5 of 7

Five of seven change place. Accuracy and AUC are answering different questions: accuracy grades one threshold, AUC grades the whole ranking. Both are legitimate; neither is “the” score. Name the metric before you name the winner, and choose the metric from the problem.

4.5 And what the seven cost

Try it yourself

Rerun the protocol with `load_wine` in place of `load_breast_cancer` and with `n_repeats=10`. Two things should happen: every standard deviation grows, because wine has 178 rows and each fold holds only 36; and the ordering of the middle five models shuffles. If your conclusion changes when you change the seed, you did not have a conclusion.

5 Choosing a method, and defending the choice

Section 4 showed that on a clean, well-separated, moderately sized table the choice of model barely matters. That is the common case and it is worth knowing. This section is about the uncommon cases, where the choice matters enormously — and every row of Table 3 is grounded in a measurement rather than in received wisdom.




Figure 4: Fit and predict time per fold, log scale — a measurement on one machine, not a property of the algorithms. The countable version: gradient boosting fits 100 trees and 1406 nodes and scores no better than logistic regression’s 31 coefficients. k -NN is the mirror image — nothing fitted at all, and the highest predict cost of the seven, because it does all its work at predict time.




Figure 5: Four regimes in which the models are not interchangeable. Left to right: a sparse truth with more columns than rows; a curved decision boundary; one per cent positives; and no labels at all.

Reading the first row properly

The sparse-truth row is the one students most often get backwards, so it is worth spelling out. There were 200 columns and 50 rows, so ordinary least squares is already under-determined: `LinearRegression` returns the minimum-norm solution, which is ridge in the limit $\lambda \rightarrow 0$. That solution reached test $R^2 = +0.2926$. Cross-validated ridge chose the smallest α in its grid and did essentially the same thing, $+0.2492$. Cross-validated lasso reached $+0.9543$ and kept 31 columns, five of which were the five real ones.

Ridge did not fail because ridge is worse. It failed because the truth was *sparse* and ridge cannot express sparsity — on a swept path its smallest coefficient anywhere was 1.455×10^{-3} , never zero. Reverse the experiment, with a dense truth and correlated columns, and ridge wins instead. Neither penalty dominates, which is exactly why both are on the syllabus.

Two answers that are always wrong

“I tried them all and picked the best score.” That is a hyper-parameter search with the evaluation set as its objective, and Section 6 prices it at $+0.0515$ on pure noise, wrong on 17 of 20 trials.

“A neural network, because it is the most powerful.” On these 569 patients the MLP scored 0.9767 and logistic regression scored 0.9767 — identical to four decimals — and the MLP fitted 3201 weights against logistic regression’s 31, exactly $103\times$ as many. Capacity you do not need is capacity you have to regularise, tune and defend.

Situation	Try first	The measurement behind the advice
Few rows, many columns, and you believe the truth is sparse	lasso; elastic net if the columns are correlated	$n = 50$, $p = 200$, five real signals: OLS test $R^2 = +0.2926$, cross-validated ridge $+0.2492$, cross-validated lasso $+0.9543$ — and the lasso kept 5 of the 5 real columns. Ridge cannot help here because it never returns a zero
You must explain the model to someone	penalised logistic regression, or a shallow tree	on breast_cancer , thirty readable coefficients score $\text{AUC } 0.9947 \pm 0.0042$ against gradient boosting's 0.9919 ± 0.0059 . The readable model cost nothing at all — it was marginally ahead
The boundary is clearly non-convex	RBF SVM, k -NN, or a forest	two moons, 5-fold accuracy: logistic 0.8550, tree 0.9300, forest 0.9433, k -NN 0.9500, SVM 0.9517 . The linear model is nine points behind and no amount of tuning closes it
Heavy class imbalance	any model — but change the metric and the threshold	at 1% positives, always-say-no scores accuracy 0.9900; the model that finds 19 of 24 positives scores 0.8275. Average precision baseline is the prevalence, 0.0100, not 0.5
No labels at all	k -means or DBSCAN, after deciding the distance	k -means on wine : ARI 0.3711 raw, 0.8975 standardised. One column carried 99.77% of the raw variance. The pre-processing decision was worth more than the algorithm
Very many rows, and latency matters	a linear model; a forest if you can wait	40,000 rows: logistic fits in 0.04 s and predicts in under 0.01 s; the forest fits in 10.86 s; the SVM fits in 5.32 s and takes 2.54 s to predict 10,000 rows

Table 3: The choosing table. Each row names a property of the *problem*, not a preference.

5.1 How to phrase a defensible choice

A good answer has three parts, in this order.

1. **A property of the problem.** “There are 80 rows and 12,000 columns, and the clinician needs a short list of genes.”
2. **The family that property implies, with the mechanism.** “Lasso, because the L_1 penalty has a corner at zero and therefore returns exact zeros; ridge shrinks but never deletes.”
3. **What would refute it.** “If the true signal is dense rather than sparse — many small effects — the lasso will delete real predictors and ridge will beat it. I would check by comparing cross-validated scores for both, and by looking at how stable the selected set is across bootstrap resamples.”

The third part is the one almost everybody omits, and it is the one that distinguishes a student who has understood the course.

6 The five mistakes that cost the most

These five are ranked by what they actually cost, in marks and in accuracy. Four of them are score inflations and can be drawn on one axis (Figure 6); the fifth is not a score at all.

6.1 Mistake 1: fitting a transformer outside the pipeline

```
rng = np.random.default_rng(24)
```




Figure 6: Four mistakes with their measured price. Orange is the number the mistake produces; blue is the honest number. The fourth row runs the other way: forgetting the scaler *costs* you 0.0633 rather than inflating anything.

```
X = rng.normal(size=(100, 500)) # 500 columns of pure noise
y = rng.integers(0, 2, size=100) # coin-flip labels: the truth is 0.5

sel = SelectKBest(f_classif, k=10).fit(X, y) # <-- sees every row
wrong = cross_val_score(LogisticRegression(max_iter=5000),
                        sel.transform(X), y, cv=cv)

right = cross_val_score(Pipeline([("k", SelectKBest(f_classif, k=10)),
                                   ("m", LogisticRegression(max_iter=5000))]),
                        X, y, cv=cv)

for r in range(20): # a sweep: 20 FRESH noise problems,
    rg = np.random.default_rng(1000 + r) # one stream, twenty draws
    ... # the spread across them is the result
```

```
500 pure-noise columns, coin-flip labels; truth is 0.5000
selection outside the pipeline 0.8000
selection inside the pipeline 0.5100
overstatement +0.2900
over 20 fresh noise problems: leaky 0.7535, honest 0.5000
the leaky protocol scored higher on 20 of 20
```

There is no signal in that data at all. The labels are coin flips. And the leaky protocol reports an accuracy of 0.80, with a straight face, on *twenty out of twenty* fresh problems.

The mechanism is worth saying out loud, because once you have said it you will never make the mistake again. `SelectKBest` was shown the labels of every row, including the rows that later became the validation fold. Among 500 coin-flip columns, ten will look predictive by chance — and once those ten have been chosen *using the answers*, the validation fold is no longer unseen. The cross-validation that follows is measuring how well a model fits noise it was pointed at.

Caution

The same error wears several costumes, and the course measured all of them: scaling before the split, +0.0034; imputing before the split, nothing measurable; deciding which rows to drop before the split, +0.1702 of R^2 ; oversampling before the split, $F_1 = 0.8942$ against an honest 0.1133, because 100% of the minority test rows were byte-identical copies of training rows. The size of the damage varies enormously. The rule does not.




Figure 7: One model, one set of scores, one test set. ROC AUC is invariant to the class balance and looks excellent. Average precision, whose baseline is the prevalence rather than 0.5, is the honest summary at this prevalence.

6.2 Mistake 2: tuning and evaluating on the same folds

```
gs = GridSearchCV(pipe, grid_of_22_models, cv=inner).fit(X, y)
print(gs.best_score_) # a MAXIMUM over 22 noisy estimates

nested = cross_val_score(GridSearchCV(pipe, grid_of_22_models, cv=inner),
                          X, y, cv=outer) # an ESTIMATE of generalisation
```

```
the grid is 22 model-and-setting combinations
breast_cancer  best_score_ 0.9824  nested 0.9824 +/- 0.0125
               optimism    -0.0000
100x30 pure noise, coin-flip labels, 20 trials; truth 0.5000
  best_score_ 0.5750  nested 0.5235  optimism +0.0515
  best_score_ exceeded the nested estimate on 17 of 20
```

Report both halves of that, because both are instructive. On `breast_cancer` — an easy, well-separated, 569-row problem — the optimism was **zero to four decimal places**. On 100 rows of pure noise the same grid reported 0.5750 where the truth is 0.5000.

Key idea

`best_score_` is the maximum of a set of noisy estimates, and the maximum of noisy estimates is biased upwards. How badly depends on how many combinations you tried and how noisy each estimate was — and you cannot tell which regime you are in *without* the nested loop. That is the argument for always running it.

6.3 Mistake 3: quoting accuracy under imbalance

```
test prevalence 0.0100, 24 positives in 2400 rows
model          acc      AUC      AP      recall
always say no   0.9900   0.5000   0.0100   0.0000
logistic, default 0.9917   0.9084   0.4690   0.2500
logistic, class_weight 0.8275   0.8938   0.3035   0.7917
the balanced model finds 19 of the 24 positives; accuracy ranks it LAST
```

Read the last row twice. The only model in the table that actually *finds* the rare class — nineteen of the twenty-four positives — scores 0.8275, and accuracy therefore ranks it behind a model that has never predicted a positive in its life. If accuracy is your objective, the optimal




Figure 8: Left, the six models with and without a scaler in the pipeline; the two tree-based rows are exactly unchanged. Right, the column-by-column decomposition of the squared distance between two unscaled patients.

policy at one per cent prevalence is to do nothing.

Caution

Quote the confusion matrix, recall and precision at a stated threshold, and average precision with its baseline. If you quote a single number for an imbalanced problem, quote average precision and say what the prevalence was — “AP 0.4690 against a baseline of 0.0100” is informative; “99.2% accurate” is not.

6.4 Mistake 4: unscaled distance-based or gradient-based methods

model	raw	scaled	gain
KNN (k=11)	0.9401	0.9683	+0.0282
SVM (RBF)	0.9086	0.9719	+0.0633
MLP (100,)	0.9244	0.9754	+0.0510
LogisticRegression	0.9543	0.9771	+0.0229
DecisionTree	0.9368	0.9368	+0.0000
RandomForest	0.9560	0.9560	+0.0000

The two zeros are exact, not rounded. A tree splits on $x_j < t$, and that predicate is unchanged by any strictly increasing transformation of x_j — so the tree is the same tree, node for node, and its predictions are identical. Everything that computes a distance, an inner product or a gradient sees a completely different problem.

Why k -NN in particular collapses

Take two **breast_cancer** patients and decompose the squared Euclidean distance between them column by column. The widest column supplies **0.904482** of it, and the twenty-seven narrowest columns share 0.007655 between them. Unscaled k -NN on that table is not a thirty-feature classifier; it is a one-feature classifier with twenty-nine decorations. That is the whole of the +0.0282.

6.5 Mistake 5: reading a coefficient as a cause

diabetes, total serum cholesterol (s1)	
fitted on its own	+0.4723




Figure 9: Standardised coefficients for the ten diabetes predictors, fitted one at a time and fitted together. Four change sign. Standardising cannot change a sign; it only makes the ten bars comparable.

```
with the other nine present -1.0900
the sign flips and the magnitude grows x2.31
predictors that change sign this way: 4 of 10 -- age, sex, s1, s3
```

Both numbers are correct. They answer different questions. $+0.4723$ is the average difference in outcome between patients whose cholesterol differs by one unit, *averaged over everything else*. -1.0900 is the same difference *among patients matched on the other nine measurements*. Four of the ten predictors change sign between the two.

To make the point without any ambiguity, build a problem where the truth is known by construction:

```
Z = rng.normal(size=4000) # the confounder
X = Z + rng.normal(scale=0.5, size=4000) # X is caused by Z, and only Z
Y = 3.0 * Z + rng.normal(scale=0.5, size=4000) # X appears nowhere here
```

```
a designed confounder: X <- Z -> Y, X does not enter Y at all
slope of Y on X alone      +2.3662 (R^2 0.7790)
slope of Y on X given Z    -0.0119
true causal effect of X on Y +0.0000
```

A regression of Y on X reports a slope of $+2.3662$ and explains 78% of the variance. The causal effect of X on Y is exactly zero. No amount of data, no better optimiser and no larger model will fix this, because it is not a statistical error — it is a question about the world that the data cannot answer.

Key idea

The defensible wording is: “holding the other predictors in this model fixed, a one-unit increase in x_j is *associated with* a change of b_j in the fitted mean of y .” Not “causes”. Not “if we increase x_j ”. And never without naming which other predictors were in the model.

7 End to end: raw arrays to a defended result

Everything above is a fragment. This section is one continuous run, in the order you would really do it, on `breast_cancer` with three per cent of its cells deleted at random so that the imputer has work to do. The positive class is **malignant**, which is the clinically sensible framing and makes the cost asymmetry real.

Raw arrays to a defended result: the order is the content




Figure 10: The nine steps, with the measured result of each. The order is not a stylistic preference: steps 2 and 4 are what make step 7 meaningful.

7.1 Steps 0 to 3: look, split, baseline, one pipeline

```

X, y0 = load_breast_cancer(return_X_y=True)
y = 1 - y0 # positive = malignant
Xm = X.copy(); Xm[rng.random(X.shape) < 0.03] = np.nan

Xtr, Xte, ytr, yte = train_test_split(Xm, y, test_size=0.30,
                                      stratify=y, random_state=24)
base = DummyClassifier(strategy="most_frequent").fit(Xtr, ytr)

pipe = Pipeline([("imp", SimpleImputer(strategy="median")),
                 ("sc", StandardScaler()),
                 ("m", LogisticRegression(max_iter=5000,
                                          solver="liblinear"))])
  
```

```

step 0 X (569, 30), 506 missing cells (2.96%), 229 complete rows of 569
       positive class = malignant, 212 of 569 rows (0.3726)
step 1 train 398, test 171, stratified (prevalence 0.3719 / 0.3743)
step 2 majority class: accuracy 0.6257, AUC 0.5000
step 3 imputer -> scaler -> model, all three refitted inside every fold
  
```

Look at step 0 again. Three per cent of the *cells* are missing and only 229 of 569 rows are complete: listwise deletion would throw away **340** patients to remove 506 numbers. That is $(1 - p)^d$ with $p = 0.03$ and $d = 30$, which is 0.40, and the measured $229/569 = 0.402$. The exponential in the number of columns is the reason imputation exists.

Step 2 is the step most often skipped. A baseline is not a formality: at 0.6257 it tells you immediately that any accuracy below about 0.63 is worse than doing nothing, and it gives the reader something to compare your number against.

7.2 Steps 4 to 6: tune inside, score outside, then score once

```

grid = {"m__C": [0.001, 0.01, 0.1, 1, 10, 100], "m__penalty": ["l1", "l2"]}
  
```

```
inner = StratifiedKFold(5, shuffle=True, random_state=24)
outer = StratifiedKFold(5, shuffle=True, random_state=25)

nested = cross_val_score(GridSearchCV(pipe, grid, cv=inner,
                                      scoring="roc_auc"),
                        Xtr, ytr, cv=outer, scoring="roc_auc")
gs = GridSearchCV(pipe, grid, cv=inner, scoring="roc_auc").fit(Xtr, ytr)
s = gs.predict_proba(Xte)[: , 1] # the test set, touched ONCE
```

```
step 4 nested CV on the training rows: AUC 0.9937 +/- 0.0086
      (range 0.9787 .. 1.0000)
step 5 best params {'m__C': 0.1, 'm__penalty': 'l2'}, best_score_ 0.9941
      best_score_ - nested = +0.0004
step 6 held-out rows, scored once: AUC 0.9966, accuracy 0.9766
      confusion matrix TN 106 FP 1 FN 3 TP 61
      recall 0.9531 precision 0.9839
```

The optimism at step 5 is +0.0004 — negligible, because the grid has twelve entries and the problem is easy. That is a result, not a licence to skip the nested loop: you only know it is +0.0004 because you ran it.

Caution

The test set was touched exactly once, at step 6. If you look at a test score and then change anything — a hyper-parameter, a feature, a model — it is no longer a test set. It is a validation set, and you have no test set at all. This is the discipline that makes the number in your report mean something.

7.3 Step 7: the threshold is a separate decision

```
Youden J on the out-of-fold scores picks 0.4778 (J = 0.9502)
at 10:1 costs the out-of-fold minimum is at 0.2525
```

threshold	TN	FP	FN	recall	precision
0.5000 default	106	1	3	0.9531	0.9839
0.4778 Youden	106	1	3	0.9531	0.9839
0.2525 cost	101	6	1	0.9844	0.9130

Two things to notice, and one of them is a null result reported as it happened. Youden's J , chosen honestly on out-of-fold training scores, picked 0.4778 and changed **nothing at all** on the test rows. That is not a failure of the method: Youden is the equal-cost answer, the default 0.5 is close to it here, and on a well-separated problem there are simply no test rows scoring between 0.4778 and 0.5.

The cost-based threshold does move. Pricing a missed malignancy at ten false alarms takes the threshold to 0.2525 and the missed malignancies from **3 to 1**, at a price of five extra false alarms. Nothing about the model changed — the coefficients are identical — only the price list.

Key idea

A classifier outputs a score. A threshold turns that score into a decision, and the threshold is a *business* decision, not a statistical one. Choose it from the costs, on the training folds, and report which threshold you chose.

7.4 Steps 8 and 9: how sure are you, and what is the model

```
step 8 171 test rows: one patient is worth 0.0058 of accuracy;
      binomial se 0.0116, 95% interval [0.9540, 0.9993]
```

step 9 30 non-zero coefficients of 30, largest |coef| 0.4945

With 171 test rows, a single patient moves accuracy by 0.0058, and the binomial interval around 0.9766 is four and a half accuracy points wide. Any comparison between two models that differ by less than that, on this test set, is not a comparison.

7.5 What a defended result reads like

The paragraph to write

“A median-imputed, standardised, L_2 -penalised logistic regression ($C = 0.1$, chosen by a five-fold inner loop) scored a nested cross-validated AUC of 0.9937 ± 0.0086 on 398 training patients, and 0.9966 on 171 held-out patients scored once. The majority-class baseline is accuracy 0.6257 and AUC 0.5000. At a threshold of 0.2525, chosen on out-of-fold training scores for a ten-to-one miss-to-false-alarm cost, the model missed 1 of 64 malignancies at a price of 6 false alarms. The test set has 171 rows, so one patient is worth 0.0058 of accuracy and the binomial interval around the accuracy is $[0.9540, 0.9993]$.

This is one cohort from one institution. A prevalence shift, a different imaging protocol, or a column that encodes the outcome would each invalidate it, and none of those is tested here.”

That paragraph contains a protocol, a spread, a baseline, a stated threshold with its justification, a clinically meaningful count, an honest precision on the estimate, and a statement of what would change the conclusion. It is about a hundred and fifty words, and it is the deliverable of the entire course.

8 The arithmetic that recurs across the units

Work these with a pen. Every one has appeared in more than one unit, and every one is short enough to be asked for under time pressure.

8.1 Least squares on five points

Simple linear regression, by hand

$x = (1, 2, 3, 4, 5)$, $y = (2, 4, 5, 4, 5)$. Then $\bar{x} = 3$, $\bar{y} = 4$, and

$$S_{xx} = \sum (x_i - \bar{x})^2 = 10, \quad S_{xy} = \sum (x_i - \bar{x})(y_i - \bar{y}) = 6.$$

$$b_1 = \frac{S_{xy}}{S_{xx}} = 0.6, \quad b_0 = \bar{y} - b_1\bar{x} = 4 - 1.8 = 2.2.$$

Fitted values 2.8, 3.4, 4.0, 4.6, 5.2; residuals $-0.8, +0.6, +1.0, -0.6, -0.2$; SSE = 2.40. Both normal equations hold: $\sum e_i = 0$ and $\sum x_i e_i = 0$. With $S_{yy} = 6.00$, $R^2 = 1 - 2.40/6.00 = 0.6000$.

```
xbar 3.0000 ybar 4.0000 Sxx 10.0000 Sxy 6.0000
b1 0.6000 b0 2.2000 SSE 2.4000
sum e -0.0000 sum x*e -0.0000 SStot 6.0000 R^2 0.6000
LinearRegression slope 0.600000000000
```

The two normal equations are the whole derivation: differentiate $\sum (y_i - b_0 - b_1 x_i)^2$ with respect to b_0 and b_1 , set both to zero, and you get $\sum e_i = 0$ and $\sum x_i e_i = 0$. Everything else is algebra.

□

8.2 Ridge, which is the same calculation with one extra term

Ridge with no intercept

$x = (1, 2, 2)$, $y = (2, 3, 5)$, one predictor. Then $x^\top x = 9$ and $x^\top y = 2 + 6 + 10 = 18$, so

$$\hat{\beta}(\lambda) = \frac{x^\top y}{x^\top x + \lambda} = \frac{18}{9 + \lambda}.$$

lam	0	by hand	2.0000	sklearn Ridge	2.0000
lam	1	by hand	1.8000	sklearn Ridge	1.8000
lam	9	by hand	1.0000	sklearn Ridge	1.0000
lam	81	by hand	0.2000	sklearn Ridge	0.2000

Two facts fall out of that one line and both are examinable. First, $\hat{\beta} \rightarrow 0$ as $\lambda \rightarrow \infty$ but never reaches zero: ridge shrinks, it does not select. Second, $x^\top x + \lambda > 0$ for every $\lambda > 0$ whatever x is — which in the matrix case is the statement that $X^\top X + \lambda I$ is positive definite and therefore invertible even when $X^\top X$ is singular. Ridge has an answer in situations where ordinary least squares does not exist. $\square$

8.3 Entropy and information gain

The fourteen-row table

Nine “yes” and five “no”:

$$H(S) = -\frac{9}{14} \log_2 \frac{9}{14} - \frac{5}{14} \log_2 \frac{5}{14} = 0.940286 \text{ bits.}$$

Splitting on Outlook gives Sunny (2, 3), Overcast (4, 0), Rain (3, 2), so

$$\text{Rem} = \frac{5}{14}(0.970951) + \frac{4}{14}(0) + \frac{5}{14}(0.970951) = 0.693536,$$

$$\text{Gain(Outlook)} = 0.940286 - 0.693536 = 0.246750.$$

```

H(S) = H(9/14, 5/14) = 0.940286 bits
Sunny      (2 yes, 3 no)  weight 5/14  H = 0.970951
Overcast   (4 yes, 0 no)  weight 4/14  H = 0.000000
Rain       (3 yes, 2 no)  weight 5/14  H = 0.970951
Remainder(Outlook) = 0.693536  Gain(Outlook) = 0.246750
High H = 0.985228  Normal H = 0.591673
Remainder(Humidity) = 0.788450  Gain(Humidity) = 0.151836
Gini(S) = 0.459184

```

Outlook wins because it is the only attribute with a *pure child*: four Overcast days, all “yes”, entropy exactly zero. That one branch removes $\frac{4}{14} \times 0.940286 = 0.2686$ from the remainder before the other two branches contribute anything.

Caution

The weights are what stop a many-valued attribute from being free. A fourteen-valued row identifier scores gain 0.940286 — the largest possible — because every child is pure with one row in it. Gain ratio divides by the split information, $\log_2 14 = 3.807355$, a 73.7% cut, and at fourteen rows that *still* is not enough to demote the identifier. Textbooks that say gain ratio fixes the bias are being sloppy.

8.4 One forward and one backward pass

A 2-2-1 sigmoid network, nine parameters

Input $x = (1, 0.5)$, target $t = 1$, loss $L = \frac{1}{2}(a_2 - t)^2$,

$$W^{(1)} = \begin{pmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{pmatrix}, \quad b^{(1)} = (0.1, 0.1), \quad W^{(2)} = (0.5, -0.5), \quad b^{(2)} = 0.2.$$

Forward: $z^{(1)} = (0.3, 0.6)$, $a^{(1)} = (0.574443, 0.645656)$, $z_2 = 0.1643931053$, $a_2 = 0.5410059687$, $L = 0.1053377604$.

Backward, using $\sigma'(z) = a(1 - a)$:

$$\delta_2 = (a_2 - t) a_2 (1 - a_2) = -0.1139767142,$$

$$\frac{\partial L}{\partial W^{(2)}} = \delta_2 a^{(1)} = (-0.06547307, -0.07358978), \quad \frac{\partial L}{\partial b^{(2)}} = \delta_2,$$

$$\delta_1 = \delta_2 W^{(2)} \odot a^{(1)} \odot (1 - a^{(1)}) = (-0.01393128, +0.01303804), \quad \frac{\partial L}{\partial W^{(1)}} = \delta_1 x^\top.$$

```
z1 = [0.300000, 0.600000]    a1 = [0.574443, 0.645656]
z2 = 0.1643931053    a2 = 0.5410059687    L = 0.1053377604
delta2 = -0.1139767142
dL/dW2 = [-0.06547307, -0.07358978]    dL/db2 = -0.11397671
delta1 = [-0.01393128, 0.01303804]
dL/dW1 = [[-0.01393128, -0.00696564], [0.01303804, 0.00651902]]
dL/db1 = [-0.01393128, 0.01303804]
max |analytic - central difference| = 2.095e-11
one step at eta = 0.5: W2 -> [0.532737, -0.463205]
```

The check at the end is the habit worth keeping: every analytic derivative agrees with a central finite difference to 2×10^{-11} . And notice the asymmetry that makes deep learning possible at all — one forward pass and one backward pass produced *all nine* derivatives, where nine separate finite differences would have needed eighteen forward passes.

8.5 Bayes, reversing a conditional

Prevalence 1%, sensitivity 99%, specificity 95%

$$P(D | +) = \frac{P(+ | D)P(D)}{P(+ | D)P(D) + P(+ | \bar{D})P(\bar{D})} = \frac{0.99 \times 0.01}{0.99 \times 0.01 + 0.05 \times 0.99} = \frac{0.0099}{0.0594} = 0.1667.$$

As frequencies, out of 10,000 people: 100 are ill and 99 of them test positive; 9,900 are well and 495 of them test positive. So of the 594 positive results, 99 are real — one in six.

```
P(D|+) = 0.009900 / 0.059400 = 0.166667
in 10,000 people: 99 true positives and 495 false positives
odds form: prior odds 0.010101 x LR+ 19.8000 = posterior odds 0.200000
back to a probability: 0.166667
```

The odds form is worth learning because it is one multiplication: posterior odds = prior odds $\times$ likelihood ratio, where $LR^+ = \text{sensitivity} / (1 - \text{specificity}) = 19.8$. A test that multiplies your odds by twenty still leaves you at one in six if you started at one in a hundred. $\square$

8.6 One k -means iteration

Six points, $k = 2$, one Lloyd step

Points $P_1(1, 1)$, $P_2(1.5, 2)$, $P_3(3, 4)$, $P_4(5, 7)$, $P_5(3.5, 5)$, $P_6(4.5, 5)$; initial centres $c_1 = (1, 1)$ and $c_2 = (5, 7)$. *Assign*: distances to c_1 are 0, 1.1180, 3.6056, 7.2111, 4.7170, 5.3151 and to c_2 are 7.2111, 6.1033, 3.6056, 0, 2.5000, 2.0616. So $\{P_1, P_2, P_3\}$ go to cluster 1 and $\{P_4, P_5, P_6\}$ to cluster 2 — note that P_3 is **exactly equidistant** at 3.6056 and the tie has to be broken by a stated rule. *Update*: $c_1 = (1.8333, 2.3333)$, $c_2 = (4.3333, 5.6667)$.

```
new c1 = (1.8333, 2.3333)    new c2 = (4.3333, 5.6667)
W with the old centres 24.750000, with the new ones 10.666667
(fell by 14.083333)
sklearn after one iteration: c1 (1.8333, 2.3333)  c2 (4.3333, 5.6667)
```

Both halves of a Lloyd iteration can only decrease the within-cluster sum of squares — reassignment because each point moves to a nearer centre, and recentring because the mean is the minimiser of squared distance. W is bounded below and the number of partitions is finite, so the algorithm terminates. It terminates at a *local* minimum, which is why `n_init` exists. $\square$

8.7 A confusion matrix and one ROC point

Twenty patients

TP = 6, FP = 2, FN = 4, TN = 8.

$$\text{accuracy} = \frac{14}{20} = 0.7000, \quad \text{precision} = \frac{6}{8} = 0.7500, \quad \text{recall} = \frac{6}{10} = 0.6000,$$

$$\text{specificity} = \frac{8}{10} = 0.8000, \quad \text{FPR} = 0.2000, \quad F_1 = \frac{2(0.75)(0.6)}{0.75 + 0.6} = 0.6667.$$

The ROC point is $(\text{FPR}, \text{TPR}) = (0.2000, 0.6000)$ and Youden's $J = 0.6000 + 0.8000 - 1 = 0.4000$.

```
TP 6  FP 2  FN 4  TN 8   (n = 20)
accuracy 0.7000 precision 0.7500 recall 0.6000 specificity 0.8000
FPR 0.2000  F1 0.6667  ROC point (0.2000, 0.6000)  Youden J 0.4000
ten scored rows -> AUC 0.8400
concordant pairs 21 of 25 = 0.8400
staircase (FPR, TPR):
(0.0,0.0) (0.0,0.2) (0.0,0.4) (0.2,0.4) (0.2,0.6) (0.2,0.8)
(0.4,0.8) (0.4,1.0) (0.6,1.0) (0.8,1.0) (1.0,1.0)
```

The staircase came from ten scored rows, five positive and five negative: sort by score descending, step *up* by $1/5$ for a positive and *right* by $1/5$ for a negative. The area is 0.8400, and independently 21 of the 25 positive–negative pairs are correctly ordered — the same number, because AUC *is* the probability that a randomly chosen positive outscore a randomly chosen negative.

Caution

`sklearn.metrics.confusion_matrix` returns the **transpose** of the layout most textbooks use: top-left is TN and bottom-right is TP. Unpack it as `tn, fp, fn, tp = cm.ravel()` and you will never get it wrong.

8.8 One logistic gradient step, and one linkage divergence

```
p = [0.5000, 0.5000, 0.5000]  X'(p - y) = [-0.5000, -2.0000]
```

```
one step at eta = 0.5: w = [0.2500, 1.0000]
log-likelihood -2.0794 -> -1.6402

single   heights [1.      2.      2.2361 3.1623 4.      ]
complete heights [1.      2.2361 4.      4.4721 8.6023]
```

At $w = 0$ every predicted probability is exactly 0.5, so the gradient is $X^\top(p - y)$ with $p = \frac{1}{2}\mathbf{1}$ — pure arithmetic, no exponentials needed. And the two linkage height sequences agree only on the first merge, which is the general fact: with all clusters singletons, the minimum, maximum and mean over a set of *one* cross distance are the same number, so *the first merge is always the smallest entry of the original matrix, whichever linkage you chose*.

9 The remaining internal assessment

This section states only what is already in the course handbook. Nothing here is new, and no dates are given, because the calendar is announced on the LMS and in class rather than fixed in a document.

Theory quizzes. Ten quizzes, T1–T10, each of **40 minutes**, make up 10 marks. T10 is the last and covers Unit VII’s algorithms: hierarchical, partitioning and density-based clustering. Each is marked out of 10 and the component is the total scaled to 10. The quizzes come in several versions; every version asks equally demanding questions on the same topics, for the same marks, against the same rubric.

Laboratory. Ten lab assignments (10 marks) and ten lab quizzes (10 marks). LA10 covers the last two experiments — clustering medical images and crop disease classification — and LQ10 follows it. Each lab quiz runs for **60 minutes at the machine** and asks you to read, trace, debug and modify *your own submitted code*. Each assignment is out of 10 with the usual choice structure: any 2 of 3 easy questions, any 1 of 2 medium ones, and a compulsory consolidation part.

Project. 15 marks in two pieces. The **repository submission** is 5 marks awarded to the group — identical for every member — for the completeness, correctness, reproducibility, documentation *and commit history* of the GitHub repository, which is the only thing submitted. The **project quiz** is 10 marks and is *marked individually* — a written paper on your own group’s project, authored in one set per theme, with every member sitting it separately and no oral component. The remaining milestones are the final submission, which is the repository URL, and then the individual written quiz.

Theory assignments. Two assignments, 5 marks in total. Each has a submission and a short in-class assessment test on what you submitted.

What the project quiz is actually asking

Your project, in your words: what the problem was, what data you had, what you cleaned and why, which models you compared, under what protocol, with what spread, what you chose and why, and what you would do differently. Every one of those is a question this course has spent a semester answering. If you can run your own pipeline end to end and explain each line of it, you are prepared.

Caution

Two thirds of the project mark is individual, and group marks are not divided by contribution. A strong contributor in a weak group can still score well; a passenger in a strong group cannot, because the group mark is only 5 of the 15.

10 Summary

1. **The course is one template with four choices.** A loss, an algorithm, a hypothesis class and a penalty — plus the columns you manufactured, which decide the ceiling, and the held-out data, which is the only evidence that any of it worked.
2. **Seven models on one clean dataset were separated by less than the fold-to-fold noise.** Best AUC 0.9956, second 0.9943, gap 0.0013, fold standard deviation 0.0046, and the nominal winner lost six folds of twenty. Five of the seven change rank if you change the metric.
3. **Choose from a property of the problem.** Sparse truth and $p \gg n$: lasso, +0.9543 against ridge's +0.2492. A curved boundary: SVM or k -NN, 0.9517 against a linear model's 0.8550. Heavy imbalance: change the metric and the threshold, because always-say-no scores 0.9900. No labels: spend your effort on the distance, worth $0.3711 \rightarrow 0.8975$ on **wine**.
4. **Five mistakes, priced.** A transformer outside the pipeline, +0.2900 of pure fiction. Tuning and evaluating on the same folds, +0.0515 on noise. Accuracy under imbalance, which ranks doing nothing above the model that works. No scaler in front of an RBF SVM, -0.0633 . And a coefficient read as a cause, +2.3662 where the true effect is exactly zero.
5. **A defended result** names its protocol, its spread, its baseline, its threshold and its caveats — and it comes from a test set touched exactly once.
6. **The honest finding of this whole lecture** is that the method is rarely what decides the answer. The data, the columns and the protocol are. That is not a disappointing conclusion; it is where the work is.

11 Practice questions

Fifteen questions spanning the whole course, weighted towards the by-hand computations that recur across units. Work them with a pen before reading Section 12.

1. **(Unit V, least squares.)** For $x = (0, 1, 2, 3, 4)$ and $y = (1, 3, 2, 5, 4)$, compute $\bar{x}$, $\bar{y}$, S_{xx} , S_{xy} , b_1 , b_0 , the five residuals, SSE and R^2 . Verify both normal equations.
2. **(Unit II, loss.)** You have readings $(4, 5, 6, 7, 100)$. Find the constant that minimises the mean squared error and the constant that minimises the mean absolute error. Which would you report as a summary of the sensor, and why?
3. **(Unit V, regularization.)** With a single predictor, $x = (2, 2, 1)$ and $y = (4, 6, 3)$, and no intercept, write $\hat{\beta}(\lambda)$ and evaluate it at $\lambda = 0, 3, 9$. Then state, in one sentence, why $x^\top x + \lambda$ can never be zero for $\lambda > 0$ and what that implies in the matrix case.
4. **(Unit VI, trees.)** A node holds 20 rows: 12 “yes” and 8 “no”. An attribute A splits it into $(8, 2)$ and $(4, 6)$. Compute $H(S)$, both child entropies, the remainder and the information gain. Compute the Gini index of the parent as well.

5. **(Unit VI, trees.)** A second attribute B splits the same node into four children of five rows each, with class counts $(5, 0)$, $(4, 1)$, $(2, 3)$ and $(1, 4)$. Compute its gain, compare with A , then compute the split information and the gain ratio for both. What does the comparison illustrate?
6. **(Unit VI, networks.)** A 1–1–1 network with sigmoid activations has $w_1 = 0.5$, $b_1 = 0$, $w_2 = 1.0$, $b_2 = 0$. For input $x = 2$ and target $t = 0$ with $L = \frac{1}{2}(a_2 - t)^2$, compute the forward pass and all four derivatives. Then take one gradient step at $\eta = 1$ and state the new loss.
7. **(Unit VI, Bayes.)** A disease has prevalence 2%. A test has sensitivity 0.90 and specificity 0.90. Compute $P(D \mid +)$ algebraically, as natural frequencies out of 10,000, and in odds form. Then recompute with prevalence 20% and comment.
8. **(Unit VII, k -means.)** Points $A(1, 1)$, $B(2, 1)$, $C(4, 3)$, $D(5, 4)$; initial centres $c_1 = (1, 1)$, $c_2 = (2, 1)$. Carry out two complete Lloyd iterations, showing the assignment and the new centres each time, and report W after each update.
9. **(Unit VI, evaluation.)** A model on 200 patients gives $TP = 30$, $FP = 20$, $FN = 10$, $TN = 140$. Compute accuracy, precision, recall, specificity, F_1 and the ROC point. Then compute the accuracy of a majority-class classifier, and say which of the two you would deploy.
10. **(Unit VI, ROC.)** Six rows scored 0.9, 0.8, 0.7, 0.4, 0.3, 0.1 with true labels 1, 0, 1, 1, 0, 0. Draw the ROC staircase, compute the AUC by trapezoids, and confirm it by counting concordant pairs.
11. **(Unit IV, protocol.)** A colleague reports 96% cross-validated accuracy from this script.

```
X = SimpleImputer().fit_transform(X)
X = StandardScaler().fit_transform(X)
X = SelectKBest(k=15).fit_transform(X, y)
print(cross_val_score(SVC(C=10, gamma=0.01), X, y, cv=5).mean())
```

12. **(Unit IV, missing data.)** A table has 25 columns and each cell is missing independently with probability 0.04. What fraction of rows survives listwise deletion? How many columns could you afford before fewer than half the rows survive?
13. **(CO4, choosing.)** For each situation, name one method, give one measured reason, and name one observation that would tell you the choice was wrong. (a) 120 rows, 9,000 columns, and a biologist needs a short list. (b) A payments system, 4 million rows, decision required in 2 ms. (c) 5,000 emails, 60 of them phishing, a missed phish costs 50 false alarms. (d) 200,000 unlabelled sensor traces.
14. **(CO3, honesty.)** Model A scores 0.9412 and model B scores 0.9388 over ten folds, with fold standard deviations 0.0201 and 0.0187 and a paired-difference standard deviation of 0.0155. What do you report? What would you need in order to claim A is better?
15. **(CO1–CO4, synthesis.)** In no more than 200 words, place ridge regression, a random forest and DBSCAN inside the four-question frame of Section 2: what is minimised, how, over what class, and how you would know it worked. For DBSCAN, say precisely which of the four questions has no answer and why.

12 Solutions

1. $\bar{x} = 2$, $\bar{y} = 3$. Deviations in x : $(-2, -1, 0, 1, 2)$; in y : $(-2, 0, -1, 2, 1)$. So $S_{xx} = 4 + 1 + 0 + 1 + 4 = 10$ and $S_{xy} = 4 + 0 + 0 + 2 + 2 = 8$. Hence $b_1 = 8/10 = 0.8$ and $b_0 = 3 - 0.8(2) = 1.4$. Fitted: 1.4, 2.2, 3.0, 3.8, 4.6; residuals $-0.4, +0.8, -1.0, +1.2, -0.6$. $SSE = 0.16 + 0.64 + 1.00 + 1.44 + 0.36 = 3.60$. $S_{yy} = 4 + 0 + 1 + 4 + 1 = 10$, so $R^2 = 1 - 3.6/10 = 0.6400$. Normal equations: $\sum e_i = 0$ exactly, and $\sum x_i e_i = 0(-0.4) + 1(0.8) + 2(-1.0) + 3(1.2) + 4(-0.6) = 0.8 - 2.0 + 3.6 - 2.4 = 0$.

2. MSE is minimised by the mean, $122/5 = 24.4$; MAE by the median, 6. Report the median, and then report that you did: the mean lies above every observation except the corrupted one, which is the diagnostic that the reading at 100 is not a measurement. Squared error is minimised by the mean precisely because it pays the square of a large deviation, so one outlier can outvote four good readings.

3. $x^\top x = 4 + 4 + 1 = 9$ and $x^\top y = 8 + 12 + 3 = 23$, so $\hat{\beta}(\lambda) = 23/(9 + \lambda)$, giving 2.5556, 1.9167 and 1.2778 at $\lambda = 0, 3, 9$. Since $x^\top x = \|x\|^2 \geq 0$ and $\lambda > 0$, the denominator is strictly positive. In the matrix case the same argument reads $v^\top (X^\top X + \lambda I)v = \|Xv\|^2 + \lambda \|v\|^2 > 0$ for every $v \neq 0$, so $X^\top X + \lambda I$ is positive definite and therefore invertible — *whatever* X is, including $p > n$. $\square$

4. $H(S) = H(0.6, 0.4) = -0.6 \log_2 0.6 - 0.4 \log_2 0.4 = 0.970951$. Child 1 is (8, 2): $H = H(0.8, 0.2) = 0.721928$. Child 2 is (4, 6): $H = H(0.4, 0.6) = 0.970951$. Remainder $= 0.5(0.721928) + 0.5(0.970951) = 0.846439$. Gain $= 0.970951 - 0.846439 = 0.124512$. Gini of the parent $= 1 - 0.6^2 - 0.4^2 = 0.480000$.

5. Child entropies: $(5, 0) \rightarrow 0$; $(4, 1) \rightarrow H(0.8, 0.2) = 0.721928$; $(2, 3) \rightarrow H(0.4, 0.6) = 0.970951$; $(1, 4) \rightarrow 0.721928$. Each weight is $5/20 = 0.25$, so the remainder is $0.25(0 + 0.721928 + 0.970951 + 0.721928) = 0.603702$ and $\text{Gain}(B) = 0.970951 - 0.603702 = 0.367249$, comfortably ahead of A 's 0.124512. Split information: for A , two children of ten, $SI = \log_2 2 = 1$, so gain ratio $= 0.124512$; for B , four children of five, $SI = \log_2 4 = 2$, so gain ratio $= 0.183625$. B still wins, but its advantage has fallen from a factor of 2.95 to a factor of 1.47. That is exactly the correction gain ratio exists to make: plain gain rewards attributes for having many values, and the penalty grows with $\log_2$ of the number of children.

6. Forward: $z_1 = 0.5(2) = 1.0$, $a_1 = \sigma(1) = 0.731059$; $z_2 = 0.731059$, $a_2 = \sigma(0.731059) = 0.675038$; $L = \frac{1}{2}(0.675038)^2 = 0.227838$. Backward: $\delta_2 = (a_2 - 0)a_2(1 - a_2) = 0.675038 \times 0.675038 \times 0.324962 = 0.148086$. Then $\partial L / \partial w_2 = \delta_2 a_1 = 0.108250$ and $\partial L / \partial b_2 = 0.148086$. $\delta_1 = \delta_2 w_2 a_1 (1 - a_1) = 0.148086 \times 1.0 \times 0.731059 \times 0.268941 = 0.029117$, so $\partial L / \partial w_1 = \delta_1 x = 0.058233$ and $\partial L / \partial b_1 = 0.029117$. One step at $\eta = 1$: $w_2 = 0.891750$, $b_2 = -0.148086$, $w_1 = 0.441767$, $b_1 = -0.029117$. Recomputing gives $a_1 = 0.706930$, $z_2 = 0.482332$, $a_2 = 0.618288$ and $L = 0.191140$ — down from 0.227838, as a descent step must be for a small enough η .

7. At 2%: $P(D \mid +) = (0.9)(0.02)/[(0.9)(0.02) + (0.1)(0.98)] = 0.018/0.116 = 0.155172$. Frequencies out of 10,000: 200 ill, 180 test positive; 9,800 well, 980 test positive; $180/1160 = 0.155172$. Odds: prior odds $= 0.02/0.98 = 0.020408$, $LR^+ = 0.9/0.1 = 9$, so posterior odds $= 0.183673$ and $p = 0.183673/1.183673 = 0.155172$. At 20%: prior odds 0.25, posterior odds 2.25, $p = 0.692308$. The same test, unchanged, moves you to 16% or to 69% depending only on the prior. That is the whole content of the base-rate result: a test result is evidence, not a verdict, and evidence multiplies odds.

8. *Iteration 1.* $d(A, c_1) = 0$, $d(A, c_2) = 1$; $A \rightarrow 1$. $d(B, c_1) = 1$, $d(B, c_2) = 0$; $B \rightarrow 2$. $d(C, c_1) = \sqrt{9+4} = 3.6056$, $d(C, c_2) = \sqrt{4+4} = 2.8284$; $C \rightarrow 2$. $d(D, c_1) = \sqrt{16+9} = 5$, $d(D, c_2) = \sqrt{9+9} = 4.2426$; $D \rightarrow 2$. New centres: $c_1 = (1, 1)$, $c_2 = (\frac{11}{3}, \frac{8}{3}) = (3.6667, 2.6667)$. $W = 0 + [(2 - 3.6667)^2 + (1 - 2.6667)^2] + [(4 - 3.6667)^2 + (3 - 2.6667)^2] + [(5 - 3.6667)^2 + (4 - 2.6667)^2] = 5.5556 + 0.2222 + 3.5556 = 9.3333$. *Iteration 2.* $d(A, c_1) = 0$, $d(A, c_2) = \sqrt{7.1111 + 2.7778} = 3.1447$; $A \rightarrow 1$. $d(B, c_1) = 1$, $d(B, c_2) = \sqrt{2.7778 + 2.7778} = 2.3570$; $B \rightarrow 1$. $d(C, c_1) = 3.6056$, $d(C, c_2) = 0.4714$; $C \rightarrow 2$. $d(D, c_1) = 5$, $d(D, c_2) = 1.8856$; $D \rightarrow 2$. New centres: $c_1 = (1.5, 1)$, $c_2 = (4.5, 3.5)$, and $W = 0.5 + 0.5 = 1.0$. The next iteration changes nothing, so the algorithm has converged.

9. Accuracy = $170/200 = 0.8500$; precision = $30/50 = 0.6000$; recall = $30/40 = 0.7500$; specificity = $140/160 = 0.8750$; $F_1 = 2(0.6)(0.75)/1.35 = 0.6667$; ROC point (0.1250, 0.7500). The majority class is the negative one, 160 of 200, so a majority-class classifier scores 0.8000 — only five points behind, with recall 0. Deploy the real model, and *report the recall and the confusion matrix rather than the accuracy*, because the accuracy gap of 0.05 badly understates the difference between finding 30 cases and finding none.

10. Sorted by score: (0.9, 1), (0.8, 0), (0.7, 1), (0.4, 1), (0.3, 0), (0.1, 0). With three positives and three negatives each step is $1/3$. The staircase runs $(0, 0) \rightarrow (0, \frac{1}{3}) \rightarrow (\frac{1}{3}, \frac{1}{3}) \rightarrow (\frac{1}{3}, \frac{2}{3}) \rightarrow (\frac{1}{3}, 1) \rightarrow (\frac{2}{3}, 1) \rightarrow (1, 1)$. Area: the region from FPR = 0 to $\frac{1}{3}$ has height $\frac{1}{3}$, contributing $\frac{1}{9}$; from $\frac{1}{3}$ to 1 the height is 1, contributing $\frac{2}{3}$. Total = $\frac{1}{9} + \frac{2}{3} = \frac{7}{9} = 0.7778$. Concordant pairs: the positive at 0.9 beats all three negatives (3); the positive at 0.7 beats 0.3 and 0.1 (2); the positive at 0.4 beats 0.3 and 0.1 (2). That is 7 of 9, = 0.7778. $\square$

11. Four problems. (i) the imputer is fitted on all rows; (ii) the scaler is fitted on all rows; (iii) the selector is fitted on all rows *and all labels* — this is the serious one, worth up to +0.2900 on data with no signal at all; (iv) C and γ have evidently been chosen somehow, and if they were chosen by looking at this same score then the number is a maximum rather than an estimate. There is also no `random_state`, no stratification and no spread reported. The selector leak matters most, because its damage scales with the number of columns and does not shrink as n grows. The fix is one object: `Pipeline([("imp", SimpleImputer()), ("sc", StandardScaler()), ("sel", SelectKBest(k=15)), ("m", SVC())])` inside a `GridSearchCV`, itself inside `cross_val_score`.

12. A row survives with probability $(1 - 0.04)^{25} = 0.96^{25} = 0.3604$, so about 36%. For half the rows to survive we need $0.96^d \geq 0.5$, that is $d \leq \ln 0.5 / \ln 0.96 = 16.98$, so **16 columns**. Listwise deletion is exponential in the number of columns, which is why it becomes unusable on wide tables at missingness rates that look harmless per cell.

13. (a) Lasso or elastic net: with $p \gg n$ it is the only family that returns an explicit list, and on a measured $n = 50$, $p = 200$ problem it scored +0.9543 against ridge's +0.2492 while recovering all five real columns. *Refuted if* the selected set is unstable across bootstrap resamples, or if a dense-truth check shows ridge doing better — in which case the signal is many small effects, not a few large ones. (b) Logistic regression: it fitted 30,000 rows in 0.04 s and predicts in constant time, whereas k -NN's predict cost grows with n and an RBF SVM took 2.54 s for 10,000 rows. *Refuted if* the boundary is strongly non-linear — check by comparing against an RBF SVM offline; if the gap is large, keep the linear model but engineer features. (c) Any reasonable classifier, but scored by average precision with a baseline of $60/5000 = 0.012$, and with the threshold chosen from the 50 : 1 cost rather than left at 0.5. *Refuted if* recall at the chosen threshold is still near zero, which would mean the score carries no usable ranking — check the ROC AUC before blaming the threshold. (d) k -means for compactness or DBSCAN if

you expect noise and irregular shapes — after standardising, which was worth $0.3711 \rightarrow 0.8975$ on **wine**. *Refuted* if the silhouette is near zero at every k , or if the clustering changes completely under a different but equally defensible distance; either says the structure you are reporting is an artefact of your choices.

14. Report: “A 0.9412 ± 0.0201 , B 0.9388 ± 0.0187 over ten folds; the difference is $+0.0024$ with a paired-difference standard deviation of 0.0155 , so the two are indistinguishable on this evidence.” The standard error of the mean paired difference is $0.0155/\sqrt{10} = 0.0049$, so $t = 0.0024/0.0049 = 0.49$ — nowhere near significance. To claim A is better you would need either many more folds (the standard error falls as $1/\sqrt{k}$, so detecting 0.0024 at $t = 2$ needs about 167 folds) or a genuinely larger effect. In practice the right move is to stop trying to separate them and choose on cost, latency or interpretability.

15. *Ridge*: minimises squared error plus $\lambda\|\beta\|_2^2$; solved in closed form as $(X^\top X + \lambda I)^{-1}X^\top y$, which exists for every $\lambda > 0$; over the class of linear functions of x ; and you would know it worked from held-out R^2 , with λ chosen by cross-validation inside the fold. *Random forest*: minimises node impurity greedily, one split at a time, with no global objective; “how” is recursive partitioning on bootstrap samples with a random column subset at each split; the class is averages of axis-aligned piecewise-constant functions; and you would know from held-out accuracy or AUC, with out-of-bag error available free because each row is unused by about 37% of the trees. *DBSCAN*: there is **no answer to question one** — it minimises nothing. It is a definition (core, border and noise points, given ϵ and `min_samples`) plus a connectivity rule, not an optimisation. Question four has no answer either, in the strict sense: with no labels there is nothing to be right about, so the best available evidence is an internal index such as the silhouette, plus stability under perturbation, plus a stated justification for the distance. Those are arguments, not scores.

13 References and further reading

All links were checked on 24 August 2026. Free resources are marked [free]. This is a revision lecture, so the list below is deliberately broad rather than deep — one entry per part of the course.

Textbooks

- James, Witten, Hastie, Tibshirani and Taylor, *An Introduction to Statistical Learning with Python*. [free] The single best revision text for this course. Ch. 2 is the four-question frame of Section 2; Ch. 3 and 6 are Unit V; Ch. 4 and 8–9 are Unit VI; Ch. 5 is the cross-validation discipline of Section 6; Ch. 12 is Unit VII. <https://www.statlearning.com/>
- Han, Kamber and Pei, *Data Mining: Concepts and Techniques*, 3rd edition, Morgan Kaufmann, 2012. **The prescribed text**. Ch. 3 is Unit IV, Ch. 8–9 are classification and its evaluation, Ch. 10 is Unit VII. §8.5 is the confusion-matrix and ROC material of Section 8.
- Müller and Guido, *Introduction to Machine Learning with Python*, O’Reilly, 2016. The practical companion: Ch. 2 is the model zoo of Section 4, Ch. 3 is preprocessing and clustering, Ch. 5 is evaluation, Ch. 6 is pipelines — read Ch. 6 before your project submission. Notebooks [free]: https://github.com/amueller/introduction_to_ml_with_python
- Bishop, *Pattern Recognition and Machine Learning*, Springer, 2006. The derivations, when you want them stated properly: §1.5 loss, §3.1 linear regression and regularization, §4.3 logistic regression, Ch. 5 networks and backpropagation, §7.1 support vector machines.

- Hastie, Tibshirani and Friedman, *The Elements of Statistical Learning*, 2e, Springer. [free] §7 is the definitive treatment of why `best_score_` is optimistic and what nested cross-validation fixes — §7.10.2, “the wrong and right way to do cross-validation”, is Section 6.1 of these notes. <https://hastie.su.domains/ElemStatLearn/>
- Murphy, *Probabilistic Machine Learning: An Introduction*, MIT Press, 2022. [free] The unifying view of equation (1): Ch. 4 on estimation, Ch. 5 on decision theory — which is where thresholds and costs belong. <https://probml.github.io/pml-book/book1.html>
- Mitchell, *Machine Learning*, McGraw Hill, 1997. Still the clearest short account of ID3 and of the concept-learning framing. <https://www.cs.cmu.edu/~tom/mlbook.html>

Online courses

- **NPTEL, Introduction to Machine Learning** — Prof. Balaraman Ravindran, IIT Madras; the reference course for the semester, and its week structure maps onto these units almost exactly. [free] <https://nptel.ac.in/courses/106106139>
- **NPTEL, Introduction to Machine Learning** — Prof. Sudeshna Sarkar, IIT Kharagpur. A second treatment worked at the board, useful if a derivation did not land the first time. [free] <https://nptel.ac.in/courses/106105152>
- **NPTEL, Data Mining** — Prof. Pabitra Mitra, IIT Kharagpur. Follows the Han–Kamber–Pei chapter structure, so it is the right revision course for Units IV, VI and VII. [free] <https://nptel.ac.in/courses/106105174>
- **SWAYAM** — the national portal hosting NPTEL, with certification cycles. [free] <https://swayam.gov.in/>
- **Google Machine Learning Crash Course** — short, practical, and the classification module is a good half-hour on thresholds, precision, recall and ROC. [free] <https://developers.google.com/machine-learning/crash-course>
- **Kaggle Learn, Intro to Machine Learning** — the fastest way to rebuild the end-to-end habit of Section 7 if it has gone rusty. [free] <https://www.kaggle.com/learn/intro-to-machine-learning>

Videos

- StatQuest with Josh Starmer, *Machine Learning Fundamentals: Bias and Variance* — the clearest ten minutes on why fitting the training data better is not the objective. Watch before rereading regularization. [free] <https://youtu.be/EuBBz3bI-aA>
- StatQuest with Josh Starmer, *Machine Learning Fundamentals: Cross Validation* — the protocol of Section 4, drawn. [free] <https://youtu.be/fSyztGwwBVw>
- StatQuest with Josh Starmer, *ROC and AUC, Clearly Explained!* — the staircase of Section 8 built one threshold at a time. [free] <https://youtu.be/4jRBRDbJemM>
- StatQuest with Josh Starmer, *StatQuest: K-means clustering* — the Lloyd iteration you worked by hand, animated, including why `n_init` matters. [free] <https://youtu.be/4b5d3muPQmA>

Documentation and interactive explainers

- scikit-learn, *Common pitfalls and recommended practices* — read this once a year. It is Section 6 written by the library authors, with the data-leakage section stating explicitly that the transformer must be inside the pipeline. [free] https://scikit-learn.org/stable/common_pitfalls.html
- scikit-learn user guide, 3.1. *Cross-validation* and 3.4. *Metrics and scoring* — the definitive statement of what each splitter and each scorer computes, including the nested cross-validation example. [free] https://scikit-learn.org/stable/modules/model_evaluation.html
- scikit-learn example, *Classifier comparison* — the picture behind Section 5: ten classifiers on three synthetic problems, showing which boundaries each family can and cannot draw. [free] https://scikit-learn.org/stable/auto_examples/classification/plot_classifier_comparison.html
- scikit-learn user guide, 2.3. *Clustering* — the comparison table of which algorithm suits which geometry and which scales to what. [free] <https://scikit-learn.org/stable/modules/clustering.html>
- **gradkit** — the course’s own from-scratch automatic differentiation and neural-network library, in pure NumPy, with every derivative derived in a docstring beside the code that implements it and verified numerically in the test suite. If the backward pass of Section 8 still feels like magic, read docs/forward-and-backward.md and then src/gradkit/function.py. [free] <https://github.com/tali7c/gradkit>
- MLU-Explain, *Bias-Variance Tradeoff* and *Train, Test and Validation Sets* — two short interactive pages that make the spread of Section 4 visible. [free] <https://mlu-explain.github.io/bias-variance/>
- MLU-Explain, *ROC and AUC* — drag the threshold and watch the confusion matrix and the curve move together. [free] <https://mlu-explain.github.io/roc-auc/>
- R2D3, *A Visual Introduction to Machine Learning* — the best single page for explaining a decision tree, over-fitting and the train/test split to someone who is not on this course. Useful for the project quiz. [free] <https://r2d3.us/visual-intro-to-machine-learning-part-1/>

Articles

- L. Breiman, “Random Forests”, *Machine Learning* **45**(1), 5–32, 2001. The paper that introduced the column sampling of Table 1, and the source of the correlation-versus-strength argument that explains why individual trees get worse while the forest gets better. <https://doi.org/10.1023/A:1010933404324>
- P. J. Rousseeuw, “Silhouettes: a graphical aid to the interpretation and validation of cluster analysis”, *Journal of Computational and Applied Mathematics* **20**, 53–65, 1987. The index you compute by hand when there is no label column. [https://doi.org/10.1016/0377-0427\(87\)90125-7](https://doi.org/10.1016/0377-0427(87)90125-7)
- L. Hubert and P. Arabie, “Comparing partitions”, *Journal of Classification* **2**, 193–218, 1985. Why the adjusted Rand index is the one to quote: two random three-way labellings of iris score Rand 0.5567 and adjusted Rand –0.0010. <https://doi.org/10.1007/BF01908075>

Sources: Müller and Guido, *Introduction to Machine Learning with Python*, O'Reilly, 2016.; Bishop, *Pattern Recognition and Machine Learning*, Springer, 2006.; Murphy, *Machine Learning: A Probabilistic Perspective*, MIT Press, 2012.; Han, Kamber and Pei, *Data Mining: Concepts and Techniques* (3e), Morgan Kaufmann, 2012.